Package 'nimble'

October 9, 2021

Title MCMC, Particle Filtering, and Programmable Hierarchical Modeling

Description A system for writing hierarchical statistical models largely
compatible with 'BUGS' and 'JAGS', writing nimbleFunctions to operate models
and do basic R-style math, and compiling both models and nimbleFunctions via
custom-generated C++. 'NIMBLE' includes default methods for MCMC, Monte Carlo
Expectation Maximization, and some other tools. The nimbleFunction system makes
it easy to do things like implement new MCMC samplers from R, customize the
assignment of samplers to different parts of a model from R, and compile the
new samplers automatically via C++ alongside the samplers 'NIMBLE' provides.
'NIMBLE' extends the 'BUGS'/'JAGS' language by making it extensible: New
distributions and functions can be added, including as calls to external
compiled code. Although most people think of MCMC as the main goal of the
'BUGS'/'JAGS' language for writing models, one can use 'NIMBLE' for writing
arbitrary other kinds of model-generic algorithms as well. A full User Manual is
available at <https://r-nimble.org>.

Version 0.12.1

Date 2021-10-10

Maintainer Christopher Paciorek <paciorek@stat.berkeley.edu>

Depends R (>= 3.1.2)

Imports methods,igraph,coda,R6

Suggests testthat

URL https://r-nimble.org, https://github.com/nimble-dev/nimble

SystemRequirements GNU make

License BSD_3_clause + file LICENSE | GPL (>= 2)

Copyright See COPYRIGHTS file.

Note For convenience, the package includes the necessary header files
for the Eigen distribution. (This is all that is needed to use
that functionality.) You can use an alternative installation of
Eigen on your system or the one we provide. The license for the
Eigen code is very permissive and allows us to distribute it
with this package. See <http://eigen.tuxfamily.org/index.php?
title=Main_Page> and also the License section on that page.

Encoding UTF-8

Collate config.R all_utils.R options.R distributions_inputList.R
distributions_processInputList.R
distributions_implementations.R BUGS_BUGSdecl.R BUGS_contexts.R
BUGS_nimbleGraph.R BUGS_modelDef.R BUGS_model.R
BUGS_graphNodeMaps.R BUGS_readBUGS.R BUGS_macros.R
BUGS_testBUGS.R BUGS_getDependencies.R BUGS_utils.R
BUGS_mathCompatibility.R externalCalls.R genCpp_exprClass.R
genCpp_operatorLists.R genCpp_RparseTree2exprClasses.R
genCpp_initSizes.R genCpp_buildIntermediates.R
genCpp_processSpecificCalls.R genCpp_sizeProcessing.R
genCpp_toEigenize.R genCpp_insertAssertions.R genCpp_maps.R
genCpp_liftMaps.R genCpp_eigenization.R genCpp_addDebugMarks.R
genCpp_generateCpp.R RCfunction_core.R RCfunction_compile.R
nimbleFunction_util.R nimbleFunction_core.R
nimbleFunction_nodeFunction.R nimbleFunction_nodeFunctionNew.R
nimbleFunction_Rexecution.R nimbleFunction_compile.R
nimbleFunction_keywordProcessing.R nimbleList_core.R
types_util.R types_symbolTable.R types_modelValues.R
types_modelValuesAccessor.R types_modelVariableAccessor.R
types_nimbleFunctionList.R types_nodeFxnVector.R
types_numericLists.R cppDefs_utils.R cppDefs_variables.R
cppDefs_core.R cppDefs_namedObjects.R cppDefs_ADtools.R
cppDefs_BUGSmodel.R cppDefs_RCfunction.R
cppDefs_nimbleFunction.R cppDefs_nimbleList.R
cppDefs_modelValues.R cppDefs_cppProject.R
cppDefs_outputCppFromRparseTree.R cppInterfaces_utils.R
cppInterfaces_models.R cppInterfaces_modelValues.R
cppInterfaces_nimbleFunctions.R cppInterfaces_otherTypes.R
nimbleProject.R initializeModel.R CAR.R MCMC_utils.R
MCMC_configuration.R MCMC_build.R MCMC_run.R MCMC_samplers.R
MCMC_conjugacy.R MCMC_autoBlock.R MCMC_RJ.R MCMC_WAIC.R
MCEM_build.R crossValidation.R BNP_distributions.R
BNP_samplers.R NF_utils.R miscFunctions.R makevars.R
setNimbleInternalFunctions.R registration.R nimble-package.r
zzz.R

RoxygenNote 7.1.1

NeedsCompilation yes

Author Perry de Valpine [aut],
Christopher Paciorek [aut, cre],
Daniel Turek [aut],
Nick Michaud [aut],
Cliff Anderson-Bergman [aut],
Fritz Obermeyer [aut],
Claudia Wehrhahn Cortes [aut] (Bayesian nonparametrics system),
Abel Rodrìguez [aut] (Bayesian nonparametrics system),
Duncan Temple Lang [aut] (packaging configuration),

Sally Paganin [aut] (reversible jump MCMC),
Joshua Hug [aut] (WAIC),
Jagadish Babu [ctb] (code for the compilation system for an early
version of NIMBLE),
Lauren Ponisio [ctb] (contributions to the cross-validation code),
Peter Sujan [ctb] (multivariate t distribution code)

R **topics documented:**

ADNimbleList   *EXPERIMENTAL Data type for the return value of* nimDerivs

---

Description

nimbleList definition for the type of nimbleList returned by nimDerivs.

Usage

    ADNimbleList

Format

An object of class list of length 1.

Fields

value The value of the function evaluated at the given input arguments.

gradient The gradient of the function evaluated at the given input arguments.

hessian The Hessian of the function evaluated at the given input arguments.

thirdDerivs Currently unused.

See Also

nimDerivs

---

any_na   *Determine if any values in a vector are NA or NaN*

Description

NIMBLE language functions that can be used in either compiled or uncompiled nimbleFunctions
to detect if there are any NA or NaN values in a vector.

Usage

    any_na(x)

    any_nan(x)

Arguments

x   vector of values

Author(s)

NIMBLE Development Team

---

as.carAdjacency   *Convert CAR structural parameters to adjacency, weights, num format*

Description

This will convert alternate representations of CAR process structure into (adj, weights, num) form required by dcar_normal.

Usage

```
as.carAdjacency(...)
```

Arguments

...   Either: a symmetric matrix of unnormalized weights, or two lists specifying adjacency indices and the corresponding unnormalized weights.

Details

Two alternate representations are handled:

A single matrix argument will be interpreted as a matrix of symmetric unnormalized weights;

Two lists will be interpreted as (the first) a list of numeric vectors specifying the adjacency (neighboring) indices of each CAR process component, and (the second) a list of numeric vectors giving the unnormalized weights for each of these neighboring relationships.

Author(s)

Daniel Turek

See Also

CAR-Normal

---

as.carCM   *Convert weights vector to parameters of* dcar_proper *distributio*

Description

Convert weights vector to `C` and `M` parameters of dcar_proper distribution

Usage

```
as.carCM(adj, weights, num)
```

Arguments

`adj`	vector of indices of the adjacent locations (neighbors) of each spatial location. This is a sparse representation of the full adjacency matrix.
`weights`	vector of symmetric unnormalized weights associated with each pair of adjacent locations, of the same length as adj. This is a sparse representation of the full (unnormalized) weight matrix.
`num`	vector giving the number of neighbors of each spatial location, with length equal to the total number of locations.

Details

Given a symmetric matrix of unnormalized weights, this function will calculate corresponding values for the `C` and `M` arguments suitable for use in the dcar_proper distribution. This function can be used to transition between usage of dcar_normal and dcar_proper, since dcar_normal uses the `adj`, `weights`, and `num` arguments, while dcar_proper requires adj, num, and also the `C` and `M` as returned by this function.

Here, `C` is a sparse vector representation of the row-normalized adjacency matrix, and `M` is a vector containing the conditional variance for each region. The resulting values of `C` and `M` are guaranteed to satisfy the symmetry constraint imposed on C and M, that $M^{-1}C$ is symmetric, where M is a diagonal matrix and C is the row-normalized adjacency matrix.

Value

A named list with elements `C` and `M`. These may be used as the `C` and `M` arguments to the dcar_proper distribution.

Author(s)

Daniel Turek

See Also

[CAR-Normal](), [CAR-Proper]()

---

asRow	*Turn a numeric vector into a single-row or single-column matrix*

Description

Turns a numeric vector into a matrix that has 1 row or 1 column. Part of NIMBLE language.

Usage

```
asRow(x)

asCol(x)
```

Arguments

x	Numeric to be turned into a single row or column matrix

Details

In the NIMBLE language, some automatic determination of how to turn vectors into single-row or single-column matrices is done. For example, in A %*% x, where A is a matrix and x a vector, x will be turned into a single-column matrix unless it is known at compile time that A is a single column, in which case x will be turned into a single-row matrix. However, if it is desired that x be turned into a single row but A cannot be determined at compile time to be a single column, then one can use A %*% asRow(x) to force this conversion.

Author(s)

Perry de Valpine

---

autoBlock	*Automated parameter blocking procedure for efficient MCMC sampling*

Description

The automated parameter blocking algorithm is no longer actively maintained. In some cases, it may not operate correctly with more recent system features and/or distributions.

Usage

```
autoBlock(
  Rmodel,
  autoIt = 20000,
  run = list("all", "default"),
  setSeed = TRUE,
  verbose = FALSE,
  makePlots = FALSE,
  round = TRUE
)
```

Arguments

Rmodel	A NIMBLE model object, created from `nimbleModel`.
autoIt	The number of MCMC iterations to run intermediate MCMC algorithms, through the course of the procedure. Default 20,000.
run	List of additional MCMC algorithms to compare against the automated blocking MCMC. These may be specified as: the character string 'all' to denote blocking all continuous-valued nodes; the character string 'default' to denote NIMBLE's default MCMC configuration; a named list element consisting of a quoted code block, which when executed returns an MCMC configuration object for comparison; a custom-specificed blocking scheme, specified as a named list element which itself is a list of character vectors, where each character vector specifies the nodes in a particular block. Default is c('all', 'default').
setSeed	Logical specificying whether to call set.seed(0) prior to beginning the blocking procedure. Default TRUE.
verbose	Logical specifying whether to output considerable details of the automated block procedure, through the course of execution. Default FALSE.
makePlots	Logical specifying whether to plot the hierarchical clustering dendrograms, through the course of execution. Default FALSE.
round	Logical specifying whether to round the final output results to two decimal places. Default TRUE.

Details

Runs NIMBLE's automated blocking procedure for a given model object, to dynamically determine a blocking scheme of the continuous-valued model nodes. This blocking scheme is designed to produce efficient MCMC sampling (defined as number of effective samples generated per second of algorithm runtime). See Turek, et al (2015) for details of this algorithm. This also (optionally) compares this blocked MCMC against several static MCMC algorithms, including all univariate sampling, blocking of all continuous-valued nodes, NIMBLE's default MCMC configuration, and custom-specified blockings of parameters.

This method allows for fine-tuned usage of the automated blocking procedure. However, the main entry point to the automatic blocking procedure is intended to be through either buildMCMC(..., autoBlock = TRUE), or configureMCMC(..., autoBlock = TRUE).

Value

Returns a named list containing elements:

- summary: A data frame containing a numerical summary of the performance of all MCMC algorithms (including that from automated blocking)
- autoGroups: A list specifying the parameter blockings converged on by the automated blocking procedure
- conf: A NIMBLE MCMC configuration object corresponding to the results of the automated blocking procedure

Author(s)

Daniel Turek

References

Turek, D., de Valpine, P., Paciorek, C., and Anderson-Bergman, C. (2015). Automated Parameter Blocking for Efficient Markov-Chain Monte Carlo Sampling. <arXiv:1503.05621>.

See Also

configureMCMC buildMCMC

BUGSdeclClass-class   *BUGSdeclClass contains the information extracted from one BUGS declaration*

Description

BUGSdeclClass contains the information extracted from one BUGS declaration

---

buildAuxiliaryFilter   *Placeholder for buildAuxiliaryFilter*

Description

This function has been moved to the 'nimbleSMC' package

Usage

```
buildAuxiliaryFilter(...)
```

Arguments

...   arguments

---

buildBootstrapFilter   *Placeholder for buildBootstrapFilter*

Description

This function has been moved to the 'nimbleSMC' package

Usage

```
buildBootstrapFilter(...)
```

Arguments

...   arguments

---

buildEnsembleKF   *Placeholder for buildEnsembleKF*

Description

This function has been moved to the 'nimbleSMC' package

Usage

```
buildEnsembleKF(...)
```

Arguments

...   arguments

---

buildIteratedFilter2   *Placeholder for buildIteratedFilter2*

Description

This function has been moved to the 'nimbleSMC' package

Usage

```
buildIteratedFilter2(...)
```

Arguments

...   arguments

---

buildLiuWestFilter   *Placeholder for buildLiuWestFilter*

Description

This function has been moved to the 'nimbleSMC' package

Usage

```
buildLiuWestFilter(...)
```

Arguments

...	arguments

---

buildMCEM   *Builds an MCEM algorithm from a given NIMBLE model*

Description

Takes a NIMBLE model and builds an MCEM algorithm for it. The user must specify which latent nodes are to be integrated out in the E-Step. All other stochastic non-data nodes will be maximized over. If the nodes do not have positive density on the entire real line, then box constraints can be used to enforce this. The M-step is done by a nimble MCMC sampler. The E-step is done by a call to R's optim with method = 'L-BFGS-B' if the nodes are constrainted, or method = 'BFGS' if the nodes are unconstrained.

Usage

```
buildMCEM(
  model,
  latentNodes,
  burnIn = 500,
  mcmcControl = list(adaptInterval = 100),
  boxConstraints = list(),
  buffer = 10^-6,
  alpha = 0.25,
  beta = 0.25,
  gamma = 0.05,
  C = 0.001,
  numReps = 300,
  forceNoConstraints = FALSE,
  verbose = TRUE
)
```

Arguments

model	a nimble model
latentNodes	character vector of the names of the stochastic nodes to integrated out. Names can be expanded, but don't need to be. For example, if the model contains x[1],x[2] and x[3] then one could provide either latentNodes = c('x[1]','x[2]','x[3]') or latentNodes = 'x'.
burnIn	burn-in used for MCMC sampler in E step
mcmcControl	list passed to configureMCMC, which builds the MCMC sampler. See help(configureMCMC) for more details
boxConstraints	list of box constraints for the nodes that will be maximized over. Each constraint is a list in which the first element is a character vector of node names to which the constraint applies and the second element is a vector giving the lower and upper limits. Limits of -Inf or Inf are allowed. Any nodes that are not given constrains will have their constraints automatically determined by NIMBLE
buffer	A buffer amount for extending the boxConstraints. Many functions with boundary constraints will produce NaN or -Inf when parameters are on the boundary. This problem can be prevented by shrinking the boundary a small amount.
alpha	probability of a type one error - here, the probability of accepting a parameter estimate that does not increase the likelihood. Default is 0.25.
beta	probability of a type two error - here, the probability of rejecting a parameter estimate that does increase the likelihood. Default is 0.25.
gamma	probability of deciding that the algorithm has converged, that is, that the difference between two Q functions is less than C, when in fact it has not. Default is 0.05.
C	determines when the algorithm has converged - when C falls above a (1-gamma) confidence interval around the difference in Q functions from time point t-1 to time point t, we say the algorithm has converged. Default is 0.001.
numReps	number of bootstrap samples to use for asymptotic variance calculation.
forceNoConstraints	
	avoid any constraints even from parameter bounds implicit in the model structure (e.g., from dunif or dgamma distributions); setting this to TRUE might allow MCEM to run when the bounds of a parameter being maximized over depend on another parameter.
verbose	logical indicating whether to print additional logging information.

Details

buildMCEM calls the NIMBLE compiler to create the MCMC and objective function as nimbleFunctions. If the given model has already been used in compiling other nimbleFunctions, it is possible you will need to create a new copy of the model for buildMCEM to use. Uses an ascent-based MCEM algorithm, which includes rules for automatically increasing the number of MC samples as iterations increase, and for determining when convergence has been reached. Constraints for parameter values can be provided. If constraints are not provided, they will be automatically determined by NIMBLE. Initial values for the parameters are taken to be the values in the model at the time buildMCEM is called, unless the values in the compiled model are changed before running the MCEM.

Value

an R list with two elements:

- run A function that when called runs the MCEM algorithm. This function takes the arguments listed in run Arguments below.
- estimateCov An EXPERIMENTAL function that when called estimates the asymptotic covariance of the parameters. The covariance is estimated using the method of Louis (1982). This function takes the arguments listed in estimateCov Arguments below.

run **Arguments**

- initM starting number of iterations for the algorithm.

estimateCov **Arguments**

- MLEs named vector of MLE values. Must have a named MLE value for each stochastic, non-data, non-latent node. If the run() method has alread been called, MLEs do not need to be provided.
- useExistingSamples logical argument. If TRUE and the run() method has previously been called, the covariance estimation will use MCMC samples from the last step of the MCEM algorithm. Otherwise, an MCMC algorithm will be run for 10,000 iterations, and those samples will be used. Defaults to FALSE.

Author(s)

Clifford Anderson-Bergman and Nicholas Michaud

References

Caffo, Brian S., Wolfgang Jank, and Galin L. Jones (2005). Ascent-based Monte Carlo expectation-maximization. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 67(2), 235-251.

Louis, Thomas A (1982). Finding the Observed Information Matrix When Using the EM Algorithm. *Journal of the Royal Statistical Society. Series B (Statistical Methodology)*, 44(2), 226-233.

Examples

```
## Not run:
pumpCode <- nimbleCode({
 for (i in 1:N){
     theta[i] ~ dgamma(alpha,beta);
     lambda[i] <- theta[i]*t[i];
     x[i] ~ dpois(lambda[i])
 }
 alpha ~ dexp(1.0);
 beta ~ dgamma(0.1,1.0);
})

pumpConsts <- list(N = 10,
             t = c(94.3, 15.7, 62.9, 126, 5.24,
```

```
                31.4, 1.05, 1.05, 2.1, 10.5))

pumpData <- list(x = c(5, 1, 5, 14, 3, 19, 1, 1, 4, 22))

pumpInits <- list(alpha = 1, beta = 1,
              theta = rep(0.1, pumpConsts$N))
pumpModel <- nimbleModel(code = pumpCode, name = 'pump', constants = pumpConsts,
                   data = pumpData, inits = pumpInits)

# Want to maximize alpha and beta (both which must be positive) and integrate over theta
box = list( list(c('alpha','beta'), c(0, Inf)))

pumpMCEM <- buildMCEM(model = pumpModel, latentNodes = 'theta[1:10]',
                       boxConstraints = box)
MLEs <- pumpMCEM$run(initM = 1000)
cov <- pumpMCEM$estimateCov()

## End(Not run)
# Could also use latentNodes = 'theta' and buildMCEM() would figure out this means 'theta[1:10]'
```

---

buildMCMC   *Create an MCMC function from a NIMBLE model, or an MCMC con-*
                             figuration object

Description

First required argument, which may be of class MCMCconf (an MCMC configuration object), or
inherit from class modelBaseClass (a NIMBLE model object). Returns an uncompiled executable
MCMC function. See details.

Usage

```
buildMCMC(conf, ...)
```

Arguments

conf   An MCMC configuration object of class MCMCconf that specifies the model,
              samplers, monitors, and thinning intervals for the resulting MCMC function.
              See configureMCMC for details of creating MCMC configuration objects. Alter-
              natively, conf may a NIMBLE model object, in which case an MCMC function
              corresponding to the default MCMC configuration for this model is returned.

...   Additional arguments to be passed to configureMCMC if conf is a NIMBLE
              model object (see help(configureMCMC)).

Details

Calling buildMCMC(conf) will produce an uncompiled MCMC function object. The uncompiled MCMC function will have arguments:

niter: The number of iterations to run the MCMC.

thin: The thinning interval for the `monitors` that were specified in the MCMC configuration. If this argument is provided at MCMC runtime, it will take precedence over the `thin` interval that was specified in the MCMC configuration. If omitted, the `thin` interval from the MCMC configuration will be used.

thin2: The thinning interval for the second set of monitors (`monitors2`) that were specified in the MCMC configuration. If this argument is provided at MCMC runtime, it will take precedence over the `thin2` interval that was specified in the MCMC configuration. If omitted, the `thin2` interval from the MCMC configuration will be used.

reset: Boolean specifying whether to reset the internal MCMC sampling algorithms to their initial state (in terms of self-adapting tuning parameters), and begin recording posterior sample chains anew. Specifying `reset = FALSE` allows the MCMC algorithm to continue running from where it left off, appending additional posterior samples to the already existing sample chains. Generally, `reset = FALSE` should only be used when the MCMC has already been run (default = TRUE).

resetMV: Boolean specifying whether to begin recording posterior sample chains anew. This argument is only considered when using `reset = FALSE`. Specifying `reset = FALSE, resetMV = TRUE` allows the MCMC algorithm to continue running from where it left off, but without appending the new posterior samples to the already existing samples, i.e. all previously obtained samples will be erased. This option can help reduce memory usage during burn-in (default = FALSE).

nburnin: Number of initial, pre-thinning, MCMC iterations to discard (default = 0).

time: Boolean specifying whether to record runtimes of the individual internal MCMC samplers. When `time = TRUE`, a vector of runtimes (measured in seconds) can be extracted from the MCMC using the method `mcmc$getTimes()` (default = FALSE).

progressBar: Boolean specifying whether to display a progress bar during MCMC execution (default = TRUE). The progress bar can be permanently disabled by setting the system option `nimbleOptions(MCMCprogressBar = FALSE)`.

Samples corresponding to the `monitors` and `monitors2` from the MCMCconf are stored into the interval variables mvSamples and mvSamples2, respectively. These may be accessed and converted into R matrix or list objects via: `as.matrix(mcmc$mvSamples) as.list(mcmc$mvSamples) as.matrix(mcmc$mvSamples2) as.list(mcmc$mvSamples2)`

The uncompiled MCMC function may be compiled to a compiled MCMC object, taking care to compile in the same project as the R model object, using: `Cmcmc <-compileNimble(Rmcmc,project = Rmodel)`

The compiled function will function identically to the uncompiled object, except acting on the compiled model object.

Calculating WAIC

Please see `help(waic)` for more information.

Author(s)

Daniel Turek

References

Watanabe, S. (2010). Asymptotic equivalence of Bayes cross validation and widely applicable information criterion in singular learning theory. *Journal of Machine Learning Research* 11: 3571-3594.

Gelman, A., Hwang, J. and Vehtari, A. (2014). Understanding predictive information criteria for Bayesian models. *Statistics and Computing* 24(6): 997-1016.

Ariyo, O., Quintero, A., Munoz, J., Verbeke, G. and Lesaffre, E. (2019). Bayesian model selection in linear mixed models for longitudinal data. *Journal of Applied Statistics* 47: 890-913.

See Also

configureMCMC runMCMC nimbleMCMC

Examples

```
## Not run:
code <- nimbleCode({
    mu ~ dnorm(0, 1)
    x ~ dnorm(mu, 1)
    y ~ dnorm(x, 1)
})
Rmodel <- nimbleModel(code, data = list(y = 0))
conf <- configureMCMC(Rmodel, monitors = c('mu', 'x'), enableWAIC = TRUE)
Rmcmc <- buildMCMC(conf)
Cmodel <- compileNimble(Rmodel)
Cmcmc <- compileNimble(Rmcmc, project=Rmodel)
Cmcmc$run(10000)
samples <- as.matrix(Cmcmc$mvSamples)
samplesAsList <- as.list(Cmcmc$mvSamples)
head(samples)
waicInfo <- Cmcmc$getWAIC()
waicInfo$WAIC
waicInfo$pWAIC

## End(Not run)
```

---

calculateWAIC   *Calculating WAIC using an offline algorithm*

Description

In addition to the core online algorithm, NIMBLE implements an offline WAIC algorithm that can be computed on the results of an MCMC. In contrast to NIMBLE's built-in online WAIC, offline WAIC can compute only conditional WAIC and does not allow for grouping data nodes.

Usage

```
calculateWAIC(mcmc, model, nburnin = 0, thin = 1)
```

Arguments

mcmc	An MCMC object (compiled or uncompiled) or matrix or dataframe of MCMC samples as the first argument of calculateWAIC.
model	A model (compiled or uncompiled) as the second argument of calculateWAIC. Only required if mcmc is a matrix/dataframe of samples.
nburnin	The number of pre-thinning MCMC samples to remove from the beginning of the posterior samples for offline WAIC calculation via calculateWAIC (default = 0). These samples are discarded in addition to any burn-in specified when running the MCMC.
thin	Thinning factor interval to apply to the samples for offline WAIC calculation using calculateWAIC (default = 1, corresponding to no thinning).

Details

The ability to calculate WAIC post hoc after all MCMC sampling has been done has certain advantages (e.g., allowing a user to exclude additional burnin samples beyond that specified initially for the MCMC) in addition to providing compatibility with versions of NIMBLE before 0.12.0. This functionality includes the ability to call the calculateWAIC function on an MCMC object or matrix of samples after running an MCMC and without setting up the MCMC initially to use WAIC, provided the necessary variables are monitored.

See help(waic) for details on using NIMBLE's recommended online algorithm for WAIC.

Offline WAIC (WAIC computed after MCMC sampling)

As an alternative to online WAIC, NIMBLE also provides a function, calculateWAIC, that can be called on an MCMC object or a matrix of samples, after running an MCMC. This function does not require that one set enableWAIC = TRUE nor WAIC = TRUE when calling runMCMC. The function checks that the necessary variables were monitored in the MCMC and returns an error if they were not. This function behaves identically to the calculateWAIC method of an MCMC object. Note that this function cannot be used when using nimbleMCMC as the underlying MCMC object is not available to the user in that case.

The calculateWAIC function requires either an MCMC object or a matrix (or dataframe) of posterior samples plus a model object. In addition, one can provide optional burnin and thin arguments.

In addition, for compatibility with older versions of NIMBLE (prior to v0.12.0), one can also use the calculateWAIC method of the MCMC object to calculate WAIC after all sampling has been completed.

The calculateWAIC() method accepts a single argument, nburnin, equivalent to the nburnin argument of the calculateWAIC function described above.

The calculateWAIC method can only be used if the enableWAIC argument to configureMCMC or to buildMCMC is set to TRUE, or if the NIMBLE option enableWAIC is set to TRUE. If a user attempts to call calculateWAIC without having set enableWAIC = TRUE (either in the call to configureMCMC, or buildMCMC, or as a NIMBLE option), an error will occur.

The calculateWAIC function and method calculate the WAIC based on Equations 5, 12, and 13 in Gelman et al. (2014) (i.e., using *p*WAIC2).

Note that there is not a unique value of WAIC for a model. The calculateWAIC function and method only provide the conditional WAIC, namely the version of WAIC where all parameters

directly involved in the likelihood are treated as $theta$ for the purposes of Equation 5 from Gelman et al. (2014). As a result, the user must set the MCMC monitors (via the `monitors` argument) to include all stochastic nodes that are parents of any data nodes; by default the MCMC monitors are only the top-level nodes of the model. For more detail on the use of different predictive distributions, see Section 2.5 from Gelman et al. (2014) or Ariyo et al. (2019). Also note that WAIC relies on a partition of the observations, i.e., 'pointwise' prediction. In `calculateWAIC` the sum over log pointwise predictive density values treats each data node as contributing a single value to the sum. When a data node is multivariate, that data node contributes a single value to the sum based on the joint density of the elements in the node. Note that if one wants the WAIC calculation via `calculateWAIC` to be based on the joint predictive density for each group of observations (e.g., grouping the observations from each person or unit in a longitudinal data context), one would need to use a multivariate distribution for the observations in each group (potentially by writing a user-defined distribution).

For more control over and flexibility in how WAIC is calculated, see `help(waic)`.

Author(s)

Joshua Hug and Christopher Paciorek

References

Watanabe, S. (2010). Asymptotic equivalence of Bayes cross validation and widely applicable information criterion in singular learning theory. *Journal of Machine Learning Research* 11: 3571-3594.

Gelman, A., Hwang, J. and Vehtari, A. (2014). Understanding predictive information criteria for Bayesian models. *Statistics and Computing* 24(6): 997-1016.

Ariyo, O., Quintero, A., Munoz, J., Verbeke, G. and Lesaffre, E. (2019). Bayesian model selection in linear mixed models for longitudinal data. *Journal of Applied Statistics* 47: 890-913.

Vehtari, A., Gelman, A. and Gabry, J. (2017). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing* 27: 1413-1432.

Hug, J.E. and Paciorek, C.J. (2021). A numerically stable online implementation and exploration of WAIC through variations of the predictive density, using NIMBLE. *arXiv e-print* <arXiv:2106.13359>.

See Also

waic configureMCMC buildMCMC runMCMC nimbleMCMC

Examples

```
code <- nimbleCode({
  for(j in 1:J) {
    for(i in 1:n)
      y[j, i] ~ dnorm(mu[j], sd = sigma)
    mu[j] ~ dnorm(mu0, sd = tau)
  }
  tau ~ dunif(0, 10)
  sigma ~ dunif(0, 10)
})
J <- 5
```

```
n <- 10
y <- matrix(rnorm(J*n), J, n)
Rmodel <- nimbleModel(code, constants = list(J = J, n = n), data = list(y = y),
                      inits = list(tau = 1, sigma = 1))

## Make sure the needed variables are monitored.
## Only conditional WAIC without data grouping is available via this approach.
conf <- configureMCMC(Rmodel, monitors = c('mu', 'sigma'))
## Not run:
Cmodel <- compileNimble(Rmodel)
Rmcmc <- buildMCMC(conf)
Cmcmc <- compileNimble(Rmcmc, project = Rmodel)
output <- runMCMC(Cmcmc, niter = 1000)
calculateWAIC(Cmcmc)            # Can run on the MCMC object
calculateWAIC(output, Rmodel)  # Can run on the samples directly

## Apply additional burnin (additional to any burnin already done in the MCMC.
calculateWAIC(Cmcmc, burnin = 500)

## End(Not run)
```

CAR-Normal   *The CAR-Normal Distribution*

Description

Density function and random generation for the improper (intrinsic) Gaussian conditional autoregressive (CAR) distribution.

Usage

```
dcar_normal(
  x,
  adj,
  weights = adj/adj,
  num,
  tau,
  c = CAR_calcNumIslands(adj, num),
  zero_mean = 0,
  log = FALSE
)

rcar_normal(
  n = 1,
  adj,
  weights = adj/adj,
  num,
  tau,
  c = CAR_calcNumIslands(adj, num),
```

```
    zero_mean = 0
)
```

Arguments

x	vector of values.
adj	vector of indices of the adjacent locations (neighbors) of each spatial location. This is a sparse representation of the full adjacency matrix.
weights	vector of symmetric unnormalized weights associated with each pair of adjacent locations, of the same length as adj. If omitted, all weights are taken to be one.
num	vector giving the number of neighboring locations of each spatial location, with length equal to the total number of locations.
tau	scalar precision of the Gaussian CAR prior.
c	integer number of constraints to impose on the improper density function. If omitted, c is calculated as the number of disjoint groups of spatial locations in the adjacency structure, which implicitly assumes a first-order CAR process for each group. Note that c should be equal to the number of eigenvalues of the precision matrix that are zero. For example, if the neighborhood structure is based on a second-order Markov random field in one dimension then the matrix has two zero eigenvalues and in two dimensions it has three zero eigenvalues. See Rue and Held (2005) and the NIMBLE User Manual for more information.
zero_mean	integer specifying whether to set the mean of all locations to zero during MCMC sampling of a node specified with this distribution in BUGS code (default 0). This argument is used only in BUGS model code when specifying models in NIMBLE. If 0, the overall process mean is included implicitly in the value of each location in a BUGS model; if 1, then during MCMC sampling, the mean of all locations is set to zero at each MCMC iteration, and a separate intercept term should be included in the BUGS model. Note that centering during MCMC as implemented in NIMBLE follows the ad hoc approach of **WinBUGS** and does not sample under the constraint that the mean is zero as discussed on p. 36 of Rue and Held (2005). See 'Details'.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations.

Details

When specifying a CAR distribution in BUGS model code, the zero_mean parameter should be specified as either 0 or 1 (rather than TRUE or FALSE).

Note that because the distribution is improper, rcar_normal does not generate a sample from the distribution. However, as discussed in Rue and Held (2005), it is possible to generate a sample from the distribution under constraints imposed based on the eigenvalues of the precision matrix that are zero.

Value

dcar_normal gives the density, while rcar_normal returns the current process values, since this distribution is improper.

Author(s)

Daniel Turek

References

Banerjee, S., Carlin, B.P., and Gelfand, A.E. (2015). *Hierarchical Modeling and Analysis for Spatial Data*, 2nd ed. Chapman and Hall/CRC.

Rue, H. and L. Held (2005). *Gaussian Markov Random Fields*, Chapman and Hall/CRC.

See Also

CAR-Proper, Distributions for other standard distributions

Examples

```
x <- c(1, 3, 3, 4)
num <- c(1, 2, 2, 1)
adj <- c(2, 1,3, 2,4, 3)
weights <- c(1, 1, 1, 1, 1, 1)
lp <- dcar_normal(x, adj, weights, num, tau = 1)
```

CAR-Proper   *The CAR-Proper Distribution*

Description

Density function and random generation for the proper Gaussian conditional autoregressive (CAR) distribution.

Usage

```
dcar_proper(
  x,
  mu,
  C = CAR_calcC(adj, num),
  adj,
  num,
  M = CAR_calcM(num),
  tau,
  gamma,
  evs = CAR_calcEVs3(C, adj, num),
  log = FALSE
)

rcar_proper(
  n = 1,
  mu,
```

```
    C = CAR_calcC(adj, num),
    adj,
    num,
    M = CAR_calcM(num),
    tau,
    gamma,
    evs = CAR_calcEVs3(C, adj, num)
)
```

Arguments

x	vector of values.
mu	vector of the same length as x, specifying the mean for each spatial location.
C	vector of the same length as adj, giving the weights associated with each pair of neighboring locations. See 'Details'.
adj	vector of indices of the adjacent locations (neighbors) of each spatial location. This is a sparse representation of the full adjacency matrix.
num	vector giving the number of neighboring locations of each spatial location, with length equal to the number of locations.
M	vector giving the diagonal elements of the conditional variance matrix, with length equal to the number of locations. See 'Details'.
tau	scalar precision of the Gaussian CAR prior.
gamma	scalar representing the overall degree of spatial dependence. See 'Details'.
evs	vector of eigenvalues of the adjacency matrix implied by C, adj, and num. This parameter should not be provided; it will always be calculated using the adjacency information.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations.

Details

If both C and M are omitted, then all weights are taken as one, and corresponding values of C and M are generated.

The C and M parameters must jointly satisfy a symmetry constraint: that M^(-1) %*% C is symmetric, where M is a diagonal matrix and C is the full weight matrix that is sparsely represented by the parameter vector C.

For a proper CAR model, the value of gamma must lie within the inverse minimum and maximum eigenvalues of M^(-0.5) %*% C %*% M^(0.5), where M is a diagonal matrix and C is the full weight matrix. These bounds can be calculated using the deterministic functions carMinBound(C,adj,num,M) and carMaxBound(C,adj,num,M), or simultaneously using carBounds(C,adj,num,M). In the case where C and M are omitted (all weights equal to one), the bounds on gamma are necessarily (-1, 1).

Value

dcar_proper gives the density, and rcar_proper generates random deviates.

Author(s)

Daniel Turek

References

Banerjee, S., Carlin, B.P., and Gelfand, A.E. (2015). *Hierarchical Modeling and Analysis for Spatial Data*, 2nd ed. Chapman and Hall/CRC.

See Also

[CAR-Normal,](#) [Distributions](#) for other standard distributions

Examples

```
x <- c(1, 3, 3, 4)
mu <- rep(3, 4)
adj <- c(2, 1,3, 2,4, 3)
num <- c(1, 2, 2, 1)

## omitting C and M uses all weights = 1
dcar_proper(x, mu, adj = adj, num = num, tau = 1, gamma = 0.95)

## equivalent to above: specifying all weights = 1,
## then using as.carCM to generate C and M arguments
weights <- rep(1, 6)
CM <- as.carCM(adj, weights, num)
C <- CM$C
M <- CM$M
dcar_proper(x, mu, C, adj, num, M, tau = 1, gamma = 0.95)

## now using non-unit weights
weights <- c(2, 2, 3, 3, 4, 4)
CM2 <- as.carCM(adj, weights, num)
C2 <- CM2$C
M2 <- CM2$M
dcar_proper(x, mu, C2, adj, num, M2, tau = 1, gamma = 0.95)
```

---

carBounds	*Calculate bounds for the autocorrelation parameter of the* `dcar_proper` *distribution*

Description

Calculate the lower and upper bounds for the `gamma` parameter of the `dcar_proper` distribution

Usage

```
carBounds(C, adj, num, M)
```

Arguments

C	vector of the same length as adj, giving the normalized weights associated with each pair of neighboring locations.
adj	vector of indices of the adjacent locations (neighbors) of each spatial location. This is a sparse representation of the full adjacency matrix.
num	vector giving the number of neighboring locations of each spatial location, with length equal to the number of locations.
M	vector giving the diagonal elements of the conditional variance matrix, with length equal to the number of locations.

Details

Bounds for gamma are the inverse of the minimum and maximum eigenvalues of: $M^{(-0.5)}CM^{(0.5)}$. The lower and upper bounds are returned in a numeric vector.

Value

A numeric vector containing the bounds (minimum and maximum allowable values) for the gamma parameter of the dcar_proper distribution.

Author(s)

Daniel Turek

See Also

[CAR-Proper](), [carMinBound](), [carMaxBound]()

---

carMaxBound	*Calculate the upper bound for the autocorrelation parameter of the* dcar_proper *distribution*

Description

Calculate the upper bound for the gamma parameter of the dcar_proper distribution

Usage

```
carMaxBound(C, adj, num, M)
```

Arguments

C	vector of the same length as `adj`, giving the normalized weights associated with each pair of neighboring locations.
adj	vector of indices of the adjacent locations (neighbors) of each spatial location. This is a sparse representation of the full adjacency matrix.
num	vector giving the number of neighboring locations of each spatial location, with length equal to the number of locations.
M	vector giving the diagonal elements of the conditional variance matrix, with length equal to the number of locations.

Details

Bounds for gamma are the inverse of the minimum and maximum eigenvalues of $M^{(-0.5)}CM^{(0.5)}$.

Value

The upper bound (maximum allowable value) for the gamma parameter of the `dcar_proper` distribution.

Author(s)

Daniel Turek

See Also

[CAR-Proper](), [carMinBound](), [carBounds]()

---

carMinBound	*Calculate the lower bound for the autocorrelation parameter of the* `dcar_proper` *distribution*

Description

Calculate the lower bound for the gamma parameter of the `dcar_proper` distribution

Usage

```
carMinBound(C, adj, num, M)
```

Arguments

C	vector of the same length as `adj`, giving the normalized weights associated with each pair of neighboring locations.
adj	vector of indices of the adjacent locations (neighbors) of each spatial location. This is a sparse representation of the full adjacency matrix.

num	vector giving the number of neighboring locations of each spatial location, with length equal to the number of locations.
M	vector giving the diagonal elements of the conditional variance matrix, with length equal to the number of locations.

Details

Bounds for gamma are the inverse of the minimum and maximum eigenvalues of: $M^{(-0.5)}CM^{(0.5)}$.

Value

The lower bound (minimum allowable value) for the gamma parameter of the dcar_proper distribution.

Author(s)

Daniel Turek

See Also

[CAR-Proper](), [carMaxBound](), [carBounds]()

CAR_calcNumIslands   *Calculate number of islands based on a CAR adjacency matrix.*

Description

Calculate number of islands (distinct connected groups) based on a CAR adjacency matrix.

Usage

```
CAR_calcNumIslands(adj, num)
```

Arguments

adj	vector of indices of the adjacent locations (neighbors) of each spatial location. This is a sparse representation of the full adjacency matrix.
num	vector giving the number of neighbors of each spatial location, with length equal to the total number of locations.

Author(s)

Daniel Turek

See Also

[CAR-Normal]()

Categorical *The Categorical Distribution*

Description

Density and random generation for the categorical distribution

Usage

```
dcat(x, prob, log = FALSE)

rcat(n = 1, prob)
```

Arguments

x	non-negative integer-value numeric value.
prob	vector of probabilities, internally normalized to sum to one.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations.

Details

See the BUGS manual for mathematical details.

Value

dcat gives the density and rcat generates random deviates.

Author(s)

Christopher Paciorek

See Also

Distributions for other standard distributions

Examples

```
probs <- c(1/4, 1/10, 1 - 1/4 - 1/10)
x <- rcat(n = 30, probs)
dcat(x, probs)
```

---

checkInterrupt   *Check for interrupt (e.g. Ctrl-C) during nimbleFunction execution.*
                                  Part of the NIMBLE language.

Description

Check for interrupt (e.g. Ctrl-C) during nimbleFunction execution. Part of the NIMBLE language.

Usage

```
checkInterrupt()
```

Details

During execution of nimbleFunctions that take a long time, it is nice to occassionally check if the user has entered an interrupt and bail out of execution if so. This function does that. During uncompiled nimbleFunction execution, it does nothing. During compiled execution, it calls R_checkUserInterrupt() of the R headers.

Author(s)

Perry de Valpine

ChineseRestaurantProcess

                              The Chinese Restaurant Process Distribution

---

Description

Density and random generation for the Chinese Restaurant Process distribution.

Usage

```
dCRP(x, conc = 1, size, log = 0)

rCRP(n, conc = 1, size)
```

Arguments

x	vector of values.
conc	scalar concentration parameter.
size	integer-valued length of x (required).
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations (only n = 1 is handled currently).

Details

The Chinese restaurant process distribution is a distribution on the space of partitions of the positive integers. The distribution with concentration parameter α equal to conc has probability function

$$f(x_i \mid x_1, \ldots, x_{i-1}) = \frac{1}{i-1+\alpha} \sum_{j=1}^{i-1} \delta_{x_j} + \frac{\alpha}{i-1+\alpha} \delta_{x^{new}},$$

where x^{new} is a new integer not in $x_1, \ldots, x_{i-1}$.

If conc is not specified, it assumes the default value of 1. The conc parameter has to be larger than zero. Otherwise, NaN are returned.

Value

dCRP gives the density, and rCRP gives random generation.

Author(s)

Claudia Wehrhahn

References

Blackwell, D., and MacQueen, J. B. (1973). Ferguson distributions via P\'olya urn schemes. *The Annals of Statistics*, 1: 353-355.

Aldous, D. J. (1985). Exchangeability and related topics. In \'*Ecole d'\'Et\'e de Probabilit\'es de Saint-Flour XIII - 1983* (pp. 1-198). Springer, Berlin, Heidelberg.

Pitman, J. (1996). Some developments of the Blackwell-MacQueen urn scheme. *IMS Lecture Notes-Monograph Series*, 30: 245-267.

Examples

```
x <- rCRP(n=1, conc = 1, size=10)
dCRP(x, conc = 1, size=10)
```

CmodelBaseClass-class   *Class* CmodelBaseClass

Description

Classes used internally in NIMBLE and not expected to be called directly by users.

CnimbleFunctionBase-class

Class CnimbleFunctionBase

Description

Classes used internally in NIMBLE and not expected to be called directly by users.

---

codeBlockClass-class   *Class* codeBlockClass

Description

Classes used internally in NIMBLE and not expected to be called directly by users.

---

compareMCMCs   *Placeholder for compareMCMCs*

Description

This function has been moved to a separate package

Usage

```
compareMCMCs(...)
```

Arguments

`...`   arguments

---

compileNimble   *compile NIMBLE models and nimbleFunctions*

Description

compile a collection of models and nimbleFunctions: generate C++, compile the C++, load the
result, and return an interface object

Usage

```
compileNimble(
  ...,
  project,
  dirName = NULL,
  projectName = "",
  control = list(),
  resetFunctions = FALSE,
  showCompilerOutput = nimbleOptions("showCompilerOutput")
)
```

Arguments

`...`	An arbitrary set of NIMBLE models and nimbleFunctions, or lists of them. If given as named parameters, those names may be used in the return list.
`project`	Optional NIMBLE model or nimbleFunction already associated with a project, which the current units for compilation should join. If not provided, a new project will be created and the current compilation units will be associated with it.
`dirName`	Optional directory name in which to generate the C++ code. If not provided, a temporary directory will be generated using R's `tempdir` function.
`projectName`	Optional character name for labeling the project if it is new
`control`	A list mostly for internal use. See details.
`resetFunctions`	Logical value stating whether nimbleFunctions associated with an existing project should all be reset for compilation purposes. See details.
`showCompilerOutput`	
	Logical value indicating whether details of C++ compilation should be printed.

Details

This is the main function for calling the NIMBLE compiler. A set of compiler calls and output will be seen. Compiling in NIMBLE does 4 things: 1. It generates C++ code files for all the model and nimbleFunction components. 2. It calls the system's C++ compiler. 3. It loads the compiled object(s) into R using `dyn.load`. And 4. it generates R objects for using the compiled model and nimbleFunctions.

When the units for compilation provided in `...` include multiple models and/or nimbleFunctions, models are compiled first, in the order in which they are provided. Groups of nimbleFunctions that were specialized from the same nimbleFunction generator (the result of a call to `nimbleFunction`, which then takes setup arguments and returns a specialized nimbleFunction) are then compiled as a group, in the order of first appearance.

The behavior of adding new compilation units to an existing project is limited. For example, one can compile a model in one call to `compileNimble` and then compile a nimbleFunction that uses the model (i.e. was given the model as a setup argument) in a second call to `compileNimble`, with the model provided as the `project` argument. Either the uncompiled or compiled model can be provided. However, compiling a second nimbleFunction and adding it to the same project will only work in limited circumstances. Basically, the limitations occur because it attempts to re-use already compiled pieces, but if these do not have all the necessary information for the new compilation, it gives up. An attempt has been made to give up in a controlled manner and provide somewhat informative messages.

When compilation is not allowed or doesn't work, try using `resetFunctions = TRUE`, which will force recompilation of all nimbleFunctions in the new call. Previously compiled nimbleFunctions will be unaffected, and their R interface objects should continue to work. The only cost is additional compilation time for the current compilation call. If that doesn't work, try re-creating the model and/or the nimbleFunctions from their generators. An alternative possible fix is to compile multiple units in one call, rather than sequentially in multiple calls.

The control list can contain the following named elements, each with `TRUE` or `FALSE`: debug, which sets a debug mode for the compiler for development purposes; debugCpp, which inserts an output message before every line of C++ code for debugging purposes; compileR, which determines

whether the R-only steps of compilation should be executed; writeCpp, which determines whether the C++ files should be generated; compileCpp, which determines whether the C++ should be compiled; loadSO, which determines whether the DLL or shared object should be loaded and interfaced; and returnAsList, which determines whether calls to the compiled nimbleFunction should return only the returned value of the call (returnAsList = FALSE) or whether a list including the input arguments, possibly modified, should be returned in a list with the returned value of the call at the end (returnAsList = TRUE). The control list is mostly for developer use, although returnAsArgs may be useful to a user. An example of developer use is that one can have the compiler write the C++ files but not compile them, then modify them by hand, then have the C++ compiler do the subsequent steps without over-writing the files.

See the NIMBLE User Manual Manual for examples

Value

If there is only one compilation unit (one model or nimbleFunction), an R interface object is returned. This object can be used like the uncompiled model or nimbleFunction, but execution will call the corresponding compiled objects or functions. If there are multiple compilation units, they will be returned as a list of interface objects, in the order provided. If names were included in the arguments, or in a list if any elements of ... are lists, those names will be used for the corresponding element of the returned list. Otherwise an attempt will be made to generate names from the argument code. For example compileNimble(A = fun1, B = fun2, project = myModel) will return a list with named elements A and B, while compileNimble(fun1, fun2, project = myModel) will return a list with named elements fun1 and fun2.

Author(s)

Perry de Valpine

---

configureMCMC   *Build the MCMCconf object for construction of an MCMC object*

Description

Creates a defaut MCMC configuration for a given model.

Usage

```
configureMCMC(
  model,
  nodes,
  control = list(),
  monitors,
  thin = 1,
  monitors2 = character(),
  thin2 = 1,
  useConjugacy = getNimbleOption("MCMCuseConjugacy"),
  onlyRW = FALSE,
```

```
    onlySlice = FALSE,
  multivariateNodesAsScalars = getNimbleOption("MCMCmultivariateNodesAsScalars"),
    enableWAIC = getNimbleOption("MCMCenableWAIC"),
    controlWAIC = list(),
    print = getNimbleOption("verbose"),
    autoBlock = FALSE,
    oldConf,
    warnNoSamplerAssigned = TRUE,
    ...
)
```

Arguments

model	A NIMBLE model object, created from [nimbleModel](#)
nodes	An optional character vector, specifying the nodes and/or variables for which samplers should be created. Nodes may be specified in their indexed form, y[1,3]. Alternatively, nodes specified without indexing will be expanded fully, e.g., x will be expanded to x[1], x[2], etc. If missing, the default value is all non-data stochastic nodes. If NULL, then no samplers are added.
control	An optional list of control arguments to sampler functions. If a control list is provided, the elements will be provided to all sampler functions which utilize the named elements given. For example, the standard Metropolis-Hastings random walk sampler ([sampler_RW](#)) utilizes control list elements adaptive, adaptInterval, and scale. (Internally it also uses targetNode, but this should not generally be provided as a control list element). The default values for control list arguments for samplers (if not otherwise provided as an argument to configureMCMC()) are in the setup code of the sampling algorithms.
monitors	A character vector of node names or variable names, to record during MCMC sampling. This set of monitors will be recorded with thinning interval thin, and the samples will be stored into the mvSamples object. The default value is all top-level stochastic nodes of the model – those having no stochastic parent nodes.
thin	The thinning interval for monitors. Default value is one.
monitors2	A character vector of node names or variable names, to record during MCMC sampling. This set of monitors will be recorded with thinning interval thin2, and the samples will be stored into the mvSamples2 object. The default value is an empty character vector, i.e. no values will be recorded.
thin2	The thinning interval for monitors2. Default value is one.
useConjugacy	A logical argument, with default value TRUE. If specified as FALSE, then no conjugate samplers will be used, even when a node is determined to be in a conjugate relationship.
onlyRW	A logical argument, with default value FALSE. If specified as TRUE, then Metropolis-Hastings random walk samplers ([sampler_RW](#)) will be assigned for all non-terminal continuous-valued nodes nodes. Discrete-valued nodes are assigned a slice sampler ([sampler_slice](#)), and terminal nodes are assigned a posterior_predictive sampler ([sampler_posterior_predictive](#)).

onlySlice   A logical argument, with default value FALSE. If specified as TRUE, then a
                         slice sampler is assigned for all non-terminal nodes. Terminal nodes are still
                         assigned a posterior_predictive sampler.

multivariateNodesAsScalars

                         A logical argument, with default value FALSE. If specified as TRUE, then non-
                         terminal multivariate stochastic nodes will have scalar samplers assigned to each
                         of the scalar components of the multivariate node. The default value of FALSE
                         results in a single block sampler assigned to the entire multivariate node. Note,
                         multivariate nodes appearing in conjugate relationships will be assigned the cor-
                         responding conjugate sampler (provided useConjugacy == TRUE), regardless of
                         the value of this argument.

enableWAIC   A logical argument, specifying whether to enable WAIC calculations for the re-
                         sulting MCMC algorithm. Defaults to the value of nimbleOptions('MCMCenableWAIC'),
                         which in turn defaults to FALSE. Setting nimbleOptions('enableWAIC' = TRUE)
                         will ensure that WAIC is enabled for all calls to configureMCMC and buildMCMC.

controlWAIC   A named list of inputs that control the behavior of the WAIC calculation. See
                         help(waic).

print   A logical argument, specifying whether to print the ordered list of default sam-
                         plers.

autoBlock   A logical argument specifying whether to use an automated blocking procedure
                         to determine blocks of model nodes for joint sampling. If TRUE, an MCMC
                         configuration object will be created and returned corresponding to the results of
                         the automated parameter blocking. Default value is FALSE.

oldConf   An optional MCMCconf object to modify rather than creating a new MCMC-
                         conf from scratch

warnNoSamplerAssigned

                         A logical argument, with default value TRUE. This specifies whether to issue a
                         warning when no sampler is assigned to a node, meaning there is no matching
                         sampler assignment rule.

...   Additional named control list elements for default samplers, or additional argu-
                         ments to be passed to the autoBlock function when autoBlock = TRUE

Details

See MCMCconf for details on how to manipulate the MCMCconf object

Author(s)

Daniel Turek

See Also

buildMCMC runMCMC nimbleMCMC

configureRJ   *Configure Reversible Jump for Variable Selection*

Description

Modifies an MCMC configuration object to perform a reversible jump MCMC sampling for variable selection, using a univariate normal proposal distribution. Users can control the mean and scale of the proposal. This function supports two different types of model specification: with and without indicator variables.

Usage

```
configureRJ(
  conf,
  targetNodes,
  indicatorNodes = NULL,
  priorProb = NULL,
  control = list(mean = NULL, scale = NULL, fixedValue = NULL)
)
```

Arguments

conf	An `MCMCconf` object.
targetNodes	A character vector, specifying the nodes and/or variables for which variable selection is to be performed. Nodes may be specified in their indexed form, `'y[1,3]'`. Alternatively, nodes specified without indexing will be expanded, e.g., `'x'` will be expanded to `'x[1]'`, `'x[2]'`, etc.
indicatorNodes	An optional character vector, specifying the indicator nodes and/or variables paired with `targetNodes`. Nodes may be specified in their indexed form, `'y[1,3]'`. Alternatively, nodes specified without indexing will be expanded, e.g., `'x'` will be expanded to `'x[1]'`, `'x[2]'`, etc. Nodes must be provided consistently with `targetNodes`. See details.
priorProb	An optional value or vector of prior probabilities for each node to be in the model. See details.
control	An optional list of control arguments: • mean. The mean of the normal proposal distribution (default = 0). • scale. The standard deviation of the normal proposal distribution (default = 1). • fixedValue. Value for the variable when it is out of the model, which can be used only when `priorProb` is provided (default = 0). If specified when `indicatorNodes` is passed, a warning is given and `fixedValue` is ignored.

Details

This function modifies the samplers in MCMC configuration object for each of the nodes provided in the `targetNodes` argument. To these elements two samplers are assigned: a reversible jump sampler to transition the variable in/out of the model, and a modified version of the original sampler, which performs updates only when the target node is already in the model.

`configureRJ` can handle two different ways of writing a NIMBLE model, either with or without indicator variables. When using indicator variables, the `indicatorNodes` argument must be provided. Without indicator variables, the `priorProb` argument must be provided. In the latter case, the user can provide a non-zero value for `fixedValue` if desired.

Note that this functionality is intended for variable selection in regression-style models but may be useful for other situations as well. At the moment, setting a variance component to zero and thereby removing a set of random effects that are explicitly part of a model will not work because MCMC sampling in that case would need to propose values for multiple parameters (the random effects), whereas the current functionality only proposes adding/removing a single model node.

Value

`NULL` `configureRJ` modifies the input MCMC configuration object in place.

Author(s)

Sally Paganin, Perry de Valpine, Daniel Turek

References

Peter J. Green. (1995). Reversible jump Markov chain Monte Carlo computation and Bayesian model determination. *Biometrika*, 82(4), 711-732.

See Also

`samplers configureMCMC`

Examples

```
## Not run:

## Linear regression with intercept and two covariates, using indicator variables

code <- nimbleCode({
  beta0 ~ dnorm(0, sd = 100)
  beta1 ~ dnorm(0, sd = 100)
  beta2 ~ dnorm(0, sd = 100)
  sigma ~ dunif(0, 100)
  z1 ~ dbern(psi)   ## indicator variable associated with beta1
  z2 ~ dbern(psi)   ## indicator variable associated with beta2
  psi ~ dunif(0, 1) ## hyperprior on inclusion probability
  for(i in 1:N) {
    Ypred[i] <- beta0 + beta1 * z1 * x1[i] + beta2 * z2 * x2[i]
    Y[i] ~ dnorm(Ypred[i], sd = sigma)
```

```
  }
})

## simulate some data
set.seed(1)
N <- 100
x1 <- runif(N, -1, 1)
x2 <- runif(N, -1, 1) ## this covariate is not included
Y <- rnorm(N, 1 + 2.5 * x1, sd = 1)

## build the model
rIndicatorModel <- nimbleModel(code, constants = list(N = N),
                               data = list(Y = Y, x1 = x1, x2 = x2),
                               inits = list(beta0 = 0, beta1 = 0, beta2 = 0, sigma = sd(Y),
                                 z1 = 1, z2 = 1, psi = 0.5))

indicatorModelConf <- configureMCMC(rIndicatorModel)

## Add reversible jump
configureRJ(conf = indicatorModelConf,         ## model configuration
            targetNodes = c("beta1", "beta2"), ## coefficients for selection
            indicatorNodes = c("z1", "z2"),    ## indicators paired with coefficients
            control = list(mean = 0, scale = 2))

indicatorModelConf$addMonitors("beta1", "beta2", "z1", "z2")

rIndicatorMCMC <- buildMCMC(indicatorModelConf)
cIndicatorModel <- compileNimble(rIndicatorModel)
cIndicatorMCMC <- compileNimble(rIndicatorMCMC, project = rIndicatorModel)

set.seed(1)
samples <- runMCMC(cIndicatorMCMC, 10000, nburnin = 6000)

## posterior probability to be included in the mode
mean(samples[ , "z1"])
mean(samples[ , "z2"])

## posterior means when in the model
mean(samples[ , "beta1"][samples[ , "z1"] != 0])
mean(samples[ , "beta2"][samples[ , "z2"] != 0])


## Linear regression with intercept and two covariates, without indicator variables

code <- nimbleCode({
  beta0 ~ dnorm(0, sd = 100)
  beta1 ~ dnorm(0, sd = 100)
  beta2 ~ dnorm(0, sd = 100)
  sigma ~ dunif(0, 100)
  for(i in 1:N) {
    Ypred[i] <- beta0 + beta1 * x1[i] + beta2 * x2[i]
    Y[i] ~ dnorm(Ypred[i], sd = sigma)
  }
```

```
})

rNoIndicatorModel <- nimbleModel(code, constants = list(N = N),
                                 data = list(Y = Y, x1 = x1, x2 = x2),
                                 inits= list(beta0 = 0, beta1 = 0, beta2 = 0, sigma = sd(Y)))

noIndicatorModelConf <- configureMCMC(rNoIndicatorModel)

## Add reversible jump
configureRJ(conf = noIndicatorModelConf,       ## model configuration
            targetNodes = c("beta1", "beta2"), ## coefficients for selection
            priorProb = 0.5,                   ## prior probability of inclusion
            control = list(mean = 0, scale = 2))

## add monitors
noIndicatorModelConf$addMonitors("beta1", "beta2")
rNoIndicatorMCMC <- buildMCMC(noIndicatorModelConf)

cNoIndicatorModel <- compileNimble(rNoIndicatorModel)
cNoIndicatorMCMC <- compileNimble(rNoIndicatorMCMC, project = rNoIndicatorModel)

set.seed(1)
samples <- runMCMC(cNoIndicatorMCMC, 10000, nburnin = 6000)

## posterior probability to be included in the mode
mean(samples[ , "beta1"] != 0)
mean(samples[ , "beta2"] != 0)

## posterior means when in the model
mean(samples[ , "beta1"][samples[ , "beta1"] != 0])
mean(samples[ , "beta2"][samples[ , "beta2"] != 0])

## End(Not run)
```

Constraint   *Constraint calculations in NIMBLE*

Description

Calculations to handle censoring

Usage

```
dconstraint(x, cond, log = FALSE)

rconstraint(n = 1, cond)
```

Arguments

x	value indicating whether cond is TRUE or FALSE
cond	logical value
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations (only n=1 is handled currently).

Details

Used for working with constraints in BUGS code. See the NIMBLE manual for additional details.

Value

dconstraint gives the density and rconstraint generates random deviates, but these are unusual as the density is 1 if x matches cond and 0 otherwise and the deviates are simply the value of cond

Author(s)

Christopher Paciorek

See Also

[Distributions](#) for other standard distributions

Examples

```
constr <- 3 > 2 && 4 > 0
x <- rconstraint(1, constr)
dconstraint(x, constr)
dconstraint(0, 3 > 4)
dconstraint(1, 3 > 4)
rconstraint(1, 3 > 4)
```

---

decide	*Makes the Metropolis-Hastings acceptance decision, based upon the input (log) Metropolis-Hastings ratio*

Description

This function returns a logical TRUE/FALSE value, indicating whether the proposed transition should be accepted (TRUE) or rejected (FALSE).

Usage

```
decide(logMetropolisRatio)
```

Arguments

logMetropolisRatio

> The log of the Metropolis-Hastings ratio, which is calculated from model prob-
> abilities and forward/reverse transition probabilities. Calculated as the ratio of
> the model probability under the proposal to that under the current values mul-
> tiplied by the ratio of the reverse transition probability to the forward transition
> probability.

Details

The Metropolis-Hastings accept/reject decisions is made as follows. If `logMetropolisRatio` is
greater than 0, accept (return `TRUE`). Otherwise draw a uniform random number between 0 and 1
and accept if it is less that `exp(logMetropolisRatio`. The proposed transition will be rejected
(return `FALSE`). If `logMetropolisRatio` is NA, NaN, or -Inf, a reject (`FALSE`) decision will be
returned.

Author(s)

Daniel Turek

---

decideAndJump   *Creates a nimbleFunction for executing the Metropolis-Hastings*
                                 jumping decision, and updating values in the model, or in a carbon
                                 copy modelValues object, accordingly.

Description

This nimbleFunction generator must be specialized to three required arguments: a model, a model-
Values, and a character vector of node names.

Usage

```
decideAndJump(model, mvSaved, target, calcNodes)
```

Arguments

model	An uncompiled or compiled NIMBLE model object.
mvSaved	A modelValues object containing identical variables and logProb variables as the model. Can be created by `modelValues(model)`.
target	A character vector providing the target node.
calcNodes	A character vector representing a set of nodes in the model (and hence also the modelValues) object.

Details

Calling decideAndJump(model, mvSaved, calcNodes) will generate a specialized nimbleFunction with four required numeric arguments:

modelLP1: The model log-probability associated with the newly proposed value(s)

modelLP0: The model log-probability associated with the original value(s)

propLP1: The log-probability associated with the proposal forward-transiton

propLP0: The log-probability associated with the proposal reverse-tranisiton

Executing this function has the following effects: – Calculate the (log) Metropolis-Hastings ratio, as logMHR = modelLP1 - modelLP0 - propLP1 + propLP0 – Make the proposal acceptance decision based upon the (log) Metropolis-Hastings ratio – If the proposal is accepted, the values and associated logProbs of all calcNodes are copied from the model object into the mvSaved object – If the proposal is rejected, the values and associated logProbs of all calcNodes are copied from the mvSaved object into the model object – Return a logical value, indicating whether the proposal was accepted

Author(s)

Daniel Turek

---

declare   *Explicitly declare a variable in run-time code of a nimbleFunction*

Description

Explicitly declare a variable in run-time code of a nimbleFunction, for cases when its dimensions cannot be inferred before it is used. Works in R and NIMBLE.

Usage

```
declare(name, def)
```

Arguments

name   Name of a variable to declare, without quotes

def   NIMBLE type declaration, of the form TYPE(nDim,sizes), where TYPE is integer, double, or logical, nDim is the number of dimensions, and sizes is an optional vector of sizes concatenated with c. If nDim is omitted, it defaults to 0, indicating a scalar. If sizes are provided, they should not be changed subsequently in the function, including by assignment. Omitting nDim results in a scalar. For logical, only scalar is currently supported.

Details

In a run-time function of a nimbleFunction (either the run function or a function provided in methods when calling nimbleFunction), the dimensionality and numeric type of a variable is inferred when possible from the statement first assigning into it. E.g. A <-B + C infers that A has numeric types, dimensions and sizes taken from B + C. However, if the first appearance of A is e.g. A[i] <-5, A must have been explicitly declared. In this case, declare(A,double(1)) would make A a 1-dimensional (i.e. vector) double.

When sizes are not set, they can be set by a call to setSize or by assignment to the whole object. Sizes are not automatically extended if assignment is made to elements beyond the current sizes. In compiled nimbleFunctions doing so can cause a segfault and crash the R session.

This part of the NIMBLE language is needed for compilation, but it also runs in R. When run in R, is works by the side effect of creating or modifying name in the calling environment.

Author(s)

NIMBLE development team

Examples

```
declare(A, logical())          ## scalar logical, the only kind allowed
declare(B, integer(2, c(10, 10))) ## 10 x 10 integer matrix
declare(C, double(3))          ## 3-dimensional double array with no sizes set.
```

---

deregisterDistributions

Remove user-supplied distributions from use in NIMBLE BUGS models

---

Description

Deregister distributional information originally supplied by the user for use in BUGS model code

Usage

```
deregisterDistributions(distributionsNames)
```

Arguments

distributionsNames

a character vector giving the names of the distributions to be dergistered

Author(s)

Christopher Paciorek

Dirichlet	*The Dirichlet Distribution*

Description

Density and random generation for the Dirichlet distribution

Usage

```
ddirch(x, alpha, log = FALSE)

rdirch(n = 1, alpha)
```

Arguments

x	vector of values.
alpha	vector of parameters of same length as x
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations (only n=1 is handled currently).

Details

See Gelman et al., Appendix A or the BUGS manual for mathematical details.

Value

ddirch gives the density and rdirch generates random deviates.

Author(s)

Christopher Paciorek

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

[Distributions](#) for other standard distributions

Examples

```
alpha <- c(1, 10, 30)
x <- rdirch(1, alpha)
ddirch(x, alpha)
```

---

distributionInfo   *Get information about a distribution*

Description

Give information about each BUGS distribution

Usage

```
getDistributionInfo(dist)

isUserDefined(dist)

pqDefined(dist)

getType(
  dist,
  params = NULL,
  valueOnly = is.null(params) && !includeParams,
  includeParams = !is.null(params)
)

getParamNames(dist, includeValue = TRUE)
```

Arguments

dist	a character vector of length one, giving the name of the distribution (as used in BUGS code), e.g. `'dnorm'`
params	an optional character vector of names of parameters for which dimensions are desired (possibly including \'value\' and alternate parameters)
valueOnly	a logical indicating whether to only return the dimension of the value of the node
includeParams	a logical indicating whether to return dimensions of parameters. If TRUE and \'params\' is NULL then dimensions of all parameters, including the dimension of the value of the node, are returned
includeValue	a logical indicating whether to return the string 'value', which is the name of the node value

Details

NIMBLE provides various functions to give information about a BUGS distribution. In some cases, functions of the same name and similar functionality operate on the node(s) of a model as well (see `help(modelBaseClass)`).

`getDistributionInfo` returns an internal data structure (a reference class object) providing various information about the distribution. The output is not very user-friendly, but does contain all of the information that NIMBLE has about the distribution.

isDiscrete tests if a BUGS distribution is a discrete distribution.

isUserDefined tests if a BUGS distribution is a user-defined distribution.

pqAvail tests if a BUGS distribution provides distribution ('p') and quantile ('q') functions.

getDimension provides the dimension of the value and/or parameters of a BUGS distribution. The return value is a numeric vector with an element for each parameter/value requested.

getType provides the type (numeric, logical, integer) of the value and/or parameters of a BUGS distribution. The return value is a character vector with an element for each parameter/value requested. At present, all quantities are stored as numeric (double) values, so this function is of little practical use but could be exploited in the future.

getParamNames provides the value and/or parameter names of a BUGS distribution.

Author(s)

Christopher Paciorek

Examples

```
distInfo <- getDistributionInfo('dnorm')
distInfo
distInfo$range

isDiscrete('dbin')

isUserDefined('dbin')

pqDefined('dgamma')
pqDefined('dmnorm')

getDimension('dnorm')
getDimension('dnorm', includeParams = TRUE)
getDimension('dnorm', c('var', 'sd'))
getDimension('dcat', includeParams = TRUE)
getDimension('dwish', includeParams = TRUE)

getType('dnorm')
getType('dnorm', includeParams = TRUE)
getType('dnorm', c('var', 'sd'))
getType('dcat', includeParams = TRUE)
getType('dwish', includeParams = TRUE)

getParamNames('dnorm', includeValue = FALSE)
getParamNames('dmnorm')
```

Double-Exponential   *The Double Exponential (Laplace) Distribution*

Description

Density, distribution function, quantile function and random generation for the double exponential
distribution, allowing non-zero location, `mu`, and non-unit scale, `sigma`, or non-unit rate, `tau`

Usage

```
ddexp(x, location = 0, scale = 1, rate = 1/scale, log = FALSE)

rdexp(n, location = 0, scale = 1, rate = 1/scale)

pdexp(
  q,
  location = 0,
  scale = 1,
  rate = 1/scale,
  lower.tail = TRUE,
  log.p = FALSE
)

qdexp(
  p,
  location = 0,
  scale = 1,
  rate = 1/scale,
  lower.tail = TRUE,
  log.p = FALSE
)
```

Arguments

x	vector of values.
location	vector of location values.
scale	vector of scale values.
rate	vector of inverse scale values.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations.
q	vector of quantiles.
lower.tail	logical; if TRUE (default) probabilities are $P[X \leq x]$; otherwise, $P[X > x]$.
log.p	logical; if TRUE, probabilities p are given by user as log(p).
p	vector of probabilities.

Details

See Gelman et al., Appendix A or the BUGS manual for mathematical details.

Value

ddexp gives the density, pdexp gives the distribution function, qdexp gives the quantile function, and rdexp generates random deviates.

Author(s)

Christopher Paciorek

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

[Distributions](#) for other standard distributions

Examples

```
x <- rdexp(50, location = 2, scale = 1)
ddexp(x, 2, 1)
```

---

eigenNimbleList   *eigenNimbleList definition*

Description

nimbleList definition for the type of nimbleList returned by [nimEigen](#).

Usage

```
eigenNimbleList
```

Format

An object of class list of length 1.

Author(s)

NIMBLE development team

See Also

[nimEigen](#)

Exponential   *The Exponential Distribution*

Description

Density, distribution function, quantile function and random generation for the exponential distribution with rate (i.e., mean of `1/rate`) or scale parameterizations.

Usage

```
dexp_nimble(x, rate = 1/scale, scale = 1, log = FALSE)

rexp_nimble(n = 1, rate = 1/scale, scale = 1)

pexp_nimble(q, rate = 1/scale, scale = 1, lower.tail = TRUE, log.p = FALSE)

qexp_nimble(p, rate = 1/scale, scale = 1, lower.tail = TRUE, log.p = FALSE)
```

Arguments

x	vector of values.
rate	vector of rate values.
scale	vector of scale values.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations.
q	vector of quantiles.
lower.tail	logical; if TRUE (default) probabilities are $P[X \leq x]$; otherwise, $P[X > x]$.
log.p	logical; if TRUE, probabilities p are given by user as log(p).
p	vector of probabilities.

Details

NIMBLE's exponential distribution functions use Rmath's functions under the hood, but are parameterized to take both rate and scale and to use 'rate' as the core parameterization in C, unlike Rmath, which uses 'scale'. See Gelman et al., Appendix A or the BUGS manual for mathematical details.

Value

`dexp_nimble` gives the density, `pexp_nimble` gives the distribution function, `qexp_nimble` gives the quantile function, and `rexp_nimble` generates random deviates.

Author(s)

Christopher Paciorek

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

[Distributions](#) for other standard distributions

Examples

```
x <- rexp_nimble(50, scale = 3)
dexp_nimble(x, scale = 3)
```

---

extractControlElement   *Extract named elements from MCMC sampler control list*

Description

Extract named elements from MCMC sampler control list

Usage

```
extractControlElement(controlList, elementName, defaultValue, error)
```

Arguments

controlList	control list object, which is passed as an argument to all MCMC sampler setup functions.
elementName	character string, giving the name of the element to be extracted from the control list.
defaultValue	default value of the control list element, giving the value to be used when the elementName does not exactly match the name of an element in the controlList.
error	character string, giving the error message to be printed if no defaultValue is provided and elementName does not match the name of an element in the controlList.

Value

The element of controlList whose name matches elementName. If no controlList name matches elementName, then defaultValue is returned.

Author(s)

Daniel Turek

flat   *The Improper Uniform Distribution*

Description

Improper flat distribution for use as a prior distribution in BUGS models

Usage

```
dflat(x, log = FALSE)

rflat(n = 1)

dhalfflat(x, log = FALSE)

rhalfflat(n = 1)
```

Arguments

x	vector of values.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations.

Value

dflat gives the pseudo-density value of 1, while rflat and rhalfflat return NaN, since one cannot simulate from an improper distribution. Similarly, dhalfflat gives a pseudo-density value of 1 when x is non-negative.

Author(s)

Christopher Paciorek

See Also

Distributions for other standard distributions

Examples

```
dflat(1)
```

---

getBound   *Get value of bound of a stochastic node in a model*

Description

Get the value of the lower or upper bound for a single stochastic node in a model.

Usage

```
getBound(model, node, bound, nodeFunctionIndex)
```

Arguments

model   A NIMBLE model object

node   The name of a stochastic node in the model

bound   Either 'lower' or 'upper' indicating the desired bound for the node

nodeFunctionIndex
                    For internal NIMBLE use only

Details

Standard usage is as a method of a model, in the form model$getBound(node,bound), but the usage as a simple function with the model as the first argument as above is also allowed.

For nodes that do not involve truncation of the distribution this will return the lower or upper bound of the distribution, which may be a constant or for a limited number of distributions a parameter or functional of a parameter (at the moment in NIMBLE, the only case where a bound is a parameter is for the uniform distribution. For nodes that are truncated, this will return the desired bound, which may be a functional of other quantities in the model or may be a constant.

---

getBUGSexampleDir   *Get the directory path to one of the classic BUGS examples installed*
                             with NIMBLE package

Description

NIMBLE comes with some of the classic BUGS examples. getBUGSexampleDir looks up the location of an example from its name.

Usage

```
getBUGSexampleDir(example)
```

Arguments

example   The name of the classic BUGS example.

Value

Character string of the fully pathed directory of the BUGS example.

Author(s)

Christopher Paciorek

See Also

[readBUGSmodel](#) for usage in creating a model from a classic BUGS example

---

getConditionallyIndependentSets

Get a list of conditionally independent sets of nodes in a nimble model

---

Description

Conditionally independent sets of nodes are typically groups of latent states whose joint probability (density) will not change even if any other non-fixed node is changed. Default fixed nodes are data nodes and parameter nodes (with no parent nodes), but this can be controlled.

Usage

```
getConditionallyIndependentSets(
  model,
  nodes,
  givenNodes,
  omit = integer(),
  inputType = c("latent", "param", "data"),
  stochOnly = TRUE,
  returnType = "names",
  returnScalarComponents = FALSE
)
```

Arguments

model   A nimble model object (uncompiled or compiled).

nodes   A vector of node names or their graph IDs that are the starting nodes from which
              conditionally independent sets of nodes should be found. If omitted, the default
              will be all latent nodes, defined as stochastic nodes that are not data and have at
              least one stochastic parent node (possible with determinstic nodes in between).
              Note that this will omit latent states that have no hyperparameters. An example
              is the first latent state in some state-space (time-series) models, which is some-
              times declared with known prior. See type because it relates to the interpretation
              of nodes.

givenNodes	A vector of node names or their graph IDs that should be considered as fixed and hence can be conditioned on. If omitted, the default will be all data nodes and all parameter nodes, the latter defined as nodes with no stochastic parent nodes (skipping over deterministic parent nodes).
omit	A vector of node names or their graph IDs that should be omitted and should block further graph exploration.
inputType	The method of graph exploration depends on what the nodes argument represents. For latent, the input nodes are interpreted as latent states, from which both parent and descendent graph exploration should be done to find nodes in the same set (nodes that are NOT conditionally independent from each other). For param, the input nodes are interpreted as parameters, so graph exploration begins from the top (input) and explores descendents. For data, the input nodes are interpreted as data nodes, so graph exploration begins from the bottom (input) explores parent nodes.
stochOnly	Logical for whether only stochastic nodes should be returned (default = TRUE). If FALSE, both deterministic and stochastic nodes are returned.
returnType	Either names for returned nodes to be node names or ids for returned nodes to be graph IDs.
returnScalarComponents	
	If FALSE (default), multivariate nodes are returned as full names (e.g. x[1:3]). If TRUE, they are returned as scalar elements (e.g. x[1], x[2], x[3]).

Details

This function returns sets of conditionally independent nodes. Multiple input nodes might be in the same set or different sets, and other nodes (not in nodes) will be included.

By default, deterministic dependencies of givenNodes are also counted as given nodes. This is relevant only for parent nodes. This allows the givenNodes to include only stochastic nodes. Say we have A -> B -> C -> D. A and D are givenNodes. C is a latent node. B is a deterministic node. By default, B is considered given. Otherwise, other branches that depend on B would be grouped in the same output set as C, but this is usually what is wanted. Any use of the resulting output must ensure that B is calculated when necessary, as usual with nimble's model-generic programming. To turn off this feature, set nimbleOptions(groupDetermWithGivenInCondIndSets = FALSE).

Value

List of nodes that are in conditionally independent sets. With each set, nodes are returned in topologically sorted order. The sets themselves are returned in topologically sorted order of their first nodes.

Author(s)

Perry de Valpine

See Also

There is a non-exported function nimble:::testConditionallyIndependentSets(model,sets,initialize = TRUE) that tests whether the conditional independence of sets is valid. It should be the case that

```
nimble:::testConditionallyIndependentSets(model,getConditionallyIndependentSets(model),initialize
= TRUE) returns TRUE.
```

---

getDefinition   *Get nimbleFunction definition*

Description

Returns a list containing the nimbleFunction definition components (setup function, run function, and other member methods) for the supplied nimbleFunction generator or specialized instance.

Usage

```
getDefinition(nf)
```

Arguments

nf   A nimbleFunction generator, or a compiled or un-compiled specialized nimble-
                 Function.

Author(s)

Daniel Turek

---

getNimbleOption   *Get NIMBLE Option*

Description

Allow the user to get the value of a global _option_ that affects the way in which NIMBLE operates

Usage

```
getNimbleOption(x)
```

Arguments

x   a character string holding an option name

Value

The value of the option.

Author(s)

Christopher Paciorek

Examples

```
getNimbleOption('verifyConjugatePosteriors')
```

---

getParam   *Get value of a parameter of a stochastic node in a model*

Description

Get the value of a parameter for any single stochastic node in a model.

Usage

```
getParam(model, node, param, nodeFunctionIndex)
```

Arguments

`model`	A NIMBLE model object
`node`	The name of a stochastic node in the model
`param`	The name of a parameter for the node
`nodeFunctionIndex`	
	For internal NIMBLE use only

Details

Standard usage is as a method of a model, in the form `model$getParam(node,param)`, but the usage as a simple function with the model as the first argument as above is also allowed.

For example, suppose node 'x[1:5]' follows a multivariate normal distribution (dmnorm) in a model declared by BUGS code. model$getParam('x[1:5]', 'mean') would return the current value of the mean parameter (which may be determined from other nodes). The parameter requested does not have to be part of the parameterization used to declare the node. Rather, it can be any parameter known to the distribution. For example, one can request the scale or rate parameter of a gamma distribution, regardless of which one was used to declare the node.

---

getSamplesDPmeasure   *Get posterior samples for a Dirichlet process measure*

Description

This function obtains posterior samples from a Dirichlet process distributed random measure of a model specified using the dCRP distribution.

Usage

```
getSamplesDPmeasure(
  MCMC,
  epsilon = 1e-04,
  setSeed = FALSE,
  progressBar = getNimbleOption("MCMCprogressBar")
)
```

Arguments

MCMC   an MCMC class object, either compiled or uncompiled.

epsilon   used for determining the truncation level of the representation of the random
                measure.

setSeed   Logical or numeric argument. If a single numeric value is provided, R's random
                number seed will be set to this value. In the case of a logical value, if TRUE, then
                R's random number seed will be set to 1. Note that specifying the argument
                setSeed = 0 does not prevent setting the RNG seed, but rather sets the random
                number generation seed to 0. Default value is FALSE.

progressBar   Logical specifying whether to display a progress bar during execution (default
                = TRUE). The progress bar can be permanently disabled by setting the system
                option nimbleOptions(MCMCprogressBar = FALSE)

Details

This function provides samples from a random measure having a Dirichlet process prior. Realizations are almost surely discrete and represented by a (finite) stick-breaking representation (Sethuraman, 1994), whose atoms (or point masses) are independent and identically distributed. This sampler can only be used with models containing a dCRP distribution.

The MCMC argument is an object of class MCMC provided by buildMCMC, or its compiled version. The MCMC should already have been run, as getSamplesDPmeasure uses the posterior samples to generate samples of the random measure. Note that the monitors associated with that MCMC must include the cluster membership variable (which has the dCRP distribution), the cluster parameter variables, all variables directly determining the dCRP concentration parameter, and any stochastic parent variables of the cluster parameter variables. See help(configureMCMC) or help(addMonitors) for information on specifying monitors for an MCMC.

The epsilon argument is optional and used to determine the truncation level of the random measure. epsilon is the tail probability of the random measure, which together with posterior samples of the concentration parameter, determines the truncation level. The default value is 1e-4.

The output is a list of matrices. Each matrix represents a sample from the random measure. In order to reduce the output's dimensionality, the weights of identical atoms are added up. The stick-breaking weights are named weights and the atoms are named based on the cluster variables in the model.

For more details about sampling the random measure and determining its truncation level, see Section 3 in Gelfand, A.E. and Kottas, A. 2002.

Author(s)

Claudia Wehrhahn and Christopher Paciorek

References

Sethuraman, J. (1994). A constructive definition of Dirichlet priors. *Statistica Sinica*, 639-650.

Gelfand, A.E. and Kottas, A. (2002). A computational approach for full nonparametric Bayesian inference under Dirichlet process mixture models. *Journal of Computational and Graphical Statistics*, 11(2), 289-305.

See Also

buildMCMC, configureMCMC,

Examples

```
## Not run:
  conf <- configureMCMC(model)
  mcmc <- buildMCMC(conf)
  cmodel <- compileNimble(model)
  cmcmc <- compileNimble(mcmc, project = model)
  runMCMC(cmcmc, niter = 1000)
  outputG <- getSamplesDPmeasure(cmcmc)

## End(Not run)
```

---

getsize   *Returns number of rows of modelValues*

Description

Returns the number of rows of NIMBLE modelValues object. Works in R and NIMBLE.

Usage

```
getsize(container)
```

Arguments

container   modelValues object

Details

See the User Manual or help(modelValuesBaseClass) for information about modelValues objects

Author(s)

Clifford Anderson-Bergman

Examples

```
mvConf <- modelValuesConf(vars = 'a', types = 'double', sizes = list(a = 1) )
mv <- modelValues(mvConf)
 resize(mv, 10)
getsize(mv)
```

---

identityMatrix   *Create an Identity matrix (Deprecated)*

Description

Returns a d-by-d identity matrix (square matrix of 0's, with 1's on the main diagnol).

Usage

```
identityMatrix(d)
```

Arguments

d   The size of the identity matrix to return, will return a d-by-d matrix

Details

This function can be used in NIMBLE run code. It is deprecated because now one can use diag(d) instead.

Value

A d-by-d identity matrix

Author(s)

Daniel Turek

Examples

```
Id <- identityMatrix(d = 3)
```

initializeModel	*Performs initialization of nimble model node values and log probabilities*

Description

Performs initialization of nimble model node values and log probabilities

Usage

```
initializeModel(model, silent = FALSE)
```

Arguments

model	A setup argument, which specializes an instance of this nimble function to a particular model.
silent	logical indicating whether to suppress logging information

Details

This nimbleFunction may be used at the beginning of nimble algorithms to perform model initialization. The intended usage is to specialize an instance of this nimbleFunction in the setup function of an algorithm, then execute that specialied function at the beginning of the algorithm run function. The specialized function takes no arguments.

Executing this function ensures that all right-hand-side only nodes have been assigned real values, that all stochastic nodes have a real value, or otherwise have their simulate() method called, that all deterministic nodes have their simulate() method called, and that all log-probabilities have been calculated with the current model values. An error results if model initialization encounters a problem, for example a missing right-hand-side only node value.

Author(s)

Daniel Turek

Examples

```
myNewAlgorithm <- nimbleFunction(
   setup = function(model, ...) {
      my_initializeModel <- initializeModel(model)
      ....
   },
   run = function(...) {
      my_initializeModel()
      ....
   }
)
```

Interval   *Interval calculations*

Description

Calculations to handle censoring

Usage

```
dinterval(x, t, c, log = FALSE)

rinterval(n = 1, t, c)
```

Arguments

x	vector of interval indices.
t	vector of values.
c	vector of one or more values delineating the intervals.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations.

Details

Used for working with censoring in BUGS code. Taking c to define the endpoints of two or more intervals (with implicit endpoints of plus/minus infinity), x (or the return value of rinterval) gives the non-negative integer valued index of the interval in which t falls. See the NIMBLE manual for additional details.

Value

dinterval gives the density and rinterval generates random deviates, but these are unusual as the density is 1 if x indicates the interval in which t falls and 0 otherwise and the deviates are simply the interval(s) in which t falls.

Author(s)

Christopher Paciorek

See Also

[Distributions](#) for other standard distributions

Examples

```
endpoints <- c(-3, 0, 3)
vals <- c(-4, -1, 1, 5)
x <- rinterval(4, vals, endpoints)
dinterval(x, vals, endpoints)
dinterval(c(1, 5, 2, 3), vals, endpoints)
```

Inverse-Gamma   *The Inverse Gamma Distribution*

Description

Density, distribution function, quantile function and random generation for the inverse gamma distribution with rate or scale (mean = scale / (shape - 1)) parameterizations.

Usage

```
dinvgamma(x, shape, scale = 1, rate = 1/scale, log = FALSE)

rinvgamma(n = 1, shape, scale = 1, rate = 1/scale)

pinvgamma(
  q,
  shape,
  scale = 1,
  rate = 1/scale,
  lower.tail = TRUE,
  log.p = FALSE
)

qinvgamma(
  p,
  shape,
  scale = 1,
  rate = 1/scale,
  lower.tail = TRUE,
  log.p = FALSE
)
```

Arguments

x	vector of values.
shape	vector of shape values, must be positive.
scale	vector of scale values, must be positive.
rate	vector of rate values, must be positive.
log	logical; if TRUE, probability density is returned on the log scale.

n	number of observations.
q	vector of quantiles.
lower.tail	logical; if TRUE (default) probabilities are $P[X \leq x]$; otherwise, $P[X > x]$.
log.p	logical; if TRUE, probabilities p are given by user as log(p).
p	vector of probabilities.

Details

The inverse gamma distribution with parameters shape $= \alpha$ and scale $= \sigma$ has density

$$f(x) = \frac{s^a}{\Gamma(\alpha)} x^{-(\alpha+1)} e^{-\sigma/x}$$

for $x \geq 0$, $\alpha > 0$ and $\sigma > 0$. (Here $\Gamma(\alpha)$ is the function implemented by R's gamma() and defined in its help.

The mean and variance are $E(X) = \frac{\sigma}{\alpha} - 1$ and $Var(X) = \frac{\sigma^2}{(\alpha-1)^2(\alpha-2)}$, with the mean defined only for $\alpha > 1$ and the variance only for $\alpha > 2$.

See Gelman et al., Appendix A or the BUGS manual for mathematical details.

Value

dinvgamma gives the density, pinvgamma gives the distribution function, qinvgamma gives the quantile function, and rinvgamma generates random deviates.

Author(s)

Christopher Paciorek

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

Distributions for other standard distributions

Examples

```
x <- rinvgamma(50, shape = 1, scale = 3)
dinvgamma(x, shape = 1, scale = 3)
```

Inverse-Wishart   *The Inverse Wishart Distribution*

Description

Density and random generation for the Inverse Wishart distribution, using the Cholesky factor of either the scale matrix or the rate matrix.

Usage

```
dinvwish_chol(x, cholesky, df, scale_param = TRUE, log = FALSE)

rinvwish_chol(n = 1, cholesky, df, scale_param = TRUE)
```

Arguments

x	vector of values.
cholesky	upper-triangular Cholesky factor of either the scale matrix (when `scale_param` is TRUE) or rate matrix (otherwise).
df	degrees of freedom.
scale_param	logical; if TRUE the Cholesky factor is that of the scale matrix; otherwise, of the rate matrix.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations (only n=1 is handled currently).

Details

See Gelman et al., Appendix A for mathematical details. The rate matrix as used here is defined as the inverse of the scale matrix, S^{-1}, given in Gelman et al.

Value

`dinvwish_chol` gives the density and `rinvwish_chol` generates random deviates.

Author(s)

Christopher Paciorek

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

Distributions for other standard distributions

Examples

```
df <- 40
ch <- chol(matrix(c(1, .7, .7, 1), 2))
x <- rwish_chol(1, ch, df = df)
dwish_chol(x, ch, df = df)
```

---

is.nf   *check if a nimbleFunction*

Description

Checks an object to determine if it is a nimbleFunction (i.e., a function created by `nimbleFunction` using setup code).

Usage

```
is.nf(f, inputIsName = FALSE, where = -1)
```

Arguments

f	object to be tested
inputIsName	logical indicating whether the function is provided as the character name of the function or the function object itself
where	Optional argument needed due to R package namespace issues but which should not need to be provided by a user.

See Also

[nimbleFunction](#) for how to create a nimbleFunction

---

is.nl   *check if a nimbleList*

Description

Checks an object to determine if it is a nimbleList (i.e., a list created by `nlDef$new()`).

Usage

```
is.nl(l)
```

Arguments

l	object to be tested

See Also

[nimbleList](#) for how to create a nimbleList

LKJ *The LKJ Distribution for the Cholesky Factor of a Correlation Matrix*

Description

Density and random generation for the LKJ distribution for the Cholesky factor of a correlation matrix.

Usage

```
dlkj_corr_cholesky(x, eta, p, log = FALSE)

rlkj_corr_cholesky(n = 1, eta, p)
```

Arguments

x	upper-triangular Cholesky factor of a correlation matrix.
eta	shape parameter.
p	size of the correlation matrix (number of rows and columns); required because random generation function has no information about dimension of matrix to generate without this argument.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations (only n=1 is handled currently).

Details

See Stan Development Team for mathematical details.

Value

dlkj_corr_cholesky gives the density and rlkj_corr_cholesky generates random deviates.

Author(s)

Christopher Paciorek

References

Stan Development Team. Stan Reference Functions, version 2.27.

See Also

[Distributions](#) for other standard distributions

Examples

```
eta <- 3
x <- rlkj_corr_cholesky(1, eta, 5)
dlkj_corr_cholesky(x, eta, 5)
```

---

makeBoundInfo   *Make an object of information about a model-bound pairing for get-Bound. Used internally*

Description

Creates a simple getBound_info object, which has a list with a boundID and a type. Unlike makeParamInfo this is more bare-bones, but keeping it for parallelism with getParam.

Usage

```
makeBoundInfo(model, nodes, bound)
```

Arguments

model	A model such as returned by [nimbleModel]().
nodes	A character string naming a stochastic nodes, such as 'mu'.
bound	A character string naming a bound of the distribution, either 'lower' or 'upper'.

Details

This is used internally by [getBound](). It is not intended for direct use by a user or even a nimble-Function programmer.

---

makeParamInfo   *Make an object of information about a model-parameter pairing for getParam. Used internally*

Description

Creates a simple getParam_info object, which has a list with a paramID and a type

Usage

```
makeParamInfo(model, nodes, param, vector = FALSE)
```

Arguments

model	A model such as returned by [nimbleModel](#).
nodes	A character string naming one one or more stochastic nodes, such as "mu", "c('mu', 'beta[2]')", or "eta[1:3, 2]". getParam only works for one node at a time, but if it is indexed (nodes[i]), then makeParamInfo sets up the information for the entire vector nodes. The processing pathway is used by the NIMBLE compiler.
param	A character string naming a parameter of the distribution followed by node, such as "mean", "rate", "lambda", or whatever parameter names are relevant for the distribution of the node.
vector	A logical indicating whether nodes should definitely be treated as a vector in compiled code, even if it has length = 1. For type consistency, the compiler needs this option. If nodes has length > 1, this argument is ignored.

Details

This is used internally by [getParam](#). It is not intended for direct use by a user or even a nimble-Function programmer.

MCMCconf-class   *Class* `MCMCconf`

Description

Objects of this class configure an MCMC algorithm, specific to a particular model. Objects are normally created by calling [configureMCMC](#). Given an MCMCconf object, the actual MCMC function can be built by calling [buildMCMC](#)(conf). See documentation below for method initialize() for details of creating an MCMCconf object.

Methods

addMonitors(..., ind = 1, print = TRUE) Adds variables to the list of monitors.

>   Arguments:

>   ...: One or more character vectors of indexed nodes, or variables, which are to be monitored. These are added onto the current monitors list.

>   print: A logical argument specifying whether to print all current monitors (default TRUE).

>   Details: See the initialize() function

addMonitors2(..., print = TRUE) Adds variables to the list of monitors2.

>   Arguments:

>   ...: One or more character vectors of indexed nodes, or variables, which are to be monitored. These are added onto the current monitors2 list.

>   print: A logical argument specifying whether to print all current monitors (default TRUE).

>   Details: See the initialize() function

`addSampler( target, type = "RW", control = list(), print = FALSE, name, scalarComponents = FALSE, silent = F`
  Adds a sampler to the list of samplers contained in the MCMCconf object.

  Arguments:

  target: The target node or nodes to be sampled. This may be specified as a character vector of model node and/or variable names. This argument is required.

  type: The type of sampler to add, specified as either a character string or a nimbleFunction object. If the character argument type='newSamplerType', then either newSamplerType or sampler_newSamplertype must correspond to a nimbleFunction (i.e. a function returned by nimbleFunction, not a specialized nimbleFunction). Alternatively, the type argument may be provided as a nimbleFunction itself rather than its name. In that case, the 'name' argument may also be supplied to provide a meaningful name for this sampler. The default value is 'RW' which specifies scalar adaptive Metropolis-Hastings sampling with a normal proposal distribution. This default will result in an error if 'target' specifies more than one target node.

  control: A list of control arguments specific to the sampler function. These will override those specified in the control list argument to configureMCMC().

  print: Logical argument, specifying whether to print the details of the newly added sampler, as well as its position in the list of MCMC samplers.

  name: Optional character string name for the sampler, which is used by the printSamplers method. If 'name' is not provided, the 'type' argument is used to generate the sampler name.

  scalarComponents: Logical argument, indicating whether the specified sampler 'type' should be assigned independently to each scalar (univariate) component of the specified 'target' node or variable. This option should only be specified as TRUE when the sampler 'type' specifies a univariate sampler.

  silent: Logical argument, specifying whether to print warning messages when assigning samplers.

  ...: Additional named arguments passed through ... will be used as additional control list elements.

  Details: A single instance of the newly configured sampler is added to the end of the list of samplers for this MCMCconf object.

  Invisibly returns a list of the current sampler configurations, which are samplerConf reference class objects.

`addSamplerOne( thisSamplerName, samplerFunction, targetOne, thisControlList, print )`
  For internal use only

`getMonitors()` Returns a character vector of the current monitors
  Details: See the initialize() function

`getMonitors2()` Returns a character vector of the current monitors2
  Details: See the initialize() function

`getSamplerDefinition(ind, print = FALSE)` Returns the nimbleFunction definition of an MCMC sampler.

  Arguments:

  ind: A numeric vector or character vector. A numeric vector may be used to specify the index of the sampler definition to return, or a character vector may be used to indicate a target node for which the sampler acting on this nodes will be printed. For example, getSamplerDefinition('x[2]') will return the definition of the sampler whose target is model node 'x[2]'. If more than one sampler function is specified, only the first is returned.

Returns a list object, containing the setup function, run function, and additional member methods for the specified nimbleFunction sampler.

getSamplerExecutionOrder() Returns a numeric vector, specifying the ordering of sampler function execution.

The indices of execution specified in this numeric vector correspond to the enumeration of samplers printed by printSamplers(), or returned by getSamplers().

getSamplers(ind) Returns a list of samplerConf objects.

Arguments:

ind: A numeric vector or character vector. A numeric vector may be used to specify the indices of the samplerConf objects to return, or a character vector may be used to indicate a set of target nodes and/or variables, for which all samplers acting on these nodes will be returned. For example, getSamplers('x') will return all samplerConf objects whose target is model node 'x', or whose targets are contained (entirely or in part) in the model variable 'x'. If omitted, then all samplerConf objects in this MCMC configuration object are returned.

initialize(model, nodes, control = list(), monitors, thin = 1, monitors2 = character(), thin2 = 1, useConj Creates a MCMC configuration for a given model. The resulting object is suitable as an argument to buildMCMC.

Arguments:

model: A NIMBLE model object, created from nimbleModel(...)

nodes: An optional character vector, specifying the nodes for which samplers should be created. Nodes may be specified in their indexed form, 'y[1, 3]', or nodes specified without indexing will be expanded fully, e.g., 'x' will be expanded to 'x[1]', 'x[2]', etc. If missing, the default value is all non-data stochastic nodes. If NULL, then no samplers are added.

control: An optional list of control arguments to sampler functions. If a control list is provided, the elements will be provided to all sampler functions which utilize the named elements given. For example, the standard Metropolis-Hastings random walk sampler (sampler_RW) utilizes control list elements 'adaptive', 'adaptInterval', 'scale'. The default values for control list arguments for samplers (if not otherwise provided as an argument to configureMCMC() or addSampler()) are contained in the setup code of each sampling algorithm.

monitors: A character vector of node names or variable names, to record during MCMC sampling. This set of monitors will be recorded with thinning interval 'thin', and the samples will be stored into the 'mvSamples' object. The default value is all top-level stochastic nodes of the model – those having no stochastic parent nodes.

monitors2: A character vector of node names or variable names, to record during MCMC sampling. This set of monitors will be recorded with thinning interval 'thin2', and the samples will be stored into the 'mvSamples2' object. The default value is an empty character vector, i.e. no values will be recorded.

thin: The thinning interval for 'monitors'. Default value is one.

thin2: The thinning interval for 'monitors2'. Default value is one.

useConjugacy: A logical argument, with default value TRUE. If specified as FALSE, then no conjugate samplers will be used, even when a node is determined to be in a conjugate relationship.

onlyRW: A logical argument, with default value FALSE. If specified as TRUE, then Metropolis-Hastings random walk samplers will be assigned for all non-terminal continuous-valued nodes nodes. Discrete-valued nodes are assigned a slice sampler, and terminal nodes are assigned a posterior_predictive sampler.

onlySlice: A logical argument, with default value FALSE. If specified as TRUE, then a slice sampler is assigned for all non-terminal nodes. Terminal nodes are still assigned a posterior_predictive sampler.

multivariateNodesAsScalars: A logical argument, with default value FALSE. If specified as TRUE, then non-terminal multivariate stochastic nodes will have scalar samplers assigned to each of the scalar components of the multivariate node. The default value of FALSE results in a single block sampler assigned to the entire multivariate node. Note, multivariate nodes appearing in conjugate relationships will be assigned the corresponding conjugate sampler (provided useConjugacy == TRUE), regardless of the value of this argument.

enableWAIC: A logical argument, specifying whether to enable WAIC calculations for the resulting MCMC algorithm. Defaults to the value of nimbleOptions('MCMCenableWAIC'), which in turn defaults to FALSE. Setting nimbleOptions('MCMCenableWAIC' = TRUE) will ensure that WAIC is enabled for all calls to 'configureMCMC' and 'buildMCMC'.

controlWAIC A named list of inputs that control the behavior of the WAIC calculation, passed as the 'control' input to 'buildWAIC'. See 'help(waic)'.

warnNoSamplerAssigned: A logical argument specifying whether to issue a warning when no sampler is assigned to a node, meaning there is no matching sampler assignment rule. Default is TRUE.

print: A logical argument specifying whether to print the montiors and samplers. Default is TRUE.

...: Additional named control list elements for default samplers, or additional arguments to be passed to the autoBlock function when autoBlock = TRUE.

printMonitors() Prints all current monitors and monitors2

   Details: See the initialize() function

printSamplers(..., ind, type, displayControlDefaults = FALSE, displayNonScalars = FALSE, displayConjuga

   Prints details of the MCMC samplers.

   Arguments:

   ...: Character node or variable names, or numeric indices. Numeric indices may be used to specify the indices of the samplers to print, or character strings may be used to indicate a set of target nodes and/or variables, for which all samplers acting on these nodes will be printed. For example, printSamplers('x') will print all samplers whose target is model node 'x', or whose targets are contained (entirely or in part) in the model variable 'x'. If omitted, then all samplers are printed.

   ind: A numeric vector or character vector. A numeric vector may be used to specify the indices of the samplers to print, or a character vector may be used to indicate a set of target nodes and/or variables, for which all samplers acting on these nodes will be printed. For example, printSamplers('x') will print all samplers whose target is model node 'x', or whose targets are contained (entirely or in part) in the model variable 'x'. If omitted, then all samplers are printed.

   type: a character vector containing sampler type names. Only samplers with one of these specified types, as printed by this printSamplers method, will be displayed. Standard regular expression mathing using is also applied.

   displayConjugateDependencies: A logical argument, specifying whether to display the dependency lists of conjugate samplers (default FALSE).

   displayNonScalars: A logical argument, specifying whether to display the values of non-scalar control list elements (default FALSE).

executionOrder: A logical argument, specifying whether to print the sampler functions in the (possibly modified) order of execution (default FALSE).

byType: A logical argument, specifying whether the nodes being sampled should be printed, sorted and organized according to the type of sampler (the sampling algorithm) which is acting on the nodes (default FALSE).

removeSampler(...) Alias for removeSamplers method

removeSamplers(..., ind, print = FALSE) Removes one or more samplers from an MCMC-conf object.

Arguments:

This function also has the side effect of resetting the sampler execution ordering so as to iterate over the remaining set of samplers, sequentially, executing each sampler once.

...: Character node names or numeric indices. Character node names specify the node names for samplers to remove, or numeric indices can provide the indices of samplers to remove.

ind: A numeric vector or character vector specifying the samplers to remove. A numeric vector may specify the indices of the samplers to be removed. Alternatively, a character vector may be used to specify a set of model nodes and/or variables, and all samplers whose 'target' is among these nodes will be removed. If omitted, then all samplers are removed.

print: A logical argument specifying whether to print the current list of samplers once the removal has been done (default FALSE).

resetMonitors() Resets the current monitors and monitors2 lists to nothing.

Details: See the initialize() function

setMonitors(..., ind = 1, print = TRUE) Sets new variables to the list of monitors.

Arguments:

...: One or more character vectors of indexed nodes, or variables, which are to be monitored. These replace the current monitors list.

print: A logical argument specifying whether to print all current monitors (default TRUE).

Details: See the initialize() function

setMonitors2(..., print = TRUE) Sets new variables to the list of monitors2.

Arguments:

...: One or more character vectors of indexed nodes, or variables, which are to be monitored. These replace the current monitors2 list.

print: A logical argument specifying whether to print all current monitors (default TRUE).

Details: See the initialize() function

setSampler(...) Alias for setSamplers method

setSamplerExecutionOrder(order, print = FALSE) Sets the ordering in which sampler functions will execute.

This allows some samplers to be "turned off", or others to execute multiple times in a single MCMC iteration. The ordering in which samplers execute can also be interleaved.

Arguments:

order: A numeric vector, specifying the ordering in which the sampler functions will execute. The indices of execution specified in this numeric vector correspond to the enumeration of samplers printed by printSamplers(), or returned by getSamplers(). If this argument is omitted, the sampler execution ordering is reset so as to sequentially execute each sampler once.

print: A logical argument specifying whether to print the current list of samplers in the modified order of execution (default FALSE).

setSamplers(..., ind, print = FALSE) Sets the ordering of the list of MCMC samplers.

> This function also has the side effect of resetting the sampler execution ordering so as to iterate over the specified set of samplers, sequentially, executing each sampler once.
>
> Arguments:
>
> ...: Character strings or numeric indices. Character names may be used to specify the node names for samplers to retain. A numeric indices may be used to specify the indicies for the new list of MCMC samplers, in terms of the current ordered list of samplers.
>
> ind: A numeric vector or character vector. A numeric vector may be used to specify the indicies for the new list of MCMC samplers, in terms of the current ordered list of samplers. For example, if the MCMCconf object currently has 3 samplers, then the ordering may be reversed by calling MCMCconf$setSamplers(3:1), or all samplers may be removed by calling MCMCconf$setSamplers(numeric(0)).
>
> Alternatively, a character vector may be used to specify a set of model nodes and/or variables, and the sampler list will modified to only those samplers acting on these target nodes.
>
> As another alternative, a list of samplerConf objects may be used as the argument, in which case this ordered list of samplerConf objects will define the samplers in this MCMC configuration object, completely over-writing the current list of samplers. No checking is done to ensure the validity of the contents of these samplerConf objects; only that all elements of the list argument are, in fact, samplerConf objects.
>
> print: A logical argument specifying whether to print the new list of samplers (default FALSE).

setThin(thin, print = TRUE) Sets the value of thin.

> Arguments:
>
> thin: The new value for the thinning interval 'thin'.
>
> print: A logical argument specifying whether to print all current monitors (default TRUE).
>
> Details: See the initialize() function

setThin2(thin2, print = TRUE) Sets the value of thin2.

> Arguments:
>
> thin2: The new value for the thinning interval 'thin2'.
>
> print: A logical argument specifying whether to print all current monitors (default TRUE).
>
> Details: See the initialize() function

Author(s)

Daniel Turek

See Also

[configureMCMC](#)

Examples

```
code <- nimbleCode({
 mu ~ dnorm(0, 1)
 x ~ dnorm(mu, 1)
})
Rmodel <- nimbleModel(code)
conf <- configureMCMC(Rmodel)
```

```
conf$setSamplers(1)
conf$addSampler(target = 'x', type = 'slice', control = list(adaptInterval = 100))
conf$addMonitors('mu')
conf$addMonitors2('x')
conf$setThin(5)
conf$setThin2(10)
conf$printMonitors()
conf$printSamplers()
```

MCMCsuite   *Placeholder for MCMCsuite*

Description

This function has been moved to a separate package

Usage

```
MCMCsuite(...)
```

Arguments

...	arguments

---

modelBaseClass-class   *Class* modelBaseClass

Description

This class underlies all NIMBLE model objects: both R model objects created from the return value of nimbleModel(), and compiled model objects. The model object contains a variety of member functions, for providing information about the model structure, setting or querying properties of the model, or accessing various internal components of the model. These member functions of the modelBaseClass are commonly used in the body of the setup function argument to nimbleFunction(), to aid in preparation of node vectors, nimbleFunctionLists, and other runtime inputs. See documentation for nimbleModel for details of creating an R model object.

Methods

calculate(nodes) See 'help(calculate)'

calculateDiff(nodes) See 'help(calculateDiff)'

check() Checks for errors in model specification and for missing values that prevent use of calculate/simulate on any nodes

checkBasics() Checks for size/dimension mismatches and for presence of NAs in model variables (the latter is not an error but a note of this is given to the user)

checkConjugacy(nodeVector, restrictLink = NULL) Determines whether or not the input nodes
 appear in conjugate relationships

 Arguments:

 nodeVector: A character vector specifying one or more node or variable names. If omitted, all
 stochastic non-data nodes are checked for conjugacy.

 Details: The return value is a named list, with an element corresponding to each conjugate
 node. The list names are the conjugate node names, and list elements are the control list
 arguments required by the corresponding MCMC conjugate sampler functions. If no model
 nodes are conjugate, an empty list is returned.

expandNodeNames(nodes, env = parent.frame(), returnScalarComponents = FALSE, returnType = "names", sort
 Takes a vector of names of nodes or variables and returns the unique and expanded names in
 the model, i.e. 'x' expands to 'x[1]', 'x[2]', ...

 Arguments:

 nodes: a vector of names of nodes (or variables) to be expanded. Alternatively, can be a vector
 of integer graph IDs, but this use is intended only for advanced users

 returnScalarComponents: should multivariate nodes (i.e. dmnorm or dmulti) be broken up
 into scalar components?

 returnType: return type. Options are 'names' (character vector) or 'ids' (graph IDs)

 sort: should names be topologically sorted before being returned?

 unique: should names be the unique names or should original ordering of nodes (after expan-
 sion of any variable names into node names) be preserved

getBound(node, bound) See 'help(getBound)'

getCode() Return the code for a model after processing if-then-else statements, expanding macros,
 and replacing some keywords (e.g. nimStep for step) to avoid R ambiguity.

getConditionallyIndependentSets(nodes, givenNodes, omit = integer(), inputType = c("latent", "param", "
 Get a list of conditionally independent sets of nodes in a nimble model.

 Conditionally independent sets of nodes are typically groups of latent states whose joint con-
 ditional probability (density) will not change even if any other non-fixed node is changed.
 Default fixed nodes are data nodes and parameter nodes (with no parent nodes), but this can
 be controlled.

 model: A nimble model object (uncompiled or compiled).

 nodes: A vector of node names or their graph IDs that are the starting nodes from which
 conditionally independent sets of nodes should be found. If omitted, the default will be all
 latent nodes, defined as stochastic nodes that are not data and have at least one stochastic
 parent node (possible with deterministic nodes in between). Note that this will omit latent
 states that have no hyperparameters. An example is the first latent state in some state-space
 (time-series) models, which is sometimes declared with known prior. See type because it
 relates to the interpretation of nodes.

 givenNodes: A vector of node names or their graph IDs that should be considered as fixed and
 hence can be conditioned on. If omitted, the default will be all data nodes and all parameter
 nodes, the latter defined as nodes with no stochastic parent nodes (skipping over deterministic
 parent nodes).

 omit: A vector of node names or their graph IDs that should be omitted and should block
 further graph exploration.

 intputType: Type of input nodes provided in nodes argument. For 'latent', the input nodes are
 interpreted as latent states, from which both downstream and upstream exploration should be

done to find nodes in the same set (nodes that are not conditionally independent from each other). For 'param', the input nodes are interpreted as parameters, so graph exploration begins from the top (input) and proceeds downstream. For 'data', the input nodes are interpreted and data nodes, so graph exploration begins from the bottom (input) and proceeds upstream.

stochOnly: Logical for whether only stochastic nodes should be returned (default = TRUE). If FALSE, both deterministic and stochastic nodes are returned.

returnType: Either 'names' for returned nodes to be node names or 'ids' for returned nodes to be graph IDs.

returnScalarComponents: If FALSE (default), multivariate nodes are returned as full names (e.g. 'x[1:3]'). If TRUE, they are returned as scalar elements (e.g. 'x[1]', 'x[2]', 'x[3]').

Details: This function returns sets of conditionally independent nodes. Multiple input nodes might be in the same set or different sets, and other nodes (not in codes) will be included.

By default, deterministic dependencies of givenNodes are also counted as given nodes. This is relevant only for parent nodes. This allows the givenNodes to include only stochastic nodes. Say we have A -> B -> C -> D. A and D are givenNodes. C is a latent node. B is a deterministic node. By default, B is considered given. Otherwise, other branches that depend on B would be grouped in the same output set as C, but this is usually not what is wanted. Any use of the resulting output must ensure that B is calculated when necessary, as usual with nimble's model-generic programming. To turn off this feature, set nimbleOptions(groupDetermWithGivenInCondIndSets = FALSE).

There is a non-exported function 'nimble:::testConditionallyIndependentSets(model, sets, initialize = TRUE)' that tests whether the conditional independence of sets is valid. It should be the case that 'nimble:::testConditionallyIndependentSets(model, getConditionallyIndependentSets(model), initialize = TRUE)' returns 'TRUE'.

Return value: List of nodes that are in conditionally independent sets. Within each set, nodes are returned in topologically sorted order. The sets themselves are returned in topologically sorted order of their first nodes.

getDependencies(nodes, omit = character(), self = TRUE, determOnly = FALSE, stochOnly = FALSE, includeDat

Returns a character vector of the nodes dependent upon the input argument nodes, sorted topologically according to the model graph. In the genealogical metaphor for a graphical model, this function returns the "children" of the input nodes. In the river network metaphor, it returns downstream nodes. By default, the returned nodes include the input nodes, include both deterministic and stochastic nodes, and stop at stochastic nodes. Additional input arguments provide flexibility in the values returned.

Arguments:

nodes: Character vector of node names, with index blocks allowed, and/or variable names, the dependents of which will be returned.

omit: Character vector of node names, which will be omitted from the nodes returned. In addition, dependent nodes subsequent to these omitted nodes will not be returned. The omitted nodes argument serves to stop the downward search within the hierarchical model structure, and excludes the specified node.

self: Logical argument specifying whether to include the input argument nodes in the return vector of dependent nodes. Default is TRUE.

determOnly: Logical argument specifying whether to return only deterministic nodes. Default is FALSE.

stochOnly: Logical argument specifying whether to return only stochastic nodes. Default is FALSE. If both determOnly and stochOnly are TRUE, no nodes will be returned.

includeData: Logical argument specifying whether to include 'data' nodes (set via nimble-Model or the setData method). Default is TRUE.

dataOnly: Logical argument specifying whether to return only 'data' nodes. Default is FALSE.

includeRHSonly: Logical argument specifying whether to include right-hand-side-only nodes (model nodes which never appear on the left-hand-side of ~ or <- in the model code). These nodes are neither stochastic nor deterministic, but instead function as variable inputs to the model. Default is FALSE.

downstream: Logical argument specifying whether the downward search through the hierarchical model structure should continue beyond the first and subsequent stochastic nodes encountered, hence returning all nodes downstream of the input nodes. Default is FALSE.

returnType: Character argument specifying type of object returned. Options are 'names' (returns character vector) and 'ids' (returns numeric graph IDs for model).

returnScalarComponenets: Logical argument specifying whether multivariate nodes should be returned as full node names (i.e. 'x[1:2]') or as scalar componenets (i.e. 'x[1]' and 'x[2]').

Details: The downward search for dependent nodes propagates through deterministic nodes, but by default will halt at the first level of stochastic nodes encountered. Use getDependenciesList for a list of one-step dependent nodes of each node in the model.

getDependenciesList(returnNames = TRUE, sort = TRUE) Returns a list of all dependent neighbor relationships. Each list element gives the one-step dependencies of one vertex, and the element name is the vertex label (integer ID or character node name)

Arguments:

returnNames: If TRUE (default), list names and element contents are returns as character node names, e.g. 'x[1]'. If FALSE, everything is returned using graph IDs, which are unique integer labels for each node.

sort: If TRUE (default), each list element is returned in topologically sorted order. If FALSE, they are returned in arbitrary order.

Details: This provides a fairly raw representation of the graph (model) structure that may be useful for inspecting what NIMBLE has created from model code.

getDimension(node, params = NULL, valueOnly = is.null(params) && !includeParams, includeParams = !is.nul

Determines the dimension of the value and/or parameters of the node

Arguments:

node: A character vector specifying a single node

params: an optional character vector of names of parameters for which dimensions are desired (possibly including 'value' and alternate parameters)

valueOnly: a logical indicating whether to only return the dimension of the value of the node

includeParams: a logical indicating whether to return dimensions of parameters. If TRUE and 'params' is NULL then dimensions of all parameters, including the dimension of the value of the node, are returned

Details: The return value is a numeric vector with an element for each parameter/value requested.

getDistribution(nodes) Returns the names of the distributions for the requested node or nodes

Arguments:

nodes: A character vector specifying one or more node or variable names.

Details: The return value is a character vector with an element for each node indicated in the input. Note that variable names are expanded to their constituent node names, so the length of the output may be longer than that of the input.

getDownstream(...) Identical to getDependencies(..., downstream = TRUE)

> Details: See documentation for member method getDependencies.

getLogProb(nodes) See 'help(getLogProb)'

getNodeNames(determOnly = FALSE, stochOnly = FALSE, includeData = TRUE, dataOnly = FALSE, includeRHSonly

> Returns a character vector of all node names in the model, in topologically sorted order. A variety of logical arguments allow for flexible subsetting of all model nodes.
>
> Arguments:
>
> determOnly: Logical argument specifying whether to return only deterministic nodes. Default is FALSE.
>
> stochOnly: Logical argument specifying whether to return only stochastic nodes. Default is FALSE.
>
> includeData: Logical argument specifying whether to include 'data' nodes (set via the member method setData). Default is TRUE.
>
> dataOnly: Logical argument specifying whether to return only 'data' nodes. Default is FALSE.
>
> includeRHSonly: Logical argument specifying whether to include right-hand-side-only nodes (model nodes which never appear on the left-hand-side of ~ or <- in the model code). Default is FALSE.
>
> topOnly: Logical argument specifying whether to return only top-level nodes from the hierarchical model structure.
>
> latentOnly: Logical argument specifying whether to return only latent (mid-level) nodes from the hierarchical model structure.
>
> endOnly: Logical argument specifying whether to return only end nodes from the hierarchical model structure.
>
> returnType: Character argument specific type object returned. Options are 'names' (returns character vector) and 'ids' (returns numeric graph IDs for model)
>
> returnScalar Componenets: Logical argument specifying whether multivariate nodes should return full node name (i.e. 'x[1:2]') or should break down into scalar componenets (i.e. 'x[1]' and 'x[2]')
>
> Details: Multiple logical input arguments may be used simultaneously. For example, 'model$getNodeNames(endOnly = TRUE, dataOnly = TRUE)' will return all end-level nodes from the model which are designated as 'data'.

getParam(node, param) See 'help(getParam)'

getParents(nodes, omit = character(), self = FALSE, determOnly = FALSE, stochOnly = FALSE, includeData = T

> Returns a character vector of the nodes on which the input nodes depend, sorted topologically according to the model graph, by default recursing and stopping at stochastic parent nodes. In the genealogical metaphor for a graphical model, this function returns the "parents" of the input nodes. In the river network metaphor, it returns upstream nodes. By default, the returned nodes omit the input nodes. Additional input arguments provide flexibility in the values returned.
>
> Arguments:
>
> nodes: Character vector of node names, with index blocks allowed, and/or variable names, the parents of which will be returned.
>
> omit: Character vector of node names, which will be omitted from the nodes returned. In addition, parent nodes beyond these omitted nodes will not be returned. The omitted nodes

argument serves to stop the upward search through the hierarchical model structure, and excludes the specified node.

self: Logical argument specifying whether to include the input argument nodes in the return vector of dependent nodes. Default is FALSE.

determOnly: Logical argument specifying whether to return only deterministic nodes. Default is FALSE.

stochOnly: Logical argument specifying whether to return only stochastic nodes. Default is FALSE. If both determOnly and stochOnly are TRUE, no nodes will be returned.

includeData: Logical argument specifying whether to include 'data' nodes (set via nimbleModel or the setData method). Default is TRUE.

dataOnly: Logical argument specifying whether to return only 'data' nodes. Default is FALSE.

includeRHSonly: Logical argument specifying whether to include right-hand-side-only nodes (model nodes which never appear on the left-hand-side of ~ or <- in the model code). These nodes are neither stochastic nor deterministic, but instead function as variable inputs to the model. Default is FALSE.

upstream: Logical argument specifying whether the upward search through the hierarchical model structure should continue beyond the first and subsequent stochastic nodes encountered, hence returning all nodes upstream of the input nodes. Default is FALSE.

immediateOnly: Logical argument specifying whether only the immediate parent nodes should be returned, even if they are deterministic. If FALSE, getParents recurses and stops at stochastic nodes. Default is FALSE.

returnType: Character argument specifying type of object returned. Options are 'names' (returns character vector) and 'ids' (returns numeric graph IDs for model).

returnScalarComponenets: Logical argument specifying whether multivariate nodes should be returned as full node names (i.e. 'x[1:2]') or as scalar componenets (i.e. 'x[1]' and 'x[2]').

Details: The upward search for dependent nodes propagates through deterministic nodes, but by default will halt at the first level of stochastic nodes encountered. Use getParentsList for a list of one-step parent nodes of each node in the model.

getParentsList(returnNames = TRUE, sort = TRUE) Returns a list of all parent neighbor relationships. Each list element gives the one-step parents of one vertex, and the element name is the vertex label (integer ID or character node name)

Arguments:

returnNames: If TRUE (default), list names and element contents are returns as character node names, e.g. 'x[1]'. If FALSE, everything is returned using graph IDs, which are unique integer labels for each node.

sort: If TRUE (default), each list element is returned in topologically sorted order. If FALSE, they are returned in arbitrary order.

Details: This provides a fairly raw representation of the graph (model) structure that may be useful for inspecting what NIMBLE has created from model code.

getVarNames(includeLogProb = FALSE, nodes) Returns the names of all variables in a model, optionally including the logProb variables

Arguments:

logProb: Logical argument specifying whether or not to include the logProb variables. Default is FALSE.

nodes: An optional character vector supplying a subset of nodes for which to extract the variable names and return the unique set of variable names

`initializeInfo()` Provides more detailed information on which model nodes are not initialized.

`isBinary(nodes)` Determines whether one or more nodes represent binary random variables

> Arguments:
>
> nodes: A character vector specifying one or more node or variable names.
>
> Details: The return value is a character vector with an element for each node indicated in the input. Note that variable names are expanded to their constituent node names, so the length of the output may be longer than that of the input.

`isData(nodes)` Returns a vector of logical TRUE / FALSE values, corresponding to the 'data' flags of the input node names.

> Arguments:
>
> nodes: A character vector of node or variable names.
>
> Details: The variable or node names specified is expanded into a vector of model node names. A logical vector is returned, indicating whether each model node has been flagged as containing 'data'. Multivariate nodes for which any elements are flagged as containing 'data' will be assigned a value of TRUE.

`isDeterm(nodes)` Determines whether one or more nodes are deterministic

> Arguments:
>
> nodes: A character vector specifying one or more node or variable names.
>
> Details: The return value is a character vector with an element for each node indicated in the input. Note that variable names are expanded to their constituent node names, so the length of the output may be longer than that of the input.

`isDiscrete(nodes)` Determines whether one or more nodes represent discrete random variables

> Arguments:
>
> nodes: A character vector specifying one or more node or variable names.
>
> Details: The return value is a character vector with an element for each node indicated in the input. Note that variable names are expanded to their constituent node names, so the length of the output may be longer than that of the input.

`isEndNode(nodes)` Determines whether one or more nodes are end nodes (nodes with no stochastic dependences)

> Arguments:
>
> nodes: A character vector specifying one or more node or variable names.
>
> Details: The return value is logical vector with an element for each node indicated in the input. Note that variable names are expanded to their constituent node names, so the length of the output may be longer than that of the input.

`isMultivariate(nodes)` Determines whether one or more nodes represent multivariate nodes

> Arguments:
>
> nodes: A character vector specifying one or more node or variable names.
>
> Details: The return value is a logical vector with an element for each node indicated in the input. Note that variable names are expanded to their constituent node names, so the length of the output may be longer than that of the input.

`isStoch(nodes)` Determines whether one or more nodes are stochastic

> Arguments:
>
> nodes: A character vector specifying one or more node or variable names.

Details: The return value is a character vector with an element for each node indicated in the input. Note that variable names are expanded to their constituent node names, so the length of the output may be longer than that of the input.

isTruncated(nodes) Determines whether one or more nodes are truncated

   Arguments:

   nodes: A character vector specifying one or more node or variable names.

   Details: The return value is a character vector with an element for each node indicated in the input. Note that variable names are expanded to their constituent nodes names, so the length of the output may be longer than that of the input

isUnivariate(nodes) Determines whether one or more nodes represent univariate random variables

   Arguments:

   nodes: A character vector specifying one or more node or variable names.

   Details: The return value is a character vector with an element for each node indicated in the input. Note that variable names are expanded to their constituent nodes names, so the length of the output may be longer than that of the input

newModel(data = NULL, inits = NULL, modelName = character(), replicate = FALSE, check = getNimbleOption("c
   Returns a new R model object, with the same model definiton (as defined from the original model code) as the existing model object.

   Arguments:

   data: A named list specifying data nodes and values, for use in the newly returned model. If not provided, the data argument from the creation of the original R model object will be used.

   inits: A named list specifying initial valuse, for use in the newly returned model. If not provided, the inits argument from the creation of the original R model object will be used.

   modelName: An optional character string, used to set the internal name of the model object. If provided, this name will propagate throughout the generated C++ code, serving to improve readability.

   replicate: Logical specifying whether to replicate all current values and data flags from the current model in the new model. If TRUE, then the data and inits arguments are not used. Default value is FALSE.

   check: A logical indicating whether to check the model object for missing or invalid values. Default is given by the NIMBLE option 'checkModel', see help on 'nimbleOptions' for details.

   calculate: A logical indicating whether to run 'calculate' on the model; this will calculate all deterministic nodes and logProbability values given the current state of all nodes. Default is TRUE. For large models, one might want to disable this, but note that deterministic nodes, including nodes introduced into the model by NIMBLE, may be NA.

   Details: The newly created model object will be identical to the original model in terms of structure and functionality, but entirely distinct in terms of the internal values.

resetData() Resets the 'data' property of ALL model nodes to FALSE. Subsequent to this call, the model will have no nodes flagged as 'data'.

setData(..., warnAboutMissingNames = TRUE) Sets the 'data' flag for specified nodes to TRUE, and also sets the value of these nodes to the value provided. This is the exclusive method for specifying 'data' nodes in a model object. When a 'data' argument is provided to 'nimbleModel()', it uses this method to set the data nodes.

Arguments:

...: Arguments may be provided as named elements with numeric values or as character names of model variables. These may be provided in a single list, a single character vector, or as multiple arguments. When a named element with a numeric value is provided, the size and dimension must match the corresponding model variable. This value will be copied to the model variable and any non-NA elements will be marked as data. When a character name is provided, the value of that variable in the model is not changed but any currently non-NA values are marked as data. Examples: setData('x', y = 1:10) will mark both x and y as data and will set the value of y to 1:10. setData(list('x', y = 1:10)) is equivalent. setData(c('x','y')) or setData('x','y') will mark both x and y as data.

Details: If a provided value (or the current value in the model when only a name is specified) contains some NA values, then the model nodes corresponding to these NAs will not have their value set, and will not be designated as 'data'. Only model nodes corresponding to numeric values in the argument list elements will be designated as data. Designating a deterministic model node as 'data' will result in an error. Designating part of a multivariate node as 'data' and part as non-data (NA) is allowed, but 'isData()' will report such a node as being 'data', calculations with the node will generally return NA, and MCMC samplers will not be assigned to such nodes.

setInits(inits) Sets initial values (or more generally, any named list of value elements) into the model

Arguments:

inits: A named list. The names of list elements must correspond to model variable names. The elements of the list must be of class numeric, with size and dimension each matching the corresponding model variable.

simulate(nodes, includeData = FALSE) See 'help(simulate)'

topologicallySortNodes(nodes, returnType = "names") Sorts the input list of node names according to the topological dependence ordering of the model structure.

Arguments:

nodes: A character vector of node or variable names, which is to be topologically sorted. Alternatively can be a numeric vector of graphIDs

returnType: character vector indicating return type. Choices are "names" or "ids"

Details: This function merely reorders its input argument. This may be important prior to calls such as model$simulate(nodes) or model$calculate(nodes), to enforce that the operation is performed in topological order.

Author(s)

Daniel Turek

See Also

[initializeModel](#)

Examples

```
code <- nimbleCode({
    mu ~ dnorm(0, 1)
```

```
    x[1] ~ dnorm(mu, 1)
    x[2] ~ dnorm(mu, 1)
})
Rmodel <- nimbleModel(code)
modelVars <- Rmodel$getVarNames()   ## returns 'mu' and 'x'
modelNodes <- Rmodel$getNodeNames()   ## returns 'mu', 'x[1]' and 'x[2]'
Rmodel$resetData()
Rmodel$setData(list(x = c(1.2, NA)))   ## flags only 'x[1]' node as data
Rmodel$isData(c('mu', 'x[1]', 'x[2]'))   ## returns c(FALSE, TRUE, FALSE)
```

---

modelDefClass-class   *Class for NIMBLE model definition*

Description

Class for NIMBLE model definition that is not usually needed directly by a user.

Details

See modelBaseClass for information about creating NIMBLE BUGS models.

---

modelInitialization   *Information on initial values in a NIMBLE model*

Description

Having uninitialized nodes in a NIMBLE model can potentially cause some algorithms to fail and can lead to poor performance in others. Here are some general guidelines on how non-intitialized variables can affect performance:

- MCMC will auto-initialize but will do so from the prior distribution. This can cause slow convergence, especially in the case of diffuse priors.

- Likewise, particle filtering methods will initialize top-level parameters from their prior distributions, which can lead to errors or poor performance in these methods.

---

modelValues   *Create a NIMBLE modelValues Object*

Description

Builds modelValues object from a model values configuration object, which can include a NIMBLE model

Usage

```
modelValues(conf, m = 1)
```

Arguments

conf   An object which includes information for building modelValues. Can either be
                a NIMBLE model (see `help(modelBaseClass)`) or the object returned from
                `modelValuesConf`

m   The number of rows to create in the modelValues object. Can later be changed
                with `resize`

Details

See the User Manual or `help(modelValuesBaseClass)` for information about manipulating NIM-
BLE modelValues object returned by this function

Author(s)

NIMBLE development team

Examples

```
#From model object:
code <- nimbleCode({
 a ~ dnorm(0,1)
 for(i in 1:3){
for(j in 1:3)
b[i,j] ~ dnorm(0,1)
}
})
Rmodel <- nimbleModel(code)
Rmodel_mv <- modelValues(Rmodel, m = 2)
#Custom modelValues object:
mvConf <- modelValuesConf(vars = c('x', 'y'),
            types = c('double', 'int'),
            sizes = list(x = 3, y = c(2,2)))
custom_mv <- modelValues(mvConf, m = 2)
custom_mv['y',]
```

---

modelValuesBaseClass-class

                    Class `modelValuesBaseClass`

Description

modelValues are NIMBLE containers built to store values from models. They can either be built directly from a model or be custom built via the `modelValuesConf` function. They consist of rows, where each row can be thought of as a set of values from a model. Like most nimble objects, and unlike most R objects, they are passed by reference instead of by value.

See the User Manual for more details.

Examples

```
mvConf <- modelValuesConf(vars = c('a', 'b'),
types = c('double', 'double'),
sizes = list(a = 1, b = c(2,2) ) )
mv <- modelValues(mvConf)
as.matrix(mv)
resize(mv, 2)
as.matrix(mv)
mv['a',1] <- 1
mv['a',2] <- 2
mv['b',1] <- matrix(0, nrow = 2, ncol = 2)
mv['b',2] <- matrix(1, nrow = 2, ncol = 2)
mv['a',]
as.matrix(mv)
basicModelCode <- nimbleCode({
a ~ dnorm(0,1)
for(i in 1:4)
b[i] ~ dnorm(0,1)
})
basicModel <- nimbleModel(basicModelCode)
basicMV <- modelValues(basicModel, m = 2) # m sets the number of rows
basicMV['b',]
```

---

modelValuesConf   *Create the confs for a custom NIMBLE modelValues object*

Description

Builds an R-based modelValues conf object

Usage

```
modelValuesConf(
  symTab,
  className,
  vars,
  types,
  sizes,
  modelDef = NA,
  where = globalenv()
)
```

Arguments

symTab	For internal use only
className	For internal use only
vars	A vector of character strings naming each variable in the modelValues object
types	A vector of character strings describing the type of data for the modelValues object. Options include 'double' (for real-valued variables) and 'int'.
sizes	A list in which the named items of the list match the var arguments and each item is a numeric vector of the dimensions
modelDef	For internal use only
where	For internal use only

Details

See the User Manual or `help(modelValuesBaseClass)` and `help(modelValues)` for information

Author(s)

Clifford Anderson-Bergman

Examples

```
#Custom modelValues object:
mvConf <- modelValuesConf(vars = c('x', 'y'),
types = c('double', 'int'),
sizes = list(x = 3, y = c(2,2)))
custom_mv <- modelValues(mvConf, m = 2)
custom_mv['y',]
```

model_macro_builder   *EXPERIMENTAL: Turn a function into a model macro builder A*
                             model macro expands one line of code in a nimbleModel into one
                             or more new lines. This supports compact programming by defin-
                             ing re-usable modules. model_macro_builder *takes as input a*
                             function that constructs new lines of model code from the origi-
                             nal line of code. It returns a function suitable for internal use by
                             nimbleModel *that arranges arguments for input function. Macros*
                             are an experimental feature and are available only after setting
                             nimbleOptions(enableModelMacros = TRUE).

Description

EXPERIMENTAL: Turn a function into a model macro builder A model macro expands one line
of code in a nimbleModel into one or more new lines. This supports compact programming by
defining re-usable modules. model_macro_builder takes as input a function that constructs new
lines of model code from the original line of code. It returns a function suitable for internal use by
nimbleModel that arranges arguments for input function. Macros are an experimental feature and
are available only after setting nimbleOptions(enableModelMacros = TRUE).

Usage

```
model_macro_builder(fun, use3pieces = TRUE, unpackArgs = TRUE)
```

Arguments

fun	A function written to construct new lines of model code.
use3pieces	(TRUE or FALSE) Should the arguments from the input line be split into pieces for the LHS (left-hand side), RHS (right-hand side, possibly further split depending on unpackArgs), and stoch (TRUE if the line uses a ~, FALSE otherwise)? See details and examples.
unpackArgs	(TRUE or FALSE) Should arguments be passed as a list (FALSE) or as separate arguments (TRUE)? See details and examples.

Details

The arguments use3pieces and unpackArgs indicate how fun expects to have arguments arranged
from an input line of code (processed by nimbleModel).

Consider the defaults use3pieces = TRUE and unpackArgs = TRUE, for a macro called macro1. In
this case, the line of model code x ~ macro1(arg1 = z[1:10],arg2 = "hello") will be passed to
fun as fun(stoch = TRUE,LHS = x,arg1 = z[1:10],arg2 = "hello").

If use3pieces = TRUE but unpackArgs = FALSE, then the RHS will be passed as is, without unpack-
ing its arguments into separate arguments to fun. In this case, x ~ macro1(arg1 = z[1:10],arg2 =
"hello") will be passed to fun as fun(stoch = TRUE,LHS = x,RHS = macro1(arg1 = z[1:10],arg2
= "hello")).

If use3pieces = FALSE and unpackArgs = FALSE, the entire line of code is passed as a single object. In this case, x ~ macro1(arg1 = z[1:10],arg2 = "hello") will be passed to fun as fun(x ~ macro1(arg1 = z[1:10],arg2 = "hello")). It is also possible in this case to pass a macro without using a ~ or <-. For example, the line macro1(arg1 = z[1:10],arg2 = "hello") will be passed to fun as fun(macro1(arg1 = z[1:10],arg2 = "hello")).

If use3pieces = FALSE and unpackArgs = TRUE, it won't make sense to anticipate a declaration using ~ or <-. Ins#' tead, arguments from an arbitrary call will be passed as separate arguments. #' For example, the line macro1(arg1 = z[1:10],arg2 = "hello") will be pa#' ssed to fun as fun(arg1 = z[1:10],arg2 = "hello").

It is extremely useful to be familiar with processing R code as an object to write fun correctly. Functions such as substitute and as.name (e.g. as.name('~')), quote, parse and deparse are particularly handy.

Multiple lines of new code should be contained in {} . Extra curly braces are not a problem. See example 2.

Macro expansion is done recursively: One macro can return code that invokes another macro.

Value

A list with a named element code that contains the replacement code.

Examples

```
nimbleOptions(enableModelMacros = TRUE)
nimbleOptions(verbose = FALSE)

## Example 1: Say one is tired of writing "for" loops.
## This macro will generate a "for" loop with dnorm declarations
all_dnorm <- model_macro_builder(
    function(stoch, LHS, RHSvar, start, end, sd = 1) {
        newCode <- substitute(
            for(i in START:END) {
                LHS[i] ~ dnorm(RHSvar[i], SD)
            },
            list(START = start,
                 END = end,
                 LHS = LHS,
                 RHSvar = RHSvar,
                 SD = sd))
        list(code = newCode)
    },
    use3pieces = TRUE,
    unpackArgs = TRUE
)

model1 <- nimbleModel(
    nimbleCode(
    {
        ## Create a "for" loop of dnorm declarations by invoking the macro
        x ~ all_dnorm(mu, start = 1, end = 10)
    }
```

```
    ))

## show code from expansion of macro
model1$getCode()
## The result should be:
## {
##     for (i in 1:10) {
##         x[i] ~ dnorm(mu[i], 1)
##     }
## }

## Example 2: Say one is tired of writing priors.
## This macro will generate a set of priors in one statement
flat_normal_priors <- model_macro_builder(
    function(...) {
        allVars <- list(...)
        priorDeclarations <- lapply(allVars,
                                function(x)
                                    substitute(VAR ~ dnorm(0, sd = 1000),
                                               list(VAR = x)))
        newCode <- quote({})
        newCode[2:(length(allVars)+1)] <- priorDeclarations
        list(code = newCode)
    },
    use3pieces = FALSE,
    unpackArgs = TRUE
)

model2 <- nimbleModel(
    nimbleCode(
    {
        flat_normal_priors(mu, beta, gamma)
    }
    ))

## show code from expansion of macro
model2$getCode()
## The result should be:
## {
##     {
##         mu ~ dnorm(0, sd = 1000)
##         beta ~ dnorm(0, sd = 1000)
##         gamma ~ dnorm(0, sd = 1000)
##     }
## }
## Extra curly braces do not matter.
```

Multinomial   *The Multinomial Distribution*

Description

Density and random generation for the multinomial distribution

Usage

```
dmulti(x, size = sum(x), prob, log = FALSE)

rmulti(n = 1, size, prob)
```

Arguments

x	vector of values.
size	number of trials.
prob	vector of probabilities, internally normalized to sum to one, of same length as x
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations (only n=1 is handled currently).

Details

See Gelman et al., Appendix A or the BUGS manual for mathematical details.

Value

`dmulti` gives the density and `rmulti` generates random deviates.

Author(s)

Christopher Paciorek

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

[Distributions](#) for other standard distributions

Examples

```
size <- 30
probs <- c(1/4, 1/10, 1 - 1/4 - 1/10)
x <- rmulti(1, size, probs)
dmulti(x, size, probs)
```

Multivariate-t   *The Multivariate t Distribution*

Description

Density and random generation for the multivariate t distribution, using the Cholesky factor of either the precision matrix (i.e., inverse scale matrix) or the scale matrix.

Usage

```
dmvt_chol(x, mu, cholesky, df, prec_param = TRUE, log = FALSE)

rmvt_chol(n = 1, mu, cholesky, df, prec_param = TRUE)
```

Arguments

x	vector of values.
mu	vector of values giving the location of the distribution.
cholesky	upper-triangular Cholesky factor of either the precision matrix (i.e., inverse scale matrix) (when prec_param is TRUE) or scale matrix (otherwise).
df	degrees of freedom.
prec_param	logical; if TRUE the Cholesky factor is that of the precision matrix; otherwise, of the scale matrix.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations (only n=1 is handled currently).

Details

See Gelman et al., Appendix A or the BUGS manual for mathematical details. The 'precision' matrix as used here is defined as the inverse of the scale matrix, Σ^{-1}, given in Gelman et al.

Value

dmvt_chol gives the density and rmvt_chol generates random deviates.

Author(s)

Peter Sujan

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

Distributions for other standard distributions

Examples

```
mu <- c(-10, 0, 10)
scalemat <- matrix(c(1, .9, .3, .9, 1, -0.1, .3, -0.1, 1), 3)
ch <- chol(scalemat)
x <- rmvt_chol(1, mu, ch, df = 1, prec_param = FALSE)
dmvt_chol(x, mu, ch, df = 1, prec_param = FALSE)
```

MultivariateNormal *The Multivariate Normal Distribution*

Description

Density and random generation for the multivariate normal distribution, using the Cholesky factor of either the precision matrix or the covariance matrix.

Usage

```
dmnorm_chol(x, mean, cholesky, prec_param = TRUE, log = FALSE)

rmnorm_chol(n = 1, mean, cholesky, prec_param = TRUE)
```

Arguments

x	vector of values.
mean	vector of values giving the mean of the distribution.
cholesky	upper-triangular Cholesky factor of either the precision matrix (when prec_param is TRUE) or covariance matrix (otherwise).
prec_param	logical; if TRUE the Cholesky factor is that of the precision matrix; otherwise, of the covariance matrix.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations (only n=1 is handled currently).

Details

See Gelman et al., Appendix A or the BUGS manual for mathematical details. The rate matrix as used here is defined as the inverse of the scale matrix, S^{-1}, given in Gelman et al.

Value

dmnorm_chol gives the density and rmnorm_chol generates random deviates.

Author(s)

Christopher Paciorek

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

[Distributions](#) for other standard distributions

Examples

```
mean <- c(-10, 0, 10)
covmat <- matrix(c(1, .9, .3, .9, 1, -0.1, .3, -0.1, 1), 3)
ch <- chol(covmat)
x <- rmnorm_chol(1, mean, ch, prec_param = FALSE)
dmnorm_chol(x, mean, ch, prec_param = FALSE)
```

---

nfMethod *access (call) a member function of a nimbleFunction*

Description

Internal function for accessing a member function (method) of a nimbleFunction. Normally a user will write nf$method(x) instead of nfMethod(nf,method)(x).

Usage

```
nfMethod(nf, methodName)
```

Arguments

nf	a specialized nimbleFunction, i.e. one that has already had setup parameters processed
methodName	a character string giving the name of the member function to call

Details

nimbleFunctions have a default member function called run, and may have other member functions provided via the methods argument to nimbleFunction. As an internal step, the NIMBLE compiler turns nf$method(x) into nfMethod(nf,method)(x), but a NIMBLE user or programmer would not normally need to use nfMethod directly.

Value

a function that can be called.

Author(s)

NIMBLE development team

---

nfVar   *Access or set a member variable of a nimbleFunction*

Description

Access or set a member variable of a specialized nimbleFunction, i.e. a variable passed to or created during the setup function that is used in run code or preserved by setupOutputs. Works in R for any variable and in NIMBLE for numeric variables.

Usage

```
nfVar(nf, varName)

nfVar(nf, varName) <- value
```

Arguments

nf   a specialized nimbleFunction, i.e. a function returned by executing a function
                returned from nimbleFunction with setup arguments

varName   a character string naming a variable in the setup function.

value   value to set the variable to.

Details

Internal way to access or set a member variable of a nimbleFunction created during setup. Normally in NIMBLE code you would use nf$var instead of nfVar(nf,var).

When nimbleFunction is called and a setup function is provided, then nimbleFunction returns a function. That function is a generator that should be called with arguments to the setup function and returns another function with run and possibly other member functions. The member functions can use objects created or passed to setup. During internal processing, the NIMBLE compiler turns some cases of nf$var into nfVar(nf,var). These provide direct access to setup variables (member data). nfVar is not typically called by a NIMBLE user or programmer.

For internal access to methods of nf, see [nfMethod](#).

For more information, see ?nimbleFunction and the NIMBLE [User Manual](#).

Value

whatever varName is in the nimbleFunction nf.

Author(s)

NIMBLE development team

Examples

```
nfGen1 <- nimbleFunction(
    setup = function(A) {
    B <- matrix(rnorm(4), nrow = 2)
    setupOutputs(B) ## preserves B even though it is not used in run-code
    },
   run = function() {
      print('This is A', A, '\n')
})

nfGen2 <- nimbleFunction(
   setup = function() {
     nf1 <- nfGen1(1000)
   },
   run = function() {
      print('accessing A:', nfVar(nf1, 'A'))
      nfVar(nf1, 'B')[2,2] <<- -1000
      print('accessing B:', nfVar(nf1, 'B'))
    })

nf2 <- nfGen2()
nf2$run()
```

nimble	*nimble*

Description

A system for writing hierarchical statistical models largely compatible with BUGS and JAGS, writing nimbleFunctions to operate models and do basic R-style math, and compiling both models and nimbleFunctions via custom-generated C++. NIMBLE includes default methods for MCMC, Monte Carlo Expectation Maximization, and some other tools. The nimbleFunction system makes it easy to do things like implement new MCMC samplers from R, customize the assignment of samplers to different parts of a model from R, and compile the new samplers automatically via C++ alongside the samplers NIMBLE provides. NIMBLE extends the BUGS/JAGS language by making it extensible: New distributions and functions can be added, including as calls to external compiled code. Although most people think of MCMC as the main goal of the BUGS/JAGS language for writing models, one can use NIMBLE for writing arbitrary other kinds of model-generic algorithms as well. A full User Manual is available at https://r-nimble.org/html_manual/cha-welcome-nimble.html.

nimble-internal	*Functions and Classes Internal to NIMBLE*

Description

Functions and classes used internally in NIMBLE and not expected to be called directly by users. Some functions and classes not intended for direct use are documented and/or exported because they are used within Reference Class methods for classes programmatically generated by NIMBLE.

Author(s)

NIMBLE Development Team

---

nimble-math   *Mathematical functions for BUGS and nimbleFunction programming*

Description

Mathematical functions for use in BUGS code and in nimbleFunction programming (i.e., nimble-Function run code). See Chapter 5 of the User Manual for more details.

Author(s)

NIMBLE Development Team

---

nimble-R-functions   *NIMBLE language functions for R-like vector construction*

Description

The functions `c`, `rep`, `seq`, `which`, `diag`, `length`, `seq_along`, `is.na`, `is.nan`, `any`, and `all` can be used in nimbleFunctions and compiled using `compileNimble`.

Usage

```
nimC(...)

nimRep(x, ...)

nimSeq(from, to, by, length.out)
```

Arguments

`...`	values to be concatenated.
`x`	vector of values to be replicated (`rep`), or logical array or vector (`which`), or object whose length is wanted (`length`), or input value (`diag`), or vector of values to be tested/checked (`is.na`, `is.nan`, `any`, `all`).
`from`	starting value of sequence.
`to`	end value of sequence.
`by`	increment of the sequence.
`length.out`	desired length of the sequence.

Details

For c, rep, seq, these functions are NIMBLE's version of similar R functions, e.g., nimRep for rep. In a nimbleFunction, either the R name (e.g., rep) or the NIMBLE name (e.g., nimRep) can be used. If the R name is used, it will be converted to the NIMBLE name. For which, length, diag, seq_along, is.na, is.nan, any, all simply use the standard name without "nim". These functions largely mimic (see exceptions below) the behavior of their R counterparts, but they can be compiled in a nimbleFunction using compileNimble.

nimC is NIMBLE's version of c and behaves identically.

nimRep is NIMBLE's version of rep. It should behave identically to rep. There are no NIMBLE versions of rep.int or rep_len.

nimSeq is NIMBLE's version of seq. It behaves like seq with support for from, to, by and length.out arguments. The along.with argument is not supported. There are no NIMBLE versions of seq.int, seq_along or seq_len, with the exception that seq_along can take a nimble-FunctionList as an argument to provide the index range of a for-loop (User Manual Ch. 13).

which behaves like the R version but without support for arr.ind or useNames arguments.

diag behaves like the R version but without support for the nrow and ncol arguments.

length behaves like the R version.

seq_along behaves like the R version.

is.na behaves like the R version but does not correctly handle NA values from R that are type 'logical', so convert these using as.numeric() before passing from R to NIMBLE.

is.nan behaves like the R version, but treats NA of type 'double' as being NaN and NA of type 'logical' as not being NaN.

any behaves like the R version but takes only one argument and treats NAs as FALSE.

all behaves like the R version but takes only one argument and treats NAs as FALSE.

---

nimbleCode   *Turn BUGS model code into an object for use in* nimbleModel *or*
                           readBUGSmodel

Description

Simply keeps model code as an R call object, the form needed by nimbleModel and optionally usable by readBUGSmodel.

Usage

```
nimbleCode(code)
```

Arguments

code   expression providing the code for the model

Details

It is equivalent to use the R function quote. nimbleCode is simply provided as a more readable alternative for NIMBLE users not familiar with quote.

Author(s)

Daniel Turek

Examples

```
code <- nimbleCode({
    x ~ dnorm(mu, sd = 1)
    mu ~ dnorm(0, sd = prior_sd)
})
```

---

nimbleExternalCall   *Create a nimbleFunction that wraps a call to external compiled code*

Description

Given C header information, a function that takes scalars or pointers can be called from a compiled nimbleFunction. If non-scalar return values are needed, an argument can be selected to behave as the return value in nimble.

Usage

```
nimbleExternalCall(
  prototype,
  returnType,
  Cfun,
  headerFile,
  oFile,
  where = getNimbleFunctionEnvironment()
)
```

Arguments

prototype	Argument type information. This can be provided as an R function using nimbleFunction type declarations or as a list of nimbleType objects.
returnType	Return object type information. This can be provided similarly to prototype as either a nimbleFunction type declaration or as a nimbleType object. In the latter case, the name will be ignored. If there is no return value, this should be void().
Cfun	Name of the external function (character).
headerFile	Name (possibly including file path) of the header file where Cfun is declared.

oFile	Name (possibly including path) of the .o file where Cfun has been compiled. Spaces in the path may cause problems.
where	An optional `where` argument passed to `setRefClass` for where the reference class definition generated for this nimbleFunction will be stored. This is needed due to R package namespace issues but should never need to be provided by a user.

Details

The only argument types allowed in Cfun are `double`, `int`, and `bool`, corresponding to `nimbleFunction` types double, `integer`, and `logical`, respectively.

If the dimensionality is greater than zero, the arguments in Cfun should be pointers. This means it will typically be necessary to pass additional integer arguments telling Cfun the size(s) of non-scalar arguments.

The return argument can only be a scalar or void. Since non-scalar arguments are passed by pointer, you can use an argument to return results from Cfun. If you wish to have a `nimbleFunction` that uses one argument of Cfun as a return object, you can wrap the result of `nimbleExternalCall` in another `nimbleFunction` that allocates the return object. This is useful for using Cfun in a `nimbleModel`. See example below.

Note that a `nimbleExternalCall` can only be executed in a compiled `nimbleFunction`, not an uncompiled one.

If you have problems with spaces in file paths (e.g. for `oFile`), try compiling everything locally by including `dirName = "."` as an argument to `compileNimble`.

Value

A `nimbleFunction` that takes the indicated input arguments, calls Cfun, and returns the result.

Author(s)

Perry de Valpine

See Also

[nimbleRcall](#) for calling arbitrary R code from compiled `nimbleFunctions`.

Examples

```
## Not run:
sink('add1.h')
cat('
 extern "C" {
 void my_internal_function(double *p, double*ans, int n);
 }
')
sink()
sink('add1.cpp')
cat('
 #include <cstdio>
```

```
 #include "add1.h"
 void my_internal_function(double *p, double *ans, int n) {
   printf("In my_internal_function\\n");
     /* cat reduces the double slash to single slash */
   for(int i = 0; i < n; i++)
     ans[i] = p[i] + 1.0;
 }
')
sink()
system('g++ add1.cpp -c -o add1.o')
Radd1 <- nimbleExternalCall(function(x = double(1), ans = double(1),
n = integer()){}, Cfun =  'my_internal_function',
headerFile = file.path(getwd(), 'add1.h'), returnType = void(),
oFile = file.path(getwd(), 'add1.o'))
## If you need to use a function with non-scalar return object in model code,
## you can wrap it  in another nimbleFunction like this:
model_add1 <- nimbleFunction(
     run = function(x = double(1)) {
         ans <- numeric(length(x))
         Radd1(x, ans, length(x))
         return(ans)
         returnType(double(1))
     })
demoCode <- nimbleCode({
     for(i in 1:4) {x[i] ~ dnorm(0,1)} ## just to get a vector
     y[1:4] <- model_add1(x[1:4])
})
demoModel <- nimbleModel(demoCode, inits = list(x = rnorm(4)),
check = FALSE, calculate = FALSE)
CdemoModel <- compileNimble(demoModel, showCompilerOutput = TRUE)

## End(Not run)
```

---

nimbleFunction   *create a nimbleFunction*

Description

create a nimbleFunction from a setup function, run function, possibly other methods, and possibly inheritance via contains

Usage

```
nimbleFunction(
  setup = NULL,
  run = function() { },
  methods = list(),
  globalSetup = NULL,
  contains = NULL,
  enableDerivs = list(),
```

```
    name = NA,
    check = getNimbleOption("checkNimbleFunction"),
    where = getNimbleFunctionEnvironment()
)
```

Arguments

setup	An optional R function definition for setup processing.
run	An optional NIMBLE function definition that executes the primary job of the nimbleFunction
methods	An optional named list of NIMBLE function definitions for other class methods.
globalSetup	For internal use only
contains	An optional object returned from [nimbleFunctionVirtual](#) that defines arguments and returnTypes for run and/or methods, to which the current nimbleFunction must conform
enableDerivs	EXPERIMENTAL A list of names of function methods to enable derivatives for. Currently only for developer use.
name	An optional name used internally, for example in generated C++ code. Usually this is left blank and NIMBLE provides a name.
check	Boolean indicating whether to check the run code for function calls that NIMBLE cannot compile. Checking can be turned off for all calls to nimbleFunction using nimbleOptions(checkNimbleFunction = FALSE).
where	An optional where argument passed to setRefClass for where the reference class definition generated for this nimbleFunction will be stored. This is needed due to R package namespace issues but should never need to be provided by a user.

Details

This is the main function for defining nimbleFunctions. A lot of information is provided in the NIMBLE User Manual, so only a brief summary will be given here.

If a setup function is provided, then nimbleFunction returns a generator: a function that when called with arguments for the setup function will execute that function and return a specialized nimbleFunction. The run and other methods can be called using $ like in other R classes, e.g. nf$run(). The methods can use objects that were created in or passed to the setup function.

If no setup function is provided, then nimbleFunction returns a function that executes the run function. It is not a generator in this case, and no other methods can be provided.

If one wants a generator but does not need any setup arguments or code, setup = TRUE can be used.

See the NIMBLE User Manual for examples.

For more information about the contains argument, see the section on nimbleFunctionLists.

Author(s)

NIMBLE development team

```
nimbleFunctionBase-class
```
                        Class `nimbleFunctionBase`

Description

Classes used internally in NIMBLE and not expected to be called directly by users.

```
nimbleFunctionList-class
```
                        Create a list of nimbleFunctions

Description

Create an empty list of nimbleFunctions that all will inherit from a base class.

Details

See the User Manual for information about creating and populating a `nimbleFunctionList`.

Author(s)

NIMBLE development team

`nimbleFunctionVirtual` *create a virtual nimbleFunction, a base class for other nimbleFunctions*

Description

define argument types and returnType for the `run` function and any `methods`, to be used in the `contains` argument of `nimbleFunction`

Usage

```
nimbleFunctionVirtual(
  contains = NULL,
  run = function() { },
  methods = list(),
  name = NA
)
```

Arguments

`contains`	Not yet functional
`run`	A NIMBLE function that will only be used to inspect its argument types and returnType.
`methods`	An optional named list of NIMBLE functions that will also only be used for inspecting argument types and returnTypes.
`name`	An optional name used internally by the NIMBLE compiled. This is usually omitted and NIMBLE provides one.

Details

See the NIMBLE [User Manual](#) section on nimbleFunctionLists for explanation of how to use a virtual nimbleFunction.

Value

An object that can be passed as the `contains` argument to `nimbleFunction` or as the argument to `nimbleFunctionList`

Author(s)

NIMBLE development team

See Also

[`nimbleFunction`](#)

`nimbleList`	*create a nimbleList*

Description

create a nimbleList from a nimbleList definition

Usage

```
nimbleList(
  ...,
  name = NA,
  predefined = FALSE,
  where = getNimbleFunctionEnvironment()
)
```

Arguments

`...`	arbitrary set of names and types for the elements of the list or a single R list of type `nimbleType`.
`name`	optional character providing a name used internally, for example in generated C++ code. Usually this is left blank and NIMBLE provides a name.
`predefined`	logical for internal use only.
`where`	optional argument passed to `setRefClass` for where the reference class definition generated for this nimbleFunction will be stored. This is needed due to R package namespace issues but should never need to be provided by a user.

Details

This function creates a definition for a nimbleList. The `types` argument defines the names, types, and dimensions of the elements of the nimbleList. Elements of nimbleLists can be either basic types (e.g., `integer`, `double`) or other nimbleList definitions. The `types` argument can be either a series of expressions of the form `name = type(dim)`, or a list of `nimbleType` objects.

`nimbleList` returns a definition, which can be used to create instances of this type of nimbleList via the `new()` member function.

Definitions can be created in R's general environment or in nimbleFunction setup code. Instances can be created using the `new()` function in R's global environment, in nimbleFunction setup code, or in nimbleFunction run code.

Instances of `nimbleList` definitions can be used as arguments to run code of nimbleFunctions, and as the return type of nimbleFunctions.

Author(s)

NIMBLE development team

Examples

```
exampleNimListDef <- nimbleList(x = integer(0), Y = double(2))

nimbleListTypes <- list(nimbleType(name = 'x', type = 'integer', dim = 0),
                        nimbleType(name = 'Y', type = 'double', dim = 2))

## this nimbleList definition is identical to the one created above
exampleNimListDef <- nimbleList(nimbleListTypes)
```

---

nimbleMCMC	*Executes one or more chains of NIMBLE's default MCMC algorithm, for a model specified using BUGS code*

Description

nimbleMCMC is designed as the most straight forward entry point to using NIMBLE's default MCMC algorithm. It provides capability for running multiple MCMC chains, specifying the number of MCMC iterations, thinning, and burn-in, and which model variables should be monitored. It also provides options to return the posterior samples, to return summary statistics calculated from the posterior samples, and to return a WAIC value.

Usage

```
nimbleMCMC(
  code,
  constants = list(),
  data = list(),
  inits,
  dimensions = list(),
  model,
  monitors,
  thin = 1,
  niter = 10000,
  nburnin = 0,
  nchains = 1,
  check = TRUE,
  setSeed = FALSE,
  progressBar = getNimbleOption("MCMCprogressBar"),
  samples = TRUE,
  samplesAsCodaMCMC = FALSE,
  summary = FALSE,
  WAIC = FALSE
)
```

Arguments

code	The quoted code expression representing the model, such as the return value from a call to nimbleCode). Not required if model is provided.
constants	Named list of constants in the model. Constants cannot be subsequently modified. For compatibility with JAGS and BUGS, one can include data values with constants and nimbleModel will automatically distinguish them based on what appears on the left-hand side of expressions in code.
data	Named list of values for the data nodes. Data values can be subsequently modified. Providing this argument also flags nodes as having data for purposes of algorithms that inspect model structure. Values that are NA will not be flagged as data.
inits	Argument to specify initial values for the model object, and for each MCMC chain. See details.
dimensions	Named list of dimensions for variables. Only needed for variables used with empty indices in model code that are not provided in constants or data.

model	A compiled or uncompiled NIMBLE model object. When provided, this model will be used to configure the MCMC algorithm to be executed, rather than using the code, constants, data and inits arguments to create a new model object. However, if also provided, the inits argument will still be used to initialize this model prior to running each MCMC chain.
monitors	A character vector giving the node names or variable names to monitor. The samples corresponding to these nodes will returned, and/or will have summary statistics calculated. Default value is all top-level stochastic nodes of the model.
thin	Thinning interval for collecting MCMC samples. Thinning occurs after the initial nburnin samples are discarded. Default value is 1.
niter	Number of MCMC iterations to run. Default value is 10000.
nburnin	Number of initial, pre-thinning, MCMC iterations to discard. Default value is 0.
nchains	Number of MCMC chains to run. Default value is 1.
check	Logical argument, specifying whether to check the model object for missing or invalid values. Default value is TRUE.
setSeed	Logical or numeric argument. If a single numeric value is provided, R's random number seed will be set to this value at the onset of each MCMC chain. If a numeric vector of length nchains is provided, then each element of this vector is provided as R's random number seed at the onset of the corresponding MCMC chain. Otherwise, in the case of a logical value, if TRUE, then R's random number seed for the ith chain is set to be i, at the onset of each MCMC chain. Note that specifying the argument setSeed = 0 does not prevent setting the RNG seed, but rather sets the random number generation seed to 0 at the beginning of each MCMC chain. Default value is FALSE.
progressBar	Logical argument. If TRUE, an MCMC progress bar is displayed during execution of each MCMC chain. Default value is defined by the nimble package option MCMCprogressBar..
samples	Logical argument. If TRUE, then posterior samples are returned from each MCMC chain. These samples are optionally returned as coda mcmc objects, depending on the samplesAsCodaMCMC argument. Default value is TRUE. See details.
samplesAsCodaMCMC	Logical argument. If TRUE, then a coda mcmc object is returned instead of an R matrix of samples, or when nchains > 1 a coda mcmc.list object is returned containing nchains mcmc objects. This argument is only used when samples is TRUE. Default value is FALSE. See details.
summary	Logical argument. When TRUE, summary statistics for the posterior samples of each parameter are also returned, for each MCMC chain. This may be returned in addition to the posterior samples themselves. Default value is FALSE. See details. z
WAIC	Logical argument. When TRUE, the WAIC (Watanabe, 2010) of the model is calculated and returned. If multiple chains are run, then a single WAIC value is calculated using the posterior samples from all chains. Default value is FALSE. Note that the version of WAIC used is the default WAIC conditional on random effects/latent states and without any grouping of data nodes. See help(waic) for more details. If a different version of WAIC is desired, use runMCMC instead of nimbleMCMC.

Details

The entry point for this function is providing the `code`, `constants`, `data` and `inits` arguments, to create a new NIMBLE model object, or alternatively providing an exisiting NIMBLE model object as the `model` argument.

At least one of `samples`, `summary` or `WAIC` must be `TRUE`, since otherwise, nothing will be returned. Any combination of these may be `TRUE`, including possibly all three, in which case posterior samples, summary statistics, and WAIC values are returned for each MCMC chain.

When `samples = TRUE`, the form of the posterior samples is determined by the `samplesAsCodaMCMC` argument, as either matrices of posterior samples, or `coda mcmc` and `mcmc.list` objects.

Posterior summary statistics are returned individually for each chain, and also as calculated from all chains combined (when `nchains > 1`).

The `inits` argument can be one of three things:

(1) a function to generate initial values, which will be executed once to initialize the model object, and once to generate initial values at the beginning of each MCMC chain, or (2) a single named list of initial values which, will be used to initialize the model object and for each MCMC chain, or (3) a list of length `nchains`, each element being a named list of initial values. The first element will be used to initialize the model object, and once element of the list will be used for each MCMC chain.

The `inits` argument may also be omitted, in which case the model will not be provided with initial values. This is not recommended.

The `niter` argument specifies the number of pre-thinning MCMC iterations, and the `nburnin` argument specifies the number of pre-thinning MCMC samples to discard. After discarding these burn-in samples, thinning of the remaining samples will take place. The total number of posterior samples returned will be floor((niter-nburnin)/thin).

Value

A list is returned with named elements depending on the arguments passed to `nimbleMCMC`, unless only one among samples, summary, and WAIC are requested, in which case only that element is returned. These elements may include `samples`, `summary`, and `WAIC`. When `nchains = 1`, posterior samples are returned as a single matrix, and summary statistics as a single matrix. When `nchains > 1`, posterior samples are returned as a list of matrices, one matrix for each chain, and summary statistics are returned as a list containing `nchains+1` matrices: one matrix corresponding to each chain, and the final element providing a summary of all chains, combined. If `samplesAsCodaMCMC` is `TRUE`, then posterior samples are provided as `coda mcmc` and `mcmc.list` objects. When `WAIC` is `TRUE`, a WAIC summary object is returned.

Author(s)

Daniel Turek

See Also

[configureMCMC](#) [buildMCMC](#) [runMCMC](#)

Examples

```
## Not run:
code <- nimbleCode({
    mu ~ dnorm(0, sd = 1000)
    sigma ~ dunif(0, 1000)
    for(i in 1:10) {
        x[i] ~ dnorm(mu, sd = sigma)
    }
})
data <- list(x = c(2, 5, 3, 4, 1, 0, 1, 3, 5, 3))
inits <- function() list(mu = rnorm(1,0,1), sigma = runif(1,0,10))
mcmc.output <- nimbleMCMC(code, data = data, inits = inits,
                          monitors = c("mu", "sigma"), thin = 10,
                          niter = 20000, nburnin = 1000, nchains = 3,
                          summary = TRUE, WAIC = TRUE)

## End(Not run)
```

---

nimbleModel   *Create a NIMBLE model from BUGS code*

Description

processes BUGS model code and optional constants, data, and initial values. Returns a NIMBLE model (see [modelBaseClass](#)) or model definition.

Usage

```
nimbleModel(
  code,
  constants = list(),
  data = list(),
  inits = list(),
  dimensions = list(),
  returnDef = FALSE,
  where = globalenv(),
  debug = FALSE,
  check = getNimbleOption("checkModel"),
  calculate = TRUE,
  name = NULL,
  userEnv = parent.frame()
)
```

Arguments

code	code for the model in the form returned by nimbleCode or (equivalently) quote
constants	named list of constants in the model. Constants cannot be subsequently modified. For compatibility with JAGS and BUGS, one can include data values with constants and nimbleModel will automatically distinguish them based on what appears on the left-hand side of expressions in code.
data	named list of values for the data nodes. Data values can be subsequently modified. Providing this argument also flags nodes as having data for purposes of algorithms that inspect model structure. Values that are NA will not be flagged as data.
inits	named list of starting values for model variables. Unlike JAGS, should only be a single list, not a list of lists.
dimensions	named list of dimensions for variables. Only needed for variables used with empty indices in model code that are not provided in constants or data.
returnDef	logical indicating whether the model should be returned (FALSE) or just the model definition (TRUE).
where	argument passed to setRefClass, indicating the environment in which the reference class definitions generated for the model and its modelValues should be created. This is needed for managing package namespace issues during package loading and does not normally need to be provided by a user.
debug	logical indicating whether to put the user in a browser for debugging. Intended for developer use.
check	logical indicating whether to check the model object for missing or invalid values. Default is given by the NIMBLE option 'checkModel'. See nimbleOptions for details.
calculate	logical indicating whether to run calculate on the model after building it; this will calculate all deterministic nodes and logProbability values given the current state of all nodes. Default is TRUE. For large models, one might want to disable this, but note that deterministic nodes, including nodes introduced into the model by NIMBLE, may be NA.
name	optional character vector giving a name of the model for internal use. If omitted, a name will be provided.
userEnv	environment in which if-then-else statements in BUGS code will be evaluated if needed information not found in constants; intended primarily for internal use only

Details

See the User Manual or help(modelBaseClass) for information about manipulating NIMBLE models created by nimbleModel, including methods that operate on models, such as getDependencies.

The user may need to provide dimensions for certain variables as in some cases NIMBLE cannot automatically determine the dimensions and sizes of variables. See the User Manual for more information.

As noted above, one may lump together constants and data (as part of the `constants` argument (unlike R interfaces to JAGS and BUGS where they are provided as the `data` argument). One may not provide lumped constants and data as the `data` argument.

For variables that are a mixture of data nodes and non-data nodes, any values passed in via `inits` for components of the variable that are data will be ignored. All data values should be passed in through `data` (or `constants` as just discussed).

Author(s)

NIMBLE development team

See Also

[readBUGSmodel](#) for creating models from BUGS-format model files

Examples

```
code <- nimbleCode({
    x ~ dnorm(mu, sd = 1)
    mu ~ dnorm(0, sd = prior_sd)
})
constants = list(prior_sd = 1)
data = list(x = 4)
Rmodel <- nimbleModel(code, constants = constants, data = data)
```

---

nimbleOptions *NIMBLE Options Settings*

Description

Allow the user to set and examine a variety of global _options_ that affect the way in which NIMBLE operates. Call `nimbleOptions()` with no arguments to see a list of available opions.

Usage

```
nimbleOptions(...)
```

Arguments

`...`   any options to be defined as one or more `name = value` pairs or as a single `list` of `name=value` pairs.

Details

`nimbleOptions` mimics `options`. Invoking `nimbleOptions()` with no arguments returns a list with the current values of the options. To access the value of a single option, one should use `getNimbleOption()`.

Value

When invoked with no arguments, returns a list with the current values of all options. When invoked with one or more arguments, returns a list of the the updated options with their updated values.

Author(s)

Christopher Paciorek

Examples

```
# Set one option:
nimbleOptions(verifyConjugatePosteriors = FALSE)

# Compactly print all options:
str(nimbleOptions(), max.level = 1)

# Save-and-restore options:
old <- nimbleOptions()                   # Saves old options.
nimbleOptions(showCompilerOutput = TRUE,
              verboseErrors = TRUE)      # Sets temporary options.
# ...do stuff...
nimbleOptions(old)                       # Restores old options.
```

---

nimbleRcall   *Make an R function callable from compiled nimbleFunctions (including nimbleModels).*

Description

Normally compiled nimbleFunctions call other compiled nimbleFunctions. nimbleRcall enables any R function (with viable argument types and return values) to be called (and evaluated in R) from compiled nimbleFunctions.

Usage

```
nimbleRcall(
  prototype,
  returnType,
  Rfun,
  where = getNimbleFunctionEnvironment()
)
```

Arguments

prototype   Argument type information for Rfun. This can be provided as an R function
                   using `nimbleFunction` type declarations or as a list of `nimbleType` objects.

returnType	Return object type information. This can be provided similarly to `prototype` as either a `nimbleFunction` type declaration or as a `nimbleType` object. In the latter case, the name will be ignored. If there is no return value this should be `void()`.
Rfun	The name of an R function to be called from compiled nimbleFunctions.
where	An optional `where` argument passed to `setRefClass` for where the reference class definition generated for this nimbleFunction will be stored. This is needed due to R package namespace issues but should never need to be provided by a user.

Details

The `nimbleFunction` returned by nimbleRcall can be used in other `nimbleFunctions`. When called from a compiled `nimbleFunction` (including from a model), arguments will be copied according to the declared types, the function named by Rfun will be called, and the returned object will be copied if necessary. The example below shows use of an R function in a compiled `nimbleModel`.

A `nimbleFunction` returned by nimbleRcall can only be used in a compiled `nimbleFunction`. Rfun itself should work in an uncompiled `nimbleFunction`.

Value

A `nimbleFunction` that wraps a call to Rfun with type-declared arguments and return object.

Author(s)

Perry de Valpine

See Also

[nimbleExternalCall](#) for calling externally provided C (or other) compiled code.

Examples

```
## Not run:
## Say we want an R function that adds 2 to every value in a vector
add2 <- function(x) {
    x + 2
}
Radd2 <- nimbleRcall(function(x = double(1)){}, Rfun = 'add2',
returnType = double(1))
demoCode <- nimbleCode({
    for(i in 1:4) {x[i] ~ dnorm(0,1)}
    z[1:4] <- Radd2(x[1:4])
})
demoModel <- nimbleModel(demoCode, inits = list(x = rnorm(4)),
check = FALSE, calculate = FALSE)
CdemoModel <- compileNimble(demoModel)

## End(Not run)
```

---

nimbleType-class   *create a nimbleType object*

Description

Create a nimbleType object, with information on the name, type, and dimension of an object to be placed in a [nimbleList](#).

Arguments

name	The name of the object, given as a character string.
type	The type of the object, given as a character string.
dim	The dimension of the object, given as an integer. This can be left blank if the object is a nimbleList.

Details

This function creates `nimbleType` objects, which can be used to define the elements of a [nimbleList](#).

The `type` argument can be chosen from among `character`, `double`, `integer`, and `logical`, or can be the name of a previously created [nimbleList](#) definition.

See the NIMBLE [User Manual](#) for additional examples.

Author(s)

NIMBLE development team

Examples

```
nimbleTypeList <- list()
nimbleTypeList[[1]] <- nimbleType(name = 'x', type = 'integer', dim = 0)
nimbleTypeList[[2]] <- nimbleType(name = 'Y', type = 'double', dim = 2)
```

---

nimCat   *cat function for use in nimbleFunctions*

Description

cat function for use in nimbleFunctions

Usage

```
nimCat(...)
```

Arguments

`...`	an arbitrary set of arguments that will be printed in sequence.

Details

`cat` in nimbleFunction run-code imitates the R function [`cat`](#). It prints its arguments in order. No newline is inserted, so include "\n" if one is desired.

When an uncompiled nimbleFunction is executed, R's `cat` is used. In a compiled nimbleFunction, a C++ output stream is used that will generally format output similarly to R's `cat`. Non-scalar numeric objects can be included, although their output will be formatted slightly different in uncompiled and compiled nimbleFunctions.

In nimbleFunction run-time code, cat is identical to `print` except the latter appends a newline at the end.

`nimCat` is the same as `cat`, and the latter is converted to the former when a nimbleFunction is defined.

See Also

[`print`](#)

Examples

```
ans <- matrix(1:4, nrow = 2) ## R code, not NIMBLE code
nimCat('Answer is ', ans) ## would work in R or NIMBLE
```

---

nimCopy	*Copying function for NIMBLE*

Description

Copies values from a NIMBLE model or modelValues object to another NIMBLE model or modelValues. Work in R and NIMBLE. The NIMBLE keyword copy is identical to `nimCopy`

Usage

```
nimCopy(
  from,
  to,
  nodes = NULL,
  nodesTo = NULL,
  row = NA,
  rowTo = NA,
  logProb = FALSE,
  logProbOnly = FALSE
)
```

Arguments

from	Either a NIMBLE model or modelValues object
to	Either a NIMBLE model or modelValues object
nodes	Vector of one or more node names of object from that will be copied from
nodesTo	Vector of one or more node names of object to that will be copied to. If nodesTo is NULL, will automatically be set to nodes
row	If from is a modelValues, the row that will be copied from
rowTo	If to is a modelValues, the row which will be copied to. If rowTo is NA, will automatically be set to row
logProb	A logical value indicating whether the log probabilities of the given nodes should also be copied (i.e. if nodes = 'x' and logProb = TRUE, then both 'x' and 'logProb_x' will be copied)
logProbOnly	A logical value indicating whether only the log probabilities of the given nodes should be copied (i.e. if nodes = 'x' and logProbOnly = TRUE, then only 'logProb_x' will be copied)

Details

This function copies values from one or more nodes (possibly including log probabilities for nodes) between models and modelValues objects. For modelValues objects, the row must be specified. This function allows one to conveniently copy multiple nodes, avoiding having to write a loop.

Author(s)

Clifford Anderson-Bergman

Examples

```
# Building model and modelValues object
simpleModelCode <- nimbleCode({
for(i in 1:100)
x[i] ~ dnorm(0,1)
})
rModel <- nimbleModel(simpleModelCode)
rModelValues <- modelValues(rModel)

#Setting model nodes
rModel$x <- rnorm(100)
#Using nimCopy in R.
nimCopy(from = rModel, to = rModelValues, nodes = 'x', rowTo = 1)

#Use of nimCopy in a simple nimbleFunction
cCopyGen <- nimbleFunction(
setup = function(model, modelValues, nodeNames){},
run = function(){
nimCopy(from = model, to = modelValues, nodes = nodeNames, rowTo = 1)
}
)
```

```
rCopy <- cCopyGen(rModel, rModelValues, 'x')
## Not run:
cModel <- compileNimble(rModel)
cCopy <- compileNimble(rCopy, project = rModel)
cModel[['x']] <- rnorm(100)

cCopy$run() ## execute the copy with the compiled function

## End(Not run)
```

---

nimDerivs   *Nimble Derivatives*

Description

EXPERIMENTAL Computes the value, gradient, and Hessian of a given `nimbleFunction` method. The R version is currently unimplemented.

Usage

```
nimDerivs(nimFxn = NA, order = nimC(0, 1, 2))
```

Arguments

nimFxn   a call to a `nimbleFunction` method with arguments included.

order   an integer vector with values within the set 0, 1, 2, corresponding to whether the function value, gradient, and Hessian should be returned respectively.

---

nimDim   *return sizes of an object whether it is a vector, matrix or array*

Description

R's regular `dim` function returns NULL for a vector. It is useful to have this function that treats a vector similarly to a matrix or array. Works in R and NIMBLE. In NIMBLE `dim` is identical to `nimDim`, not to R's `dim`

Usage

```
nimDim(obj)
```

Arguments

obj   objects for which the sizes are requested

Value

a vector of sizes in each dimension

Author(s)

NIMBLE development team

Examples

```
x <- rnorm(4)
dim(x)
nimDim(x)
y <- matrix(x, nrow = 2)
dim(y)
nimDim(y)
```

---

nimEigen   *Spectral Decomposition of a Matrix*

Description

Computes eigenvalues and eigenvectors of a numeric matrix.

Usage

```
nimEigen(x, symmetric = FALSE, only.values = FALSE)
```

Arguments

x	a numeric matrix (double or integer) whose spectral decomposition is to be computed.
symmetric	if TRUE, the matrix is guarranteed to be symmetric, and only its lower triangle (diagonal included) is used. Otherwise, the matrix is checked for symmetry. Default is FALSE.
only.values	if TRUE, only the eigenvalues are computed, otherwise both eigenvalues and eigenvectors are computed. Setting only.values = TRUE can speed up eigendecompositions, especially for large matrices. Default is FALSE.

Details

Computes the spectral decomposition of a numeric matrix using the Eigen C++ template library. In a nimbleFunction, eigen is identical to nimEigen. If the matrix is symmetric, a faster and more accurate algorithm will be used to compute the eigendecomposition. Note that non-symmetric matrices can have complex eigenvalues, which are not supported by NIMBLE. If a complex eigenvalue or a complex element of an eigenvector is detected, a warning will be issued and that element will be returned as NaN.

Additionally, returnType(eigenNimbleList()) can be used within a link{nimbleFunction} to specify that the function will return a nimbleList generated by the nimEigen function. eigenNimbleList() can also be used to define a nested nimbleList element. See the User Manual for usage examples.

Value

The spectral decomposition of x is returned as a nimbleList with elements:

- values vector containing the eigenvalues of x, sorted in decreasing order. Since x is required to be symmetric, all eigenvalues will be real numbers.
- vectors. matrix with columns containing the eigenvectors of x, or an empty matrix if only.values is TRUE.

Author(s)

NIMBLE development team

See Also

nimSvd for singular value decompositions in NIMBLE.

Examples

```
eigenvaluesDemoFunction <- nimbleFunction(
  setup = function(){
    demoMatrix <- diag(4) + 2
  },
  run = function(){
    eigenvalues <- eigen(demoMatrix, symmetric = TRUE)$values
    returnType(double(1))
    return(eigenvalues)
})
```

---

nimMatrix   *Creates matrix or array objects for use in nimbleFunctions*

Description

In a nimbleFunction, matrix and array are identical to nimMatrix and nimArray, respectively

Usage

```
nimMatrix(
  value = 0,
  nrow = NA,
  ncol = NA,
  init = TRUE,
  fillZeros = TRUE,
```

```
  recycle = TRUE,
  type = "double"
)

nimArray(
  value = 0,
  dim = c(1, 1),
  init = TRUE,
  fillZeros = TRUE,
  recycle = TRUE,
  nDim,
  type = "double"
)
```

Arguments

value	value(s) for initialization (default = 0). This can be a vector, matrix or array, but it will be used as a vector.
nrow	the number of rows in a matrix (default = 1)
ncol	the number of columns in a matrix (default = 1)
init	logical, whether to initialize values (default = TRUE)
fillZeros	logical, whether to initialize any elements not filled by (possibly recycled) value with 0 (or FALSE for nimLogical) (default = TRUE)
recycle	logical, whether value should be recycled to fill the entire contents of the new object (default = TRUE)
type	character representing the data type, i.e. 'double', 'integer', or 'logical' (default = 'double')
dim	vector of dimension sizes in an array (default = c(1,1))
nDim	number of dimensions in an array. This is only necessary for compileNimble if the length of dim cannot be determined during compilation.

Details

These functions are similar to R's [matrix](#) and [array](#) functions, but they can be used in a nimbleFunction and compiled using compileNimble. Largely for compilation purposes, finer control is provided over initialization behavior, similarly to [nimNumeric](#), [nimInteger](#), and [nimLogical](#). If init = FALSE, no initialization will be done, and value, fillZeros and recycle will be ignored. If init=TRUE and recycle=TRUE, then fillZeros will be ignored, and value will be repeated (according to R's recycling rule) as much as necessary to fill the object. If init=TRUE and recycle=FALSE, then if fillZeros=TRUE, values of 0 (or FALSE for nimLogical) will be filled in after value. Compiled code will be more efficient if unnecessary initialization is not done, but this may or may not be noticeable depending on the situation.

When used in a nimbleFunction (in run or other member function), matrix and array are immediately converted to nimMatrix and nimArray, respectively.

The nDim argument is only necessary for a use like dim <-c(2,3,4); A <-nimArray(0,dim = dim,nDim = 3). It is necessary because the NIMBLE compiler must determine during compilation that A will

be a 3-dimensional numeric array. However, the compiler doesn't know for sure what the length of `dim` will be at run time, only that it is a vector. On the other hand, `A <-nimArray(0,dim = c(2,3,4))` is allowed because the compiler can directly determine that a vector of length three is constructed inline for the `dim` argument.

Author(s)

Daniel Turek and Perry de Valpine

See Also

[nimNumeric](#) [nimInteger](#) [nimLogical](#)

---

nimNumeric   *Creates numeric, integer or logical vectors for use in nimbleFunctions*

Description

In a `nimbleFunction`, numeric, integer and logical are identical to `nimNumeric`, `nimInteger` and `nimLogical`, respectively.

Usage

```
nimNumeric(
  length = 0,
  value = 0,
  init = TRUE,
  fillZeros = TRUE,
  recycle = TRUE
)

nimInteger(
  length = 0,
  value = 0,
  init = TRUE,
  fillZeros = TRUE,
  recycle = TRUE
)

nimLogical(
  length = 0,
  value = 0,
  init = TRUE,
  fillZeros = TRUE,
  recycle = TRUE
)
```

Arguments

`length`	the length of the vector (default = 0)
`value`	value(s) for initializing the vector (default = 0). This may be a vector, matrix or array but will be used as a vector.
`init`	logical, whether to initialize elements of the vector (default = TRUE)
`fillZeros`	logical, whether to initialize any elements not filled by (possibly recycled) `value` with 0 (or FALSE for `nimLogical`) (default = TRUE)
`recycle`	logical, whether `value` should be recycled to fill the entire `length` of the new vector (default = TRUE)

Details

These functions are similar to R's [numeric](), [integer](), [logical]() functions, but they can be used in a nimbleFunction and then compiled using compileNimble. Largely for compilation purposes, finer control is provided over initialization behavior. If init = FALSE, no initialization will be done, and `value`, `fillZeros` and `recycle` will be ignored. If init=TRUE and recycle=TRUE, then `fillZeros` will be ignored, and `value` will be repeated (according to R's recycling rule) as much as necessary to fill a vector of length `length`. If init=TRUE and recycle=FALSE, then if fillZeros=TRUE, values of 0 (or FALSE for `nimLogical`) will be filled in after `value` up to length `length`. Compiled code will be more efficient if unnecessary initialization is not done, but this may or may not be noticeable depending on the situation.

When used in a `nimbleFunction` (in run or other member function), numeric, integer and logical are immediately converted to nimNumeric, nimInteger and nimLogical, respectively.

Author(s)

Daniel Turek, Chris Paciorek, Perry de Valpine

See Also

[nimMatrix](), [nimArray]()

---

nimOptim	*Nimble wrapper around R's builtin* [optim]().

Description

Nimble wrapper around R's builtin [optim]().

Usage

```
nimOptim(
  par,
  fn,
  gr = "NULL",
  ...,
  method = "Nelder-Mead",
  lower = -Inf,
  upper = Inf,
  control = nimOptimDefaultControl(),
  hessian = FALSE
)
```

Arguments

par	Initial values for the parameters to be optimized over.
fn	A function to be minimized (or maximized), with first argument the vector of parameters over which minimization is to take place. It should return a scalar result.
gr	A function to return the gradient for the "BFGS", "CG" and "L-BFGS-B" methods.
...	IGNORED
method	The method to be used. See 'Details' section of [optim](). One of: "Nelder-Mead", "BFGS", "CG", "L-BFGS-B". Note that the R methods "SANN", "Brent" are not supported.
lower	Vector or scalar of lower bounds for parameters.
upper	Vector or scalar of upper bounds for parameters.
control	A list of control parameters. See `Details` section of [optim]().
hessian	Logical. Should a Hessian matrix be returned?

Value

[optimResultNimbleList]()

See Also

[optim]()

Examples

```
## Not run:
objectiveFunction <- nimbleFunction(
    run = function(par = double(1)) {
        return(sum(par) * exp(-sum(par ^ 2) / 2))
        returnType(double(0))
    }
)
```

```
optimizer <- nimbleFunction(
    run = function(method = character(0), fnscale = double(0)) {
        control <- optimDefaultControl()
        control$fnscale <- fnscale
        par <- c(0.1, -0.1)
        return(optim(par, objectiveFunction, method = method, control = control))
        returnType(optimResultNimbleList())
    }
)
cOptimizer <- compileNimble(optimizer)
cOptimizer(method = 'BFGS', fnscale = -1)

## End(Not run)
```

nimOptimDefaultControl

Creates a deafult control *argument for* [nimOptim](.)

Description

Creates a deafult control argument for [nimOptim](.).

Usage

```
nimOptimDefaultControl()
```

Value

[optimControlNimbleList](.)

See Also

[nimOptim](.), [optim](.)

nimPrint   *print function for use in nimbleFunctions*

Description

print function for use in nimbleFunctions

Usage

```
nimPrint(...)
```

Arguments

...   an abitrary set of arguments that will be printed in sequence.

Details

The keyword `print` in nimbleFunction run-time code will be automatically turned into `nimPrint`. This is a function that prints its arguments in order using `cat` in R, or using `std::cout` in C++ code generated by compiling nimbleFunctions. Non-scalar numeric objects can be included, although their output will be formatted slightly different in uncompiled and compiled nimbleFunctions.

See Also

[cat](#)

Examples

```
ans <- matrix(1:4, nrow = 2) ## R code, not NIMBLE code
nimPrint('Answer is ', ans) ## would work in R or NIMBLE
```

| nimStop | *Halt execution of a nimbleFunction function method. Part of the NIM-BLE language* |

Description

Halt execution of a nimbleFunction function method. Part of the NIMBLE language

Usage

```
nimStop(msg)
```

Arguments

msg   Character object to be output as an error message

Details

The NIMBLE `stop` is similar to the native R `stop`, but it takes only one argument, the error message to be output. During uncompiled NIMBLE execution, `nimStop` simply calls R's stop funtion. During compiled execution it calls the error function from the R headers. `stop` is an alias for `nimStop` in the NIMBLE language

Author(s)

Perry de Valpine

---

nimSvd   *Singular Value Decomposition of a Matrix*

Description

Computes singular values and, optionally, left and right singular vectors of a numeric matrix.

Usage

```
nimSvd(x, vectors = "full")
```

Arguments

x   a symmetric numeric matrix (double or integer) whose spectral decomposition
               is to be computed.

vectors   character that determines whether to calculate left and right singular vectors.
               Can take values 'none', 'thin' or 'full'. Defaults to 'full'. See 'Details'.

Details

Computes the singular value decomposition of a numeric matrix using the Eigen C++ template
library.

The vectors character argument determines whether to compute no left and right singular vectors
('none'), thinned left and right singular vectors ('thin'), or full left and right singular vectors
('full'). For a matrix x with dimensions n and p, setting vectors = 'thin' will does the following
(quoted from eigen website): In case of a rectangular n-by-p matrix, letting m be the smaller value
among n and p, there are only m singular vectors; the remaining columns of U and V do not
correspond to actual singular vectors. Asking for thin U or V means asking for only their m first
columns to be formed. So U is then a n-by-m matrix, and V is then a p-by-m matrix. Notice that
thin U and V are all you need for (least squares) solving.

Setting vectors = 'full' will compute full matrices for U and V, so that U will be of size n-by-n,
and V will be of size p-by-p.

In a nimbleFunction, svd is identical to nimSvd.

returnType(svdNimbleList()) can be used within a link{nimbleFunction} to specify that the
function will return a nimbleList generated by the nimSvd function. svdNimbleList() can also
be used to define a nested nimbleList element. See the User Manual for usage examples.

Value

The singular value decomposition of x is returned as a nimbleList with elements:

- d length m vector containing the singular values of x, sorted in decreasing order.

- v matrix with columns containing the left singular vectors of x, or an empty matrix if vectors
  = 'none'.

- u matrix with columns containing the right singular vectors of x, or an empty matrix if
  vectors = 'none'.

Author(s)

NIMBLE development team

See Also

nimEigen for spectral decompositions.

Examples

```
singularValuesDemoFunction <- nimbleFunction(
  setup = function(){
    demoMatrix <- diag(4) + 2
  },
  run = function(){
    singularValues <- svd(demoMatrix)$d
    returnType(double(1))
    return(singularValues)
})
```

---

nodeFunctions   *calculate, calculateDiff, simulate, or get the current log probabilities (densities) a set of nodes in a NIMBLE model*

Description

calculate, calculateDiff, simulate, or get the current log probabilities (densities) of one or more nodes of a NIMBLE model and (for calculate and getLogProb) return the sum of their log probabilities (or densities). Part of R and NIMBLE.

Usage

```
calculate(model, nodes, nodeFxnVector, nodeFunctionIndex)

calculateDiff(model, nodes, nodeFxnVector, nodeFunctionIndex)

getLogProb(model, nodes, nodeFxnVector, nodeFunctionIndex)

simulate(model, nodes, includeData = FALSE, nodeFxnVector, nodeFunctionIndex)
```

Arguments

model	A NIMBLE model, either the compiled or uncompiled version
nodes	A character vector of node names, with index blocks allowed, such as 'x', 'y[2]', or 'z[1:3, 2:4]'
nodeFxnVector	An optional vector of nodeFunctions on which to operate, in lieu of model and nodes

nodeFunctionIndex

> For internal NIMBLE use only

includeData   A logical argument specifying whether data nodes should be simulated into
                   (only relevant for simulate)

Details

Standard usage is as a method of a model, in the form model$calculate(nodes), but the usage as
a simple function with the model as the first argument as above is also allowed.

These functions expands the nodes and then process them in the model in the order provided. Expanding nodes means turning 'y[1:2]' into c('y[1]','y[2]') if y is a vector of scalar nodes. Calculation is defined for a stochastic node as executing the log probability (density) calculation and for a deterministic node as calculating whatever function was provided on the right-hand side of the model declaration.

Difference calculation (calculateDiff) executes the operation(s) on the model as calculate, but it returns the sum of the difference between the new log probabilities and the previous ones.

Simulation is defined for a stochastic node as drawing a random value from its distribution, and for deterministic node as equivalent to calculate.

getLogProb collects and returns the sum of the log probabilities of nodes, using the log probability values currently stored in the model (as generated from the most recent call to calculate on each node)

These functions can be used from R or in NIMBLE run-time functions that will be compiled. When executed in R (including when an uncompiled nimbleFunction is executed), they can be slow because the nodes are expanded each time. When compiled in NIMBLE, the nodes are expanded only once during compilation, so execution will be much faster.

It is common to want the nodes to be provided in topologically sorted order, so that they will be calculated or simulated following the order of the model graph. Functions such as model$getDependencies(nodes, ...) return nodes in topologically sorted order. They can be directly sorted by model$topologicallySortNodes(nodes), but if so it is a good idea to expand names first by model$topologicallySortNodes(model$expandNodeNames(nodes))

Value

calculate and getLogProb return the sum of the log probabilities (densities) of the calculated nodes, with a contribution of 0 from any deterministic nodes

calculateDiff returns the sum of the difference between the new and old log probabilities (densities) of the calculated nodes, with a contribution of 0 from any deterministic nodes.

simulate returns NULL.

Author(s)

NIMBLE development team

optimControlNimbleList

EXPERIMENTAL Data type for the control *parameter of* nimOptim

Description

nimbleList definition for the type of nimbleList input as the control parameter to nimOptim. See optim for details.

Usage

```
optimControlNimbleList
```

Format

An object of class list of length 1.

See Also

optim, nimOptim

optimDefaultControl   *Creates a deafult* control *argument for* optim *(just an empty list).*

Description

Creates a deafult control argument for optim (just an empty list).

Usage

```
optimDefaultControl()
```

Value

an empty list.

See Also

nimOptim, optim

---

optimResultNimbleList   *EXPERIMENTAL Data type for the return value of* nimOptim

---

Description

nimbleList definition for the type of nimbleList returned by nimOptim.

Usage

```
optimResultNimbleList
```

Format

An object of class list of length 1.

Fields

par The best set of parameters found.

value The value of fn corresponding to par.

counts A two-element integer vector giving the number of calls to fn and gr respectively.

convergence An integer code. 0 indicates successful completion. Possible error codes are 1 indicates that the iteration limit maxit had been reached. 10 indicates degeneracy of the Nelder-Mead simplex. 51 indicates a warning from the "L-BFGS-B" method; see component message for further details. 52 indicates an error from the "L-BFGS-B" method; see component message for further details.

message A character string giving any additional information returned by the optimizer, or NULL.

hessian Only if argument hessian is true. A symmetric matrix giving an estimate of the Hessian at the solution found.

See Also

optim, nimOptim

---

printErrors   *Print error messages after failed compilation*

Description

Retrieves the error file from R's tempdir and prints to the screen.

Usage

```
printErrors(excludeWarnings = TRUE)
```

Arguments

excludeWarnings

logical indicating whether compiler warnings should be printed; generally such warnings can be ignored.

Author(s)

Christopher Paciorek

---

rankSample	*Generates a weighted sample (with replacement) of ranks*

Description

Takes a set of non-negative weights (do not need to sum to 1) and returns a sample with size elements of the integers 1:length(weights), where the probability of being sampled is proportional to the value of weights. An important note is that the output vector will be sorted in ascending order. Also, right now it works slightly odd syntax (see example below). Later releases of NIMBLE will contain more natural syntax.

Usage

```
rankSample(weights, size, output, silent = FALSE)
```

Arguments

weights	A vector of numeric weights. Does not need to sum to 1, but must be non-negative
size	Size of sample
output	An R object into which the values will be placed. See example below for proper use
silent	Logical indicating whether to suppress logging information

Details

If invalid weights provided (i.e. negative weights or weights sum to 1), sets output = rep(1, size) and prints warning. rankSample can be used inside nimble functions.

rankSample first samples from the joint distribution size uniform(0,1) distributions by conditionally sampling from the rank statistics. This leads to a sorted sample of uniform(0,1)'s. Then, a cdf vector is constructed from weights. Because the sample of uniforms is sorted, rankSample walks down the cdf in linear time and fills out the sample.

Author(s)

Clifford Anderson-Bergman

Examples

```
set.seed(1)
sampInts = NA #sampled integers will be placed in sampInts
rankSample(weights = c(1, 1, 2), size = 10, sampInts)
sampInts
# [1] 1 1 2 2 2 2 2 3 3 3
rankSample(weights = c(1, 1, 2), size = 10000, sampInts)
table(sampInts)
#sampInts
#   1    2    3
#2434 2492 5074

#Used in a nimbleFunction
sampGen <- nimbleFunction(setup = function(){
x = 1:2
},
run = function(weights = double(1), k = integer() ){
rankSample(weights, k, x)
returnType(integer(1))
return(x)
})
rSamp <- sampGen()
rSamp$run(1:4, 5)
#[1] 3 3 4 4 4
```

---

readBUGSmodel	*Create a NIMBLE BUGS model from a variety of input formats, in-cluding BUGS model files*

Description

readBUGSmodel processes inputs providing the model and values for constants, data, initial values of the model in a variety of forms, returning a NIMBLE BUGS R model

Usage

```
readBUGSmodel(
  model,
  data = NULL,
  inits = NULL,
  dir = NULL,
  useInits = TRUE,
  debug = FALSE,
  returnComponents = FALSE,
  check = getNimbleOption("checkModel"),
  calculate = TRUE
)
```

Arguments

model	one of (1) a character string giving the file name containing the BUGS model code, with relative or absolute path, (2) an R function whose body is the BUGS model code, or (3) the output of `nimbleCode`. If a file name, the file can contain a 'var' block and 'data' block in the manner of the JAGS versions of the BUGS examples but should not contain references to other input data files nor a const block. The '.bug' or '.txt' extension can be excluded.
data	(optional) (1) character string giving the file name for an R file providing the input constants and data as R code [assigning individual objects or as a named list], with relative or absolute path, or (2) a named list providing the input constants and data. If neither is provided, the function will look for a file named 'name_of_model-data' including extensions .R, .r, or .txt.
inits	(optional) (1) character string giving the file name for an R file providing starting values as R code [assigning individual objects or as a named list], with relative or absolute path, or (2) a named list providing the starting values. Unlike JAGS, this should provide a single set of starting values, and therefore if provided as a list should be a simple list and not a list of lists.
dir	(optional) character string giving the directory where the (optional) files are located
useInits	boolean indicating whether to set the initial values, either based on `inits` or by finding the '-inits' file corresponding to the input model file
debug	logical indicating whether to put the user in a browser for debugging when `readBUGSmodel` calls `nimbleModel`. Intended for developer use.
returnComponents	
	logical indicating whether to return pieces of model object without building the model. Default is FALSE.
check	logical indicating whether to check the model object for missing or invalid values. Default is given by the NIMBLE option 'checkModel'. See `nimbleOptions` for details.
calculate	logical indicating whether to run `calculate` on the model after building it; this will calculate all deterministic nodes and logProbability values given the current state of all nodes. Default is TRUE. For large models, one might want to disable this, but note that deterministic nodes, including nodes introduced into the model by NIMBLE, may be `NA`.

Details

Note that `readBUGSmodel` should handle most common ways of providing information on a model as used in BUGS and JAGS but does not handle input model files that refer to additional files containing data. Please see the BUGS examples provided with NIMBLE in the `classic-bugs` directory of the installed NIMBLE package or JAGS (https://sourceforge.net/projects/mcmc-jags/files/Examples/) for examples of supported formats. Also, `readBUGSmodel` takes both constants and data via the 'data' argument, unlike `nimbleModel`, in which these are distinguished. The reason for allowing both to be given via 'data' is for backwards compatibility with the BUGS examples, in which constants and data are not distinguished.

Value

returns a NIMBLE BUGS R model

Author(s)

Christopher Paciorek

See Also

[nimbleModel](#)

Examples

```
## Reading a model defined in the R session

code <- nimbleCode({
    x ~ dnorm(mu, sd = 1)
    mu ~ dnorm(0, sd = prior_sd)
})
data = list(prior_sd = 1, x = 4)
model <- readBUGSmodel(code, data = data, inits = list(mu = 0))
model$x
model[['mu']]
model$calculate('x')

## Reading a classic BUGS model

pumpModel <- readBUGSmodel('pump.bug', dir = getBUGSexampleDir('pump'))
pumpModel$getVarNames()
pumpModel$x
```

---

registerDistributions   *Add user-supplied distributions for use in NIMBLE BUGS models*

Description

Register distributional information so that NIMBLE can process user-supplied distributions in BUGS model code

Usage

```
registerDistributions(
  distributionsInput,
  userEnv = parent.frame(),
  verbose = nimbleOptions("verbose")
)
```

Arguments

`distributionsInput`

> either a list or character vector specifying the user-supplied distributions. If a list, it should be a named list of lists in the form of that shown in `nimble:::distributionsInputList` with each list having required field `BUGSdist` and optional fields `Rdist`, `altParams`, `discrete`, `pqAvail`, `types`, and with the name of the list the same as that of the density function. Alternatively, simply a character vector providing the names of the density functions for the user-supplied distributions.

`userEnv`
> environment in which to look for the nimbleFunctions that provide the distribution; this will generally not need to be set by the user as it will default to the environment from which this function was called.

`verbose`
> logical indicating whether to print additional logging information

Details

When `distributionsInput` is a list of lists, see below for more information on the structure of the list. When `distributionsInput` is a character vector, the distribution is assumed to be of standard form, with parameters assumed to be the arguments provided in the density nimbleFunction, no alternative parameterizations, and the distribution assumed to be continuous with range from minus infinity to infinity. The availability of distribution and quantile functions is inferred from whether appropriately-named functions exist in the global environment.

One usually does not need to explicitly call `registerDistributions` as it will be called automatically when the user-supplied distribution is used for the first time in BUGS code. However, if one wishes to provide alternative parameterizations, to provide a range, or to indicate a distribution is discrete, then one still must explicitly register the distribution using `registerDistributions` with the argument in the list format.

Format of the component lists when `distributionsInput` is a list of lists:

- `BUGSdist` a character string in the form of the density name (starting with 'd') followed by the names of the parameters in parentheses. When alternative parameterizations are given in `Rdist`, this should be an exhaustive list of the unique parameter names from all possible parameterizations, with the default parameters specified first.

- `Rdist` an optional character vector with one or more alternative specifications of the density; each alternative specification can be an alternative name for the density, a different ordering of the parameters, different parameter name(s), or an alternative parameterization. In the latter case, the character string in parentheses should provide a given reparameterization as comma-separated name = value pairs, one for each default parameter, where name is the name of the default parameter and value is a mathematical expression relating the default parameter to the alternative parameters or other default parameters. The default parameters should correspond to the input arguments of the nimbleFunctions provided as the density and random generation functions. The mathematical expression can use any of the math functions allowed in NIM-BLE (see the [User Manual](#)) as well as user-supplied nimbleFunctions (which must have no setup code). The names of your nimbleFunctions for the distribution functions must match the function name in the `Rdist` entry (or if missing, the function name in the `BUGSdist` entry

- `discrete` a optional logical indicating if the distribution is that of a discrete random variable. If not supplied, distribution is assumed to be for a continuous random variable.

- `pqAvail` an optional logical indicating if distribution (CDF) and quantile (inverse CDF) functions are provided as nimbleFunctions. These are required for one to be able to use truncated versions of the distribution. Only applicable for univariate distributions. If not supplied, assumed to be FALSE.

- `altParams` a character vector of comma-separated 'name = value' pairs that provide the mathematical expressions relating non-canonical parameters to canonical parameters (canonical parameters are those passed as arguments to your distribution functions). These inverse functions are used for MCMC conjugacy calculations when a conjugate relationship is expressed in terms of non-default parameters (such as the precision for normal-normal conjugacy). If not supplied, the system will still function but with a possible loss of efficiency in certain algorithms.

- `types` a character vector of comma-separated 'name = input' pairs indicating the type and dimension of the random variable and parameters (including default and alternative parameters). 'input' should take the form 'double(d)' or 'integer(d)', where 'd' is 0 for scalars, 1 for vectors, 2 for matrices. Note that since NIMBLE uses doubles for numerical calculations and the default type is double(0), one should generally use 'double' and one need only specify the type for non-scalars. 'name' should be either 'value' to indicate the random variable itself or the parameter name to indicate a given parameter.

- `range` a vector of two values giving the range of the distribution for possible use in future algorithms (not used currently). When the lower or upper limit involves a strict inequality (e.g., $x>0$), you should simply treat it as a non-strict inequality ($x>=0$, and set the lower value to 0). Also we do not handle ranges that are functions of parameters, so simply use the smallest/largest possible values given the possible parameter values. If not supplied this is taken to be (-Inf,Inf).

Author(s)

Christopher Paciorek

Examples

```
dmyexp <- nimbleFunction(
   run = function(x = double(0), rate = double(0), log = integer(0)) {
       returnType(double(0))
       logProb <- log(rate) - x*rate
       if(log) {
           return(logProb)
       } else {
           return(exp(logProb))
       }
   })
rmyexp <- nimbleFunction(
   run = function(n = integer(0), rate = double(0)) {
       returnType(double(0))
       if(n != 1) nimPrint("rmyexp only allows n = 1; using n = 1.")
       dev <- runif(1, 0, 1)
       return(-log(1-dev) / rate)
   }
   )
registerDistributions(list(
```

```
    dmyexp = list(
                BUGSdist = "dmyexp(rate, scale)",
                Rdist = "dmyexp(rate = 1/scale)",
                altParams = "scale = 1/rate",
                pqAvail = FALSE)))
code <- nimbleCode({
    y ~ dmyexp(rate = r)
    r ~ dunif(0, 100)
})
m <- nimbleModel(code, inits = list(r = 1), data = list(y = 2))
m$calculate('y')
m$r <- 2
m$calculate('y')
m$resetData()
m$simulate('y')
m$y

# alternatively, simply specify a character vector with the
# name of one or more 'd' functions
deregisterDistributions('dmyexp')
registerDistributions('dmyexp')

# or simply use in BUGS code without registration
deregisterDistributions('dmyexp')
m <- nimbleModel(code, inits = list(r = 1), data = list(y = 2))

# example of Dirichlet-multinomial registration to illustrate
# use of 'types' (note that registration is not actually needed
# in this case)
ddirchmulti <- nimbleFunction(
    run = function(x = double(1), alpha = double(1), size = double(0),
                    log = integer(0, default = 0)) {
        returnType(double(0))
        logProb <- lgamma(size) - sum(lgamma(x)) + lgamma(sum(alpha)) -
            sum(lgamma(alpha)) + sum(lgamma(alpha + x)) - lgamma(sum(alpha) +
                                                                    size)
        if(log) return(logProb)
        else return(exp(logProb))
    })

rdirchmulti <- nimbleFunction(
    run = function(n = integer(0), alpha = double(1), size = double(0)) {
        returnType(double(1))
        if(n != 1) print("rdirchmulti only allows n = 1; using n = 1.")
        p <- rdirch(1, alpha)
        return(rmulti(1, size = size, prob = p))
    })

registerDistributions(list(
    ddirchmulti = list(
        BUGSdist = "ddirchmulti(alpha, size)",
        types = c('value = double(1)', 'alpha = double(1)')
        )
```

```
      ))
```

---

resize   *Resizes a modelValues object*

Description

Adds or removes rows to a modelValues object. If rows are added to a modelValues object, the default values are NA. Works in both R and NIMBLE.

Usage

```
resize(container, k)
```

Arguments

container   modelValues object

k   number of rows that modelValues is set to

Details

See the User Manual or help(modelValuesBaseClass) for infomation about modelValues objects

Author(s)

Clifford Anderson-Bergman

Examples

```
mvConf <- modelValuesConf(vars = c('a', 'b'),
            types = c('double', 'double'),
            sizes = list(a = 1, b = c(2,2) ) )
mv <- modelValues(mvConf)
as.matrix(mv)
resize(mv, 3)
as.matrix(mv)
```

Rmatrix2mvOneVar *Set values of one variable of a modelValues object from an R matrix*

Description

Normally a modelValues object is accessed one "row" at a time. This function allows all rows for one variable to set from a matrix with one dimension more than the variable to be set.

Usage

```
Rmatrix2mvOneVar(mat, mv, varName, k)
```

Arguments

mat	Input matrix
mv	modelValues object to be modified.
varName	Character string giving the name of the variable on mv to be set
k	Number of rows to use

Details

This function may be deprecated in the future when a more natural system for interacting with modelValues objects is developed.

RmodelBaseClass-class *Class* RmodelBaseClass

Description

Classes used internally in NIMBLE and not expected to be called directly by users.

---

run.time *Time execution of NIMBLE code*

Description

Time execution of NIMBLE code

Usage

```
run.time(code)
```

Arguments

code   code to be timed

Details

Function for use in nimbleFunction run code; when nimbleFunctions are run in R, this simply wraps
`system.time`.

Author(s)

NIMBLE Development Team

---

runCrossValidate   *Perform k-fold cross-validation on a NIMBLE model fit by MCMC*

Description

Takes a NIMBLE model MCMC configuration and conducts k-fold cross-validation of the MCMC
fit, returning a measure of the model's predictive performance.

Usage

```
runCrossValidate(
  MCMCconfiguration,
  k,
  foldFunction = "random",
  lossFunction = "MSE",
  MCMCcontrol = list(),
  returnSamples = FALSE,
  nCores = 1,
  nBootReps = 200,
  silent = FALSE
)
```

Arguments

MCMCconfiguration

                a NIMBLE MCMC configuration object, returned by a call to `configureMCMC`.

k   number of folds that should be used for cross-validation.

foldFunction   one of (1) an R function taking a single integer argument i, and returning a
                character vector with the names of the data nodes to leave out of the model for
                fold i, or (2) the character string "random", indicating that data nodes will be
                randomly partitioned into k folds. Note that choosing "random" and setting k
                equal to the total number of data nodes in the model will perform leave-one-out
                cross-validation. Defaults to "random". See 'Details'.

lossFunction   one of (1) an R function taking a set of simulated data and a set of observed data, and calculating the loss from those, or (2) a character string naming one of NIMBLE's built-in loss functions. If a character string, must be one of "predictive" to use the log predictive density as a loss function or "MSE" to use the mean squared error as a loss function. Defaults to "MSE". See 'Details' for information on creating a user-defined loss function.

MCMCcontrol   (optional) an R list with parameters governing the MCMC algorithm, See 'Details' for specific parameters.

returnSamples   logical indicating whether to return all posterior samples from all MCMC runs. This can result in a very large returned object (there will be k sets of posterior samples returned). Defaults to FALSE.

nCores   number of cpu cores to use in parallelizing the CV calculation. Only MacOS and Linux operating systems support multiple cores at this time. Defaults to 1.

nBootReps   number of bootstrap samples to use when estimating the Monte Carlo error of the cross-validation metric. Defaults to 200. If no Monte Carlo error estimate is desired, nBootReps can be set to NA, which can potentially save significant computation time.

silent   Boolean specifying whether to show output from the algorithm as it's running (default = FALSE).

Details

k-fold CV in NIMBLE proceeds by separating the data in a nimbleModel into k folds, as determined by the foldFunction argument. For each fold, the corresponding data are held out of the model, and MCMC is run to estimate the posterior distribution and simultaneously impute posterior predictive values for the held-out data. Then, the posterior predictive values are compared to the known, held-out data values via the specified lossFunction. The loss values are averaged over the posterior samples for each fold, and these averaged values for each fold are then averaged over all folds to produce a single out-of-sample loss estimate. Additionally, estimates of the Monte Carlo error for each fold are returned.

Value

an R list with four elements:

- CVvalue The CV value, measuring the model's ability to predict new data. Smaller relative values indicate better model performance.

- CVstandardError The standard error of the CV value, giving an indication of the total Monte Carlo error in the CV estimate.

- foldCVInfo An list of fold CV values and standard errors for each fold.

- samples An R list, only returned when returnSamples = TRUE. The i'th element of this list will be a matrix of posterior samples from the model with the i'th fold of data left out. There will be k sets of samples.

The `foldFunction` **Argument**

If the default `'random'` method is not used, the `foldFunction` argument must be an R function that takes a single integer-valued argument i. i is guaranteed to be within the range $[1, k]$. For each integer value i, the function should return a character vector of node names corresponding to the data nodes that will be left out of the model for that fold. The returned node names can be expanded, but don't need to be. For example, if fold i is inteded to leave out the model nodes x[1], x[2] and x[3] then the function could return either c('x[1]','x[2]','x[3]') or 'x[1:3]'.

The `lossFunction` **Argument**

If you don't wish to use NIMBLE's built-in "MSE" or "predictive" loss functions, you may provide your own R function as the `lossFunction` argument to runCrossValidate. A user-supplied `lossFunction` must be an R function that takes two arguments: the first, named simulatedDataValues, will be a vector of simulated data values. The second, named actualDataValues, will be a vector of observed data values corresponding to the simulated data values in simulatedDataValues. The loss function should return a single scalar number. See 'Examples' for an example of a user-defined loss function.

The `MCMCcontrol` **Argument**

The `MCMCcontrol` argument is a list with the following elements:

- `niter` an integer argument determining how many MCMC iterations should be run for each loss value calculation. Defaults to 10000, but should probably be manually set.
- `nburnin` the number of samples from the start of the MCMC chain to discard as burn-in for each loss value calculation. Must be between 0 and `niter`. Defaults to 10

Author(s)

Nicholas Michaud and Lauren Ponisio

Examples

```
## Not run:

## We conduct CV on the classic "dyes" BUGS model.

dyesCode <- nimbleCode({
  for (i in 1:BATCHES) {
    for (j in 1:SAMPLES) {
      y[i,j] ~ dnorm(mu[i], tau.within);
    }
    mu[i] ~ dnorm(theta, tau.between);
  }

  theta ~ dnorm(0.0, 1.0E-10);
  tau.within ~ dgamma(0.001, 0.001);  sigma2.within <- 1/tau.within;
  tau.between ~ dgamma(0.001, 0.001);  sigma2.between <- 1/tau.between;
})
```

```
dyesData <- list(y = matrix(c(1545, 1540, 1595, 1445, 1595,
                              1520, 1440, 1555, 1550, 1440,
                              1630, 1455, 1440, 1490, 1605,
                              1595, 1515, 1450, 1520, 1560,
                              1510, 1465, 1635, 1480, 1580,
                              1495, 1560, 1545, 1625, 1445),
                            nrow = 6, ncol = 5))

dyesConsts <- list(BATCHES = 6,
                   SAMPLES = 5)

dyesInits <- list(theta = 1500, tau.within = 1, tau.between =  1)

dyesModel <- nimbleModel(code = dyesCode,
                         constants = dyesConsts,
                         data = dyesData,
                         inits = dyesInits)

# Define the fold function.
# This function defines the data to leave out for the i'th fold
# as the i'th row of the data matrix y.  This implies we will have
# 6 folds.

dyesFoldFunction <- function(i){
  foldNodes_i <- paste0('y[', i, ', ]')  # will return 'y[1,]' for i = 1 e.g.
  return(foldNodes_i)
}

# We define our own loss function as well.
# The function below will compute the root mean squared error.

RMSElossFunction <- function(simulatedDataValues, actualDataValues){
  dataLength <- length(simulatedDataValues) # simulatedDataValues is a vector
  SSE <- 0
  for(i in 1:dataLength){
    SSE <- SSE + (simulatedDataValues[i] - actualDataValues[i])^2
  }
  MSE <- SSE / dataLength
  RMSE <- sqrt(MSE)
  return(RMSE)
}

dyesMCMCconfiguration <- configureMCMC(dyesModel)

crossValOutput <- runCrossValidate(MCMCconfiguration = dyesMCMCconfiguration,
                                   k = 6,
                                   foldFunction = dyesFoldFunction,
                                   lossFunction = RMSElossFunction,
                                   MCMCcontrol = list(niter = 5000,
                                                      nburnin = 500))

## End(Not run)
```

---

runMCMC	*Run one or more chains of an MCMC algorithm and return samples, summary and/or WAIC*

Description

Takes as input an MCMC algorithm (ideally a compiled one for speed) and runs the MCMC with one or more chains, any returns any combination of posterior samples, posterior summary statistics, and a WAIC value.

Usage

```
runMCMC(
  mcmc,
  niter = 10000,
  nburnin = 0,
  thin,
  thin2,
  nchains = 1,
  inits,
  setSeed = FALSE,
  progressBar = getNimbleOption("MCMCprogressBar"),
  samples = TRUE,
  samplesAsCodaMCMC = FALSE,
  summary = FALSE,
  WAIC = FALSE,
  perChainWAIC = FALSE
)
```

Arguments

mcmc	A NIMBLE MCMC algorithm. See details.
niter	Number of iterations to run each MCMC chain. Default value is 10000.
nburnin	Number of initial, pre-thinning, MCMC iterations to discard. Default value is 0.
thin	Thinning interval for collecting MCMC samples, corresponding to monitors. Thinning occurs after the initial nburnin samples are discarded. Default value is 1.
thin2	Thinning interval for collecting MCMC samples, corresponding to the second, optional set of monitors2. Thinning occurs after the initial nburnin samples are discarded. Default value is 1.
nchains	Number of MCMC chains to run. Default value is 1.
inits	Optional argument to specify initial values for each chain. See details.
setSeed	Logical or numeric argument. If a single numeric value is provided, R's random number seed will be set to this value at the onset of each MCMC chain. If a numeric vector of length nchains is provided, then each element of this vector

is provided as R's random number seed at the onset of the corresponding MCMC chain. Otherwise, in the case of a logical value, if TRUE, then R's random number seed for the ith chain is set to be i, at the onset of each MCMC chain. Note that specifying the argument setSeed = 0 does not prevent setting the RNG seed, but rather sets the random number generation seed to 0 at the beginning of each MCMC chain. Default value is FALSE.

progressBar   Logical argument. If TRUE, an MCMC progress bar is displayed during execution of each MCMC chain. Default value is defined by the nimble package option MCMCprogressBar.

samples   Logical argument. If TRUE, then posterior samples are returned from each MCMC chain. These samples are optionally returned as coda mcmc objects, depending on the samplesAsCodaMCMC argument. Default value is TRUE. See details.

samplesAsCodaMCMC
    Logical argument. If TRUE, then a coda mcmc object is returned instead of an R matrix of samples, or when nchains > 1 a coda mcmc.list object is returned containing nchains mcmc objects. This argument is only used when samples is TRUE. Default value is FALSE. See details.

summary   Logical argument. When TRUE, summary statistics for the posterior samples of each parameter are also returned, for each MCMC chain. This may be returned in addition to the posterior samples themselves. Default value is FALSE. See details.

WAIC   Logical argument. When TRUE, the WAIC (Watanabe, 2010) of the model is calculated and returned. Note that in order for the WAIC to be calculated, the mcmc object must have also been created with the argument 'enableWAIC = TRUE'. If multiple chains are run, then a single WAIC value is calculated using the posterior samples from all chains. Default value is FALSE. See help(waic).

perChainWAIC   Logical argument. When TRUE and multiple chains are run, the WAIC for each chain is returned as a means of helping assess the stability of the WAIC estimate. Default value is FALSE, corresponding to a single WAIC estimate based on all of the chains.

Details

At least one of samples, summary or WAIC must be TRUE, since otherwise, nothing will be returned. Any combination of these may be TRUE, including possibly all three, in which case posterior samples and summary statistics are returned for each MCMC chain, and an overall WAIC value is calculated and returned.

When samples = TRUE, the form of the posterior samples is determined by the samplesAsCodaMCMC argument, as either matrices of posterior samples, or coda mcmc and mcmc.list objects.

Posterior summary statistics are returned individually for each chain, and also as calculated from all chains combined (when nchains > 1).

If provided, the inits argument can be one of three things:

(1) a function to generate initial values, which will be executed to generate initial values at the beginning of each MCMC chain, or (2) a single named list of initial values which, will be used for each chain, or (3) a list of length nchains, each element being a named list of initial values which be used for one MCMC chain.

The `inits` argument may also be omitted, in which case the current values in the `model` object will be used as the initial values of the first chain, and subsequent chains will begin using starting values where the previous chain ended.

Other aspects of the MCMC algorithm, such as the specific sampler assignments, must be specified in advance using the MCMC configuration object (created using `configureMCMC`), which is then used to build an MCMC algorithm (using `buildMCMC`) argument.

The `niter` argument specifies the number of pre-thinning MCMC iterations, and the `nburnin` argument specifies the number of pre-thinning MCMC samples to discard. After discarding these burn-in samples, thinning of the remaining samples will take place. The total number of posterior samples returned will be floor((niter-nburnin)/thin).

The MCMC option `mcmc$run(...,reset = FALSE)`, used to continue execution of an MCMC chain, is not available through runMCMC().

Value

A list is returned with named elements depending on the arguments passed to `nimbleMCMC`, unless this list contains only a single element, in which case only that element is returned. These elements may include `samples`, `summary`, and `WAIC`, and when the MCMC is monitoring a second set of nodes using `monitors2`, also `samples2`. When `nchains = 1`, posterior samples are returned as a single matrix, and summary statistics as a single matrix. When `nchains > 1`, posterior samples are returned as a list of matrices, one matrix for each chain, and summary statistics are returned as a list containing `nchains+1` matrices: one matrix corresponding to each chain, and the final element providing a summary of all chains, combined. If `samplesAsCodaMCMC` is `TRUE`, then posterior samples are provided as coda `mcmc` and `mcmc.list` objects. When `WAIC` is `TRUE`, a WAIC summary object is returned.

Author(s)

Daniel Turek

See Also

[configureMCMC](#) [buildMCMC](#) [nimbleMCMC](#)

Examples

```
## Not run:
code <- nimbleCode({
    mu ~ dnorm(0, sd = 1000)
    sigma ~ dunif(0, 1000)
    for(i in 1:10) {
        x[i] ~ dnorm(mu, sd = sigma)
    }
})
Rmodel <- nimbleModel(code)
Rmodel$setData(list(x = c(2, 5, 3, 4, 1, 0, 1, 3, 5, 3)))
Rmcmc <- buildMCMC(Rmodel)
Cmodel <- compileNimble(Rmodel)
Cmcmc <- compileNimble(Rmcmc, project = Rmodel)
inits <- function() list(mu = rnorm(1,0,1), sigma = runif(1,0,10))
```

```
samplesList <- runMCMC(Cmcmc, niter = 10000, nchains = 3, inits = inits)

## End(Not run)
```

---

sampler_BASE   *MCMC Sampling Algorithms*

Description

Details of the MCMC sampling algorithms provided with the NIMBLE MCMC engine

Usage

```
sampler_BASE()

sampler_posterior_predictive(model, mvSaved, target, control)

sampler_posterior_predictive_branch(model, mvSaved, target, control)

sampler_binary(model, mvSaved, target, control)

sampler_categorical(model, mvSaved, target, control)

sampler_RW(model, mvSaved, target, control)

sampler_RW_block(model, mvSaved, target, control)

sampler_RW_llFunction(model, mvSaved, target, control)

sampler_slice(model, mvSaved, target, control)

sampler_ess(model, mvSaved, target, control)

sampler_AF_slice(model, mvSaved, target, control)

sampler_crossLevel(model, mvSaved, target, control)

sampler_RW_llFunction_block(model, mvSaved, target, control)

sampler_RW_multinomial(model, mvSaved, target, control)

sampler_RW_dirichlet(model, mvSaved, target, control)

sampler_RW_wishart(model, mvSaved, target, control)

sampler_RW_lkj_corr_cholesky(model, mvSaved, target, control)
```

```
sampler_RW_block_lkj_corr_cholesky(model, mvSaved, target, control)

sampler_CAR_normal(model, mvSaved, target, control)

sampler_CAR_proper(model, mvSaved, target, control)

sampler_RJ_fixed_prior(model, mvSaved, target, control)

sampler_RJ_indicator(model, mvSaved, target, control)

sampler_RJ_toggled(model, mvSaved, target, control)

sampler_CRP_concentration(model, mvSaved, target, control)

sampler_CRP(model, mvSaved, target, control)
```

Arguments

`model`	(uncompiled) model on which the MCMC is to be run
`mvSaved`	`modelValues` object to be used to store MCMC samples
`target`	node(s) on which the sampler will be used
`control`	named list that controls the precise behavior of the sampler, with elements specific to `samplertype`. The default values for control list are specified in the setup code of each sampling algorithm. Descriptions of each sampling algorithm, and the possible customizations for each sampler (using the `control` argument) appear below.

`sampler_base`

base class for new samplers

When you write a new sampler for use in a NIMBLE MCMC (see User Manual), you must include `contains = sampler_BASE`.

binary sampler

The binary sampler performs Gibbs sampling for binary-valued (discrete 0/1) nodes. This can only be used for nodes following either a `dbern(p)` or `dbinom(p,size=1)` distribution.

The binary sampler accepts no control list arguments.

categorical sampler

The categorical sampler performs Gibbs sampling for a single node, which must follow a categorical (`dcat`) distribution.

The categorical sampler accepts no control list arguments.

RW sampler

The RW sampler executes adaptive Metropolis-Hastings sampling with a normal proposal distribution (Metropolis, 1953), implementing the adaptation routine given in Shaby and Wells, 2011. This sampler can be applied to any scalar continuous-valued stochastic node, and can optionally sample on a log scale.

The RW sampler accepts the following control list elements:

- log. A logical argument, specifying whether the sampler should operate on the log scale. (default = FALSE)

- reflective. A logical argument, specifying whether the normal proposal distribution should reflect to stay within the range of the target distribution. (default = FALSE)

- adaptive. A logical argument, specifying whether the sampler should adapt the scale (proposal standard deviation) throughout the course of MCMC execution to achieve a theoretically desirable acceptance rate. (default = TRUE)

- adaptInterval. The interval on which to perform adaptation. Every adaptInterval MCMC iterations (prior to thinning), the RW sampler will perform its adaptation procedure. This updates the scale variable, based upon the sampler's achieved acceptance rate over the past adaptInterval iterations. (default = 200)

- adaptFactorExponent. Exponent controling the rate of decay of the scale adaptation factor. See Shaby and Wells, 2011, for details. (default = 0.8)

- scale. The initial value of the normal proposal standard deviation. If adaptive = FALSE, scale will never change. (default = 1)

The RW sampler cannot be used with options log=TRUE and reflective=TRUE, i.e. it cannot do reflective sampling on a log scale.

RW_block sampler

The RW_block sampler performs a simultaneous update of one or more model nodes, using an adaptive Metropolis-Hastings algorithm with a multivariate normal proposal distribution (Roberts and Sahu, 1997), implementing the adaptation routine given in Shaby and Wells, 2011. This sampler may be applied to any set of continuous-valued model nodes, to any single continuous-valued multivariate model node, or to any combination thereof.

The RW_block sampler accepts the following control list elements:

- adaptive. A logical argument, specifying whether the sampler should adapt the scale (a coefficient for the entire proposal covariance matrix) and propCov (the multivariate normal proposal covariance matrix) throughout the course of MCMC execution. If only the scale should undergo adaptation, this argument should be specified as TRUE. (default = TRUE)

- adaptScaleOnly. A logical argument, specifying whether adaption should be done only for scale (TRUE) or also for provCov (FALSE). This argument is only relevant when adaptive = TRUE. When adaptScaleOnly = FALSE, both scale and propCov undergo adaptation; the sampler tunes the scaling to achieve a theoretically good acceptance rate, and the proposal covariance to mimic that of the empirical samples. When adaptScaleOnly = TRUE, only the proposal scale is adapted. (default = FALSE)

- adaptInterval. The interval on which to perform adaptation. Every adaptInterval MCMC iterations (prior to thinning), the RW_block sampler will perform its adaptation procedure, based on the past adaptInterval iterations. (default = 200)

- adaptFactorExponent. Exponent controling the rate of decay of the scale adaptation factor. See Shaby and Wells, 2011, for details. (default = 0.8)

- scale. The initial value of the scalar multiplier for propCov. If adaptive = FALSE, scale will never change. (default = 1)

- propCov. The initial covariance matrix for the multivariate normal proposal distribution. This element may be equal to the character string 'identity', in which case the identity matrix of the appropriate dimension will be used for the initial proposal covariance matrix. (default = 'identity')

Note that modifying elements of the control list may greatly affect the performance of this sampler. In particular, the sampler can take a long time to find a good proposal covariance when the elements being sampled are not on the same scale. We recommend providing an informed value for propCov in this case (possibly simply a diagonal matrix that approximates the relative scales), as well as possibly providing a value of scale that errs on the side of being too small. You may also consider decreasing adaptFactorExponent and/or adaptInterval, as doing so has greatly improved performance in some cases.

RW_llFunction sampler

Sometimes it is useful to control the log likelihood calculations used for an MCMC updater instead of simply using the model. For example, one could use a sampler with a log likelihood that analytically (or numerically) integrates over latent model nodes. Or one could use a sampler with a log likelihood that comes from a stochastic approximation such as a particle filter, allowing composition of a particle MCMC (PMCMC) algorithm (Andrieu et al., 2010). The RW_llFunction sampler handles this by using a Metropolis-Hastings algorithm with a normal proposal distribution and a user-provided log-likelihood function. To allow compiled execution, the log-likelihood function must be provided as a specialized instance of a nimbleFunction. The log-likelihood function may use the same model as the MCMC as a setup argument, but if so the state of the model should be unchanged during execution of the function (or you must understand the implications otherwise).

The RW_llFunction sampler accepts the following control list elements:

- adaptive. A logical argument, specifying whether the sampler should adapt the scale (proposal standard deviation) throughout the course of MCMC execution. (default = TRUE)

- adaptInterval. The interval on which to perform adaptation. (default = 200)

- scale. The initial value of the normal proposal standard deviation. (default = 1)

- llFunction. A specialized nimbleFunction that accepts no arguments and returns a scalar double number. The return value must be the total log-likelihood of all stochastic dependents of the target nodes – and, if includesTarget = TRUE, of the target node(s) themselves – or whatever surrogate is being used for the total log-likelihood. This is a required element with no default.

- includesTarget. Logical variable indicating whether the return value of llFunction includes the log-likelihood associated with target. This is a required element with no default.

slice sampler

The slice sampler performs slice sampling of the scalar node to which it is applied (Neal, 2003). This sampler can operate on either continuous-valued or discrete-valued scalar nodes. The slice sampler performs a 'stepping out' procedure, in which the slice is iteratively expanded to the left or right by an amount sliceWidth. This sampler is optionally adaptive, governed by a control list element, whereby the value of sliceWidth is adapted towards the observed absolute difference between successive samples.

The slice sampler accepts the following control list elements:

- adaptive. A logical argument, specifying whether the sampler will adapt the value of sliceWidth throughout the course of MCMC execution. (default = TRUE)

- adaptInterval. The interval on which to perform adaptation. (default = 200)

- sliceWidth. The initial value of the width of each slice, and also the width of the expansion during the iterative 'stepping out' procedure. (default = 1)

- sliceMaxSteps. The maximum number of expansions which may occur during the 'stepping out' procedure. (default = 100)

- maxContractions. The maximum number of contractions of the interval that may occur during sampling (this prevents infinite looping in unusual situations). (default = 100)

- maxContractionsWarning. A logical argument specifying whether to warn when the maximum number of contractions is reached. (default = TRUE)

ess sampler

The ess sampler performs elliptical slice sampling of a single node, which must follow either a univariate or multivariate normal distribution (Murray, 2010). The algorithm is an extension of slice sampling (Neal, 2003), generalized to context of the Gaussian distribution. An auxiliary variable is used to identify points on an ellipse (which passes through the current node value) as candidate samples, which are accepted contingent upon a likelihood evaluation at that point. This algorithm requires no tuning parameters and therefore no period of adaptation, and may result in very efficient sampling from Gaussian distributions.

The ess sampler accepts the following control list arguments.

- maxContractions. The maximum number of contractions of the interval that may occur during sampling (this prevents infinite looping in unusual situations). (default = 100)

- maxContractionsWarning. A logical argument specifying whether to warn when the maximum number of contractions is reached. (default = TRUE)

AF_slice sampler

The automated factor slice sampler conducts a slice sampling algorithm on one or more model nodes. The sampler uses the eigenvectors of the posterior covariance between these nodes as an orthogonal basis on which to perform its 'stepping Out' procedure. The sampler is adaptive in updating both the width of the slices and the values of the eigenvectors. The sampler can be applied to any set of continuous or discrete-valued model nodes, to any single continuous or discrete-valued multivariate model node, or to any combination thereof. The automated factor slice sampler accepts the following control list elements:

- sliceWidths. A numeric vector of initial slice widths. The length of the vector must be equal to the sum of the lengths of all nodes being used by the automated factor slice sampler. Defaults to a vector of 1's.
- sliceAdaptFactorMaxIter. The number of iterations for which the factors (eigenvectors) will continue to adapt to the posterior correlation. (default = 15000)
- sliceAdaptFactorInterval. The interval on which to perform factor adaptation. (default = 1000)
- sliceAdaptWidthMaxIter. The maximum number of iterations for which to adapt the widths for a given set of factors. (default = 512)
- sliceAdaptWidthTolerance. The tolerance for when widths no longer need to adapt, between 0 and 0.5. (default = 0.1)
- sliceMaxSteps. The maximum number of expansions which may occur during the 'stepping out' procedure. (default = 100)
- maxContractions. The maximum number of contractions of the interval that may occur during sampling (this prevents infinite looping in unusual situations). (default = 100)
- maxContractionsWarning. A logical argument specifying whether to warn when the maximum number of contractions is reached. (default = TRUE)

crossLevel sampler

This sampler is constructed to perform simultaneous updates across two levels of stochastic dependence in the model structure. This is possible when all stochastic descendents of node(s) at one level have conjugate relationships with their own stochastic descendents. In this situation, a Metropolis-Hastings algorithm may be used, in which a multivariate normal proposal distribution is used for the higher-level nodes, and the corresponding proposals for the lower-level nodes undergo Gibbs (conjugate) sampling. The joint proposal is either accepted or rejected for all nodes involved based upon the Metropolis-Hastings ratio. This sampler is a conjugate version of Scheme 3 in Knorr-Held and Rue (2002). It can also be seen as a Metropolis-based version of collapsed Gibbs sampling (in particular Sampler 3 of van Dyk and Park (2008)).

The requirement that all stochastic descendents of the target nodes must themselves have only conjugate descendents will be checked when the MCMC algorithm is built. This sampler is useful when there is strong dependence across the levels of a model that causes problems with convergence or mixing.

The crossLevel sampler accepts the following control list elements:

- adaptive. Logical argument, specifying whether the multivariate normal proposal distribution for the target nodes should be adaptived. (default = TRUE)
- adaptInterval. The interval on which to perform adaptation. (default = 200)
- scale. The initial value of the scalar multiplier for propCov. (default = 1)
- propCov. The initial covariance matrix for the multivariate normal proposal distribution. This element may be equal to the character string 'identity' or any positive definite matrix of the appropriate dimensions. (default = 'identity')

RW_llFunction_block sampler

Sometimes it is useful to control the log likelihood calculations used for an MCMC updater instead of simply using the model. For example, one could use a sampler with a log likelihood that analytically (or numerically) integrates over latent model nodes. Or one could use a sampler with a log

likelihood that comes from a stochastic approximation such as a particle filter, allowing composition of a particle MCMC (PMCMC) algorithm (Andrieu et al., 2010) (but see samplers listed below for NIMBLE's direct implementation of PMCMC). The `RW_llFunction_block` sampler handles this by using a Metropolis-Hastings algorithm with a multivariate normal proposal distribution and a user-provided log-likelihood function. To allow compiled execution, the log-likelihood function must be provided as a specialized instance of a nimbleFunction. The log-likelihood function may use the same model as the MCMC as a setup argument, but if so the state of the model should be unchanged during execution of the function (or you must understand the implications otherwise).

The RW_llFunction_block sampler accepts the following control list elements:

- adaptive. A logical argument, specifying whether the sampler should adapt the proposal covariance throughout the course of MCMC execution. (default is TRUE)

- adaptScaleOnly. A logical argument, specifying whether adaption should be done only for scale (TRUE) or also for provCov (FALSE). This argument is only relevant when adaptive = TRUE. When adaptScaleOnly = FALSE, both scale and propCov undergo adaptation; the sampler tunes the scaling to achieve a theoretically good acceptance rate, and the proposal covariance to mimic that of the empirical samples. When adaptScaleOnly = TRUE, only the proposal scale is adapted. (default = FALSE)

- adaptInterval. The interval on which to perform adaptation. (default = 200)

- adaptFactorExponent. Exponent controling the rate of decay of the scale adaptation factor. See Shaby and Wells, 2011, for details. (default = 0.8)

- scale. The initial value of the scalar multiplier for propCov. If adaptive = FALSE, scale will never change. (default = 1)

- propCov. The initial covariance matrix for the multivariate normal proposal distribution. This element may be equal to the character string 'identity', in which case the identity matrix of the appropriate dimension will be used for the initial proposal covariance matrix. (default = 'identity')

- llFunction. A specialized nimbleFunction that accepts no arguments and returns a scalar double number. The return value must be the total log-likelihood of all stochastic dependents of the target nodes – and, if includesTarget = TRUE, of the target node(s) themselves – or whatever surrogate is being used for the total log-likelihood. This is a required element with no default.

- includesTarget. Logical variable indicating whether the return value of llFunction includes the log-likelihood associated with target. This is a required element with no default.

RW_multinomial sampler

This sampler is designed for sampling multinomial target distributions. The sampler performs a series of Metropolis-Hastings steps between pairs of groups. Proposals are generated via a draw from a binomial distribution, whereafter the proposed number density is moved from one group to another group. The acceptance or rejection of these proposals follows a standard Metropolis-Hastings procedure. Probabilities for the random binomial proposals are adapted to a target acceptance rate of 0.5.

The `RW_multinomial` sampler accepts the following control list elements:

- adaptive. A logical argument, specifying whether the sampler should adapt the binomial proposal probabilities throughout the course of MCMC execution. (default = TRUE)

- adaptInterval. The interval on which to perform adaptation. A minimum value of 100 is required. (default = 200)

RW_dirichlet sampler

This sampler is designed for sampling non-conjugate Dirichlet distributions. The sampler performs a series of Metropolis-Hastings updates (on the log scale) to each component of a gamma-reparameterization of the target Dirichlet distribution. The acceptance or rejection of these proposals follows a standard Metropolis-Hastings procedure.

The RW_dirichlet sampler accepts the following control list elements:

- adaptive. A logical argument, specifying whether the sampler should independently adapt the scale (proposal standard deviation, on the log scale) for each componentwise Metropolis-Hasting update, to achieve a theoretically desirable acceptance rate for each. (default = TRUE)
- adaptInterval. The interval on which to perform adaptation. Every adaptInterval MCMC iterations (prior to thinning), the sampler will perform its adaptation procedure. (default = 200)
- adaptFactorExponent. Exponent controling the rate of decay of the scale adaptation factor. See Shaby and Wells, 2011, for details. (default = 0.8)
- scale. The initial value of the proposal standard deviation (on the log scale) for each component of the reparameterized Dirichlet distribution. If adaptive = FALSE, the proposal standard deviations will never change. (default = 1)

RW_wishart sampler

This sampler is designed for sampling non-conjugate Wishart and inverse-Wishart distributions. More generally, it can update any symmetric positive-definite matrix (for example, scaled covaraiance or precision matrices). The sampler performs block Metropolis-Hastings updates following a transformation to an unconstrained scale (Cholesky factorization of the original matrix, then taking the log of the main diagonal elements.

The RW_wishart sampler accepts the following control list elements:

- adaptive. A logical argument, specifying whether the sampler should adapt the scale and proposal covariance for the multivariate normal Metropolis-Hasting proposals, to achieve a theoretically desirable acceptance rate for each. (default = TRUE)
- adaptInterval. The interval on which to perform adaptation. Every adaptInterval MCMC iterations (prior to thinning), the sampler will perform its adaptation procedure. (default = 200)
- adaptFactorExponent. Exponent controling the rate of decay of the scale adaptation factor. See Shaby and Wells, 2011, for details. (default = 0.8)
- scale. The initial value of the scalar multiplier for the multivariate normal Metropolis-Hastings proposal covariance. If adaptive = FALSE, scale will never change. (default = 1)

RW_block_lkj_corr_cholesky sampler

This sampler is designed for sampling non-conjugate LKJ correlation Cholesky factor distributions. The sampler performs a blocked Metropolis-Hastings update following a transformation to an unconstrained scale (using the signed stickbreaking approach documented in Section 10.12 of the Stan Language Reference Manual, version 2.27).

The `RW_block_lkj_corr_cholesky` sampler accepts the following control list elements:

- adaptive. A logical argument, specifying whether the sampler should adapt the scale (a coefficient for the entire proposal covariance matrix) and propCov (the multivariate normal proposal covariance matrix) throughout the course of MCMC execution. If only the scale should undergo adaptation, this argument should be specified as TRUE. (default = TRUE)

- adaptScaleOnly. A logical argument, specifying whether adaption should be done only for scale (TRUE) or also for provCov (FALSE). This argument is only relevant when adaptive = TRUE. When adaptScaleOnly = FALSE, both scale and propCov undergo adaptation; the sampler tunes the scaling to achieve a theoretically good acceptance rate, and the proposal covariance to mimic that of the empirical samples. When adaptScaleOnly = TRUE, only the proposal scale is adapted. (default = FALSE)

- adaptInterval. The interval on which to perform adaptation. Every adaptInterval MCMC iterations (prior to thinning), the RW_block sampler will perform its adaptation procedure, based on the past adaptInterval iterations. (default = 200)

- adaptFactorExponent. Exponent controlling the rate of decay of the scale adaptation factor. See Shaby and Wells, 2011, for details. (default = 0.8)

- scale. The initial value of the scalar multiplier for propCov. If adaptive = FALSE, scale will never change. (default = 1)

- propCov. The initial covariance matrix for the multivariate normal proposal distribution. This element may be equal to the character string 'identity', in which case the identity matrix of the appropriate dimension will be used for the initial proposal covariance matrix. (default = 'identity')

This is the default sampler for the LKJ distribution. However, blocked samplers may perform poorly if the adaptation configuration is poorly chosen. See the comments in the RW_block section of this documentation.

RW_lkj_corr_cholesky sampler

This sampler is designed for sampling non-conjugate LKJ correlation Cholesky factor distributions. The sampler performs individual Metropolis-Hastings updates following a transformation to an unconstrained scale (using the signed stickbreaking approach documented in Section 10.12 of the Stan Language Reference Manual, version 2.27).

The `RW_lkj_corr_cholesky` sampler accepts the following control list elements:

- adaptive. A logical argument, specifying whether the sampler should adapt the scales of the univariate normal Metropolis-Hasting proposals, to achieve a theoretically desirable acceptance rate for each. (default = TRUE)

- adaptInterval. The interval on which to perform adaptation. Every adaptInterval MCMC iterations (prior to thinning), the sampler will perform its adaptation procedure. (default = 200)

- adaptFactorExponent. Exponent controling the rate of decay of the scale adaptation factor. See Shaby and Wells, 2011, for details. (default = 0.8)

- scale. The initial value of the scalar multiplier for the multivariate normal Metropolis-Hastings proposal covariance. If adaptive = FALSE, scale will never change. (default = 1)

Note that this sampler is likely run much more slowly than the blocked sampler for the LKJ distribution, as updating each single element will generally incur the full cost of updating all dependencies of the entire matrix.

CAR_normal sampler

The CAR_normal sampler operates uniquely on improper (intrinsic) Gaussian conditional autoregressive (CAR) nodes, those with a `dcar_normal` prior distribution. It internally assigns one of three univariate samplers to each dimension of the target node: a posterior predictive, conjugate, or RW sampler; however these component samplers are specialized to operate on dimensions of a `dcar_normal` distribution.

The CAR_normal sampler accepts the following control list elements:

- `carUseConjugacy`. A logical argument, specifying whether to assign conjugate samplers for conjugate components of the target node. If `FALSE`, a RW sampler would be assigned instead. (default = TRUE)

- `adaptive`. A logical argument, specifying whether any component RW samplers should adapt the scale (proposal standard deviation), to achieve a theoretically desirable acceptance rate. (default = TRUE)

- `adaptInterval`. The interval on which to perform adaptation for any component RW samplers. Every `adaptInterval` MCMC iterations (prior to thinning), component RW samplers will perform an adaptation procedure. This updates the `scale` variable, based upon the sampler's achieved acceptance rate over the past `adaptInterval` iterations. (default = 200)

- `scale`. The initial value of the normal proposal standard deviation for any component RW samplers. If adaptive = FALSE, scale will never change. (default = 1)

CAR_proper sampler

The CAR_proper sampler operates uniquely on proper Gaussian conditional autoregressive (CAR) nodes, those with a `dcar_proper` prior distribution. It internally assigns one of three univariate samplers to each dimension of the target node: a posterior predictive, conjugate, or RW sampler, however these component samplers are specialized to operate on dimensions of a `dcar_proper` distribution.

The CAR_proper sampler accepts the following control list elements:

- `carUseConjugacy`. A logical argument, specifying whether to assign conjugate samplers for conjugate components of the target node. If `FALSE`, a RW sampler would be assigned instead. (default = TRUE)

- `adaptive`. A logical argument, specifying whether any component RW samplers should adapt the scale (proposal standard deviation), to achieve a theoretically desirable acceptance rate. (default = TRUE)

- `adaptInterval`. The interval on which to perform adaptation for any component RW samplers. Every adaptInterval MCMC iterations (prior to thinning), component RW samplers will perform an adaptation procedure. This updates the scale variable, based upon the sampler's achieved acceptance rate over the past adaptInterval iterations. (default = 200)

- `scale`. The initial value of the normal proposal standard deviation for any component RW samplers. If adaptive = FALSE, scale will never change. (default = 1)

CRP sampler

The CRP sampler is designed for fitting models involving Dirichlet process mixtures. It is exclusively assigned by NIMBLE's default MCMC configuration to nodes having the Chinese Restaurant Process distribution, dCRP. It executes sequential sampling of each component of the node (i.e., the cluster membership of each element being clustered). Internally, either of two samplers can be assigned, depending on conjugate or non-conjugate structures within the model. For conjugate and non-conjugate model structures, updates are based on Algorithm 2 and Algorithm 8 in Neal (2000), respectively.

- checkConjugacy. A logical argument, specifying whether to assign conjugate samplers if valid. (default = TRUE)

- printTruncation. A logical argument, specifying whether to print a warning when the MCMC attempts to use more clusters than the maximum number specified in the model. Only relevant where the user has specified the maximum number of clusters to be less than the number of observations. (default = TRUE)

CRP_concentration sampler

The CRP_concentration sampler is designed for Bayesian nonparametric mixture modeling. It is exclusively assigned to the concentration parameter of the Dirichlet process when the model is specified using the Chinese Restaurant Process distribution, dCRP. This sampler is assigned by default by NIMBLE's default MCMC configuration and can only be used when the prior for the concentration parameter is a gamma distribution. The assigned sampler is an augmented beta-gamma sampler as discussed in Section 6 in Escobar and West (1995).

posterior_predictive and posterior_predictive_branch samplers

The posterior_predictive sampler is only appropriate for use on terminal stochastic nodes. Note that such nodes play no role in inference but have often been included in BUGS models to accomplish posterior predictive checks. NIMBLE allows posterior predictive values to be simulated independently of running MCMC, for example by writing a nimbleFunction to do so. This means that in many cases where terminal stochastic nodes have been included in BUGS models, they are not needed when using NIMBLE.

The posterior_predictive sampler functions by calling the simulate() method of relevant node, then updating model probabilities and deterministic dependent nodes. The application of a posterior_predictive sampler to any non-terminal node will result in invalid posterior inferences. The posterior_predictive sampler will automatically be assigned to all terminal, non-data stochastic nodes in a model by the default MCMC configuration, so it is uncommon to manually assign this sampler.

Closely related, the posterior_predictive_branch is only appropriate for nodes for which all downstream (dependent) nodes are non-data; that is, the branch beginning from that node is a jointly posterior predictive network of nodes. This sampler operates by calling the simulate method of all nodes in this downstream (non-data) network, which serves to improve MCMC mixing of the posterior predictive branch.

The posterior_predictive and posterior_predictive_branch samplers accept no control list arguments.

RJ_fixed_prior sampler

This sampler proposes addition/removal for variable of interest in the framework of variable se-
lection using reversible jump MCMC, with a specified prior probability of inclusion. A normal
proposal distribution is used to generate proposals for the addition of the variable. This is a spe-
cialized sampler used by `configureRJ` function, when the model code is written without using
indicator variables. See `help{configureRJ}` for details. It is not intended for direct assignment.

RJ_indicator sampler

This sampler proposes transitions of a binary indicator variable, corresponding to a variable of
interest, in the framework of variable selection using reversible jump MCMC. This is a specialized
sampler used by `configureRJ` function, when the model code is written using indicator variables.
See `help{configureRJ}` for details. It is not intended for direct assignment.

RJ_toggled sampler

This sampler operates in the framework of variable selection using reversible jump MCMC. Specif-
ically, it conditionally performs updates of the target variable of interest using the originally-
specified sampling configuration, when variable is "in the model". This is a specialized sampler
used by `configureRJ` when adding a reversible jump MCMC . See `help{configureRJ}` for de-
tails. It is not intended for direct assignment.

Particle filter samplers

As of Version 0.10.0 of NIMBLE, the `RW_PF` and `RW_PF_block` samplers live in the 'nimbleSMC'
package. Please load that package in order to use the samplers.

Author(s)

Daniel Turek

References

Andrieu, C., Doucet, A., and Holenstein, R. (2010). Particle Markov Chain Monte Carlo Methods.
Journal of the Royal Statistical Society: Series B (Statistical Methodology), 72(3), 269-342.

Escobar, M. D., and West, M. (1995). Bayesian density estimation and inference using mixtures.
Journal of the American Statistical Association, 90(430), 577-588.

Knorr-Held, L. and Rue, H. (2003). On block updating in Markov random field models for disease
mapping. *Scandinavian Journal of Statistics*, 29, 597-614.

Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., and Teller, E. (1953). Equation
of State Calculations by Fast Computing Machines. *The Journal of Chemical Physics*, 21(6), 1087-
1092.

Murray, I., Prescott Adams, R., and MacKay, D. J. C. (2010). Elliptical Slice Sampling. *arXiv
e-prints*, arXiv:1001.0175.

Neal, R. M. (2000). Markov chain sampling methods for Dirichlet process mixture models. *Journal
of Computational and Graphical Statistics*, 9(2), 249-265.

Neal, R. M. (2003). Slice Sampling. *The Annals of Statistics*, 31(3), 705-741.

Pitt, M.K. and Shephard, N. (1999). Filtering via simulation: Auxiliary particle filters. *Journal of the American Statistical Association* 94(446), 590-599.

Roberts, G. O. and S. K. Sahu (1997). Updating Schemes, Correlation Structure, Blocking and Parameterization for the Gibbs Sampler. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 59(2), 291-317.

Shaby, B. and M. Wells (2011). *Exploring an Adaptive Metropolis Algorithm*. 2011-14. Department of Statistics, Duke University.

Stan Development Team (2020). *Stan Language Reference Manual, Version 2.22, Section 10.12*.

Tibbits, M. M., Groendyke, C., Haran, M., and Liechty, J. C. (2014). Automated Factor Slice Sampling. *Journal of Computational and Graphical Statistics*, 23(2), 543-563.

van Dyk, D.A. and T. Park. (2008). Partially collapsed Gibbs Samplers. *Journal of the American Statistical Association*, 103(482), 790-796.

See Also

configureMCMC addSampler buildMCMC runMCMC

Examples

```
## y[1] ~ dbern() or dbinom():
# mcmcConf$addSampler(target = 'y[1]', type = 'binary')

# mcmcConf$addSampler(target = 'a', type = 'RW',
#    control = list(log = TRUE, adaptive = FALSE, scale = 3))
# mcmcConf$addSampler(target = 'b', type = 'RW',
#    control = list(adaptive = TRUE, adaptInterval = 200))
# mcmcConf$addSampler(target = 'p', type = 'RW',
#    control = list(reflective = TRUE))

## a, b, and c all continuous-valued:
# mcmcConf$addSampler(target = c('a', 'b', 'c'), type = 'RW_block')

# mcmcConf$addSampler(target = 'p', type = 'RW_llFunction',
#    control = list(llFunction = RllFun, includesTarget = FALSE))

# mcmcConf$addSampler(target = 'y[1]', type = 'slice',
#    control = list(adaptive = FALSE, sliceWidth = 3))
# mcmcConf$addSampler(target = 'y[2]', type = 'slice',
#    control = list(adaptive = TRUE, sliceMaxSteps = 1))

# mcmcConf$addSampler(target = 'x[1:10]', type = 'ess')   ## x[1:10] ~ dmnorm()

# mcmcConf$addSampler(target = 'x[1:5]', type = 'RW_multinomial')   ## x[1:5] ~ dmulti()

# mcmcConf$addSampler(target = 'p[1:5]', type = 'RW_dirichlet')   ## p[1:5] ~ ddirch()

## y[1] is a posterior predictive node:
# mcmcConf$addSampler(target = 'y[1]', type = 'posterior_predictive')
```

---

setAndCalculate   *Creates a nimbleFunction for setting the values of one or more model*
                                   nodes, calculating the associated deterministic dependents and log-
                                   Prob values, and returning the total sum log-probability.

Description

This nimbleFunction generator must be specialized to any model object and one or more model
nodes. A specialized instance of this nimbleFunction will set the values of the target nodes in
the specified model, calculate the associated logProbs, calculate the values of any deterministic
dependents, calculate the logProbs of any stochastic dependents, and return the sum log-probability
associated with the target nodes and all stochastic dependent nodes.

Usage

```
setAndCalculate(model, targetNodes)

setAndCalculateDiff(model, targetNodes)
```

Arguments

model   An uncompiled or compiled NIMBLE model. This argument is required.

targetNodes   A character vector containing the names of one or more nodes or variables in
               the model. This argument is required.

Details

Calling `setAndCalculate(model,targetNodes)` or `setAndCalculate(model,targetNodes)` will
return a nimbleFunction object whose `run` function takes a single, required argument:

targetValues: A vector of numeric values which will be put into the target nodes in the specified
model object. The length of this numeric vector much exactly match the number of target nodes.

The difference between `setAndCalculate` and `setAndCalculateDiff` is the return value of their
`run` functions. In the former, `run` returns the sum of the log probabilities of the `targetNodes` with
the provided `targetValues`, while the latter returns the difference between that sum with the new
`targetValues` and the previous values in the `model`.

Author(s)

Daniel Turek

Examples

```
code <- nimbleCode({ for(i in 1:3) { x[i] ~ dnorm(0,1); y[i] ~ dnorm(0, 1)}})
Rmodel <- nimbleModel(code)
my_setAndCalc <- setAndCalculate(Rmodel, c('x[1]', 'x[2]', 'y[1]', 'y[2]'))
lp <- my_setAndCalc$run(c(1.2, 1.4, 7.6, 8.9))
```

setAndCalculateOne	*Creates a nimbleFunction for setting the value of a scalar model node, calculating the associated deterministic dependents and logProb values, and returning the total sum log-probability.*

Description

This nimbleFunction generator must be specialized to any model object and any scalar model node. A specialized instance of this nimbleFunction will set the value of the target node in the specified model, calculate the associated logProb, calculate the values of any deterministic dependents, calculate the logProbs of any stochastic dependents, and return the sum log-probability associated with the target node and all stochastic dependent nodes.

Usage

```
setAndCalculateOne(model, targetNode)
```

Arguments

model	An uncompiled or compiled NIMBLE model. This argument is required.
targetNode	The character name of any scalar node in the model object. This argument is required.

Details

Calling setAndCalculateOne(model, targetNode) will return a function with a single, required argument:

targetValue: The numeric value which will be put into the target node, in the specified model object.

Author(s)

Daniel Turek

Examples

```
code <- nimbleCode({ for(i in 1:3) x[i] ~ dnorm(0, 1) })
Rmodel <- nimbleModel(code)
my_setAndCalc <- setAndCalculateOne(Rmodel, 'x[1]')
lp <- my_setAndCalc$run(2)
```

---

setSize   *set the size of a numeric variable in NIMBLE*

Description

set the size of a numeric variable in NIMBLE. This works in R and NIMBLE, but in R it usually has no effect.

Usage

```
setSize(numObj, ..., copy = TRUE, fillZeros = TRUE)
```

Arguments

numObj	This is the object to be resized
...	sizes, provided as scalars, in order, or as a single vector
copy	logical indicating whether values should be preserved (in column-major order)
fillZeros	logical indicating whether newly allocated space should be initialized with zeros (in compiled code)

Details

Note that assigning the result of `numeric`, `integer`, `logical`, `matrix`, or `array` is often as good or better than using `setSize`. For example, 'x <- matrix(nrow = 5, ncol = 5)' is equivalent to 'setSize(x, 5, 5)' but the former allows more control over initialization.

This function is part of the NIMBLE language. Its purpose is to explicitly resize a multivariate object (vector, matrix or array), currently up to 4 dimensions. Explicit resizing is not needed when an entire object is assigned to. For example, in `Y <-A %*% B`, where A and B are matrices, `Y` will be resized automatically. Explicit resizing is necessary when assignment will be by indexed elements or blocks, if the object is not already an appropriate size for the assignment. E.g. prior to `Y[5:10] <-A %*% B`, one can use setSize to ensure that `Y` has a size (length) of at least 10.

This does work in uncompiled (R) and well as compiled execution, but in some cases it is only necessary for compiled execution. During uncompiled execution, it may not catch bugs due to resizing because some R objects will be dynamically resized during assignments anyway.

If preserving values in the resized object and/or initializing new values with 0 is not necessary, then setting these arguments to FALSE will yield slightly more efficient compiled code.

Author(s)

NIMBLE development team

setupOutputs	*Explicitly declare objects created in setup code to be preserved and compiled as member data*

Description

Normally a nimbleFunction determines what objects from setup code need to be preserved for run code or other member functions. setupOutputs allows explicit declaration for cases when an object created in setup code is not used in member functions.

Arguments

...	An arbitrary set of names

Details

Normally any object created in `setup` whose name appears in `run` or another member function is included in the saved results of setup code. When the nimbleFunction is compiled, such objects will become member data of the resulting C++ class. If it is desired to force an object to become member data even if it does not appear in a member function, declare it using setupOutputs. E.g., setupOutputs(a,b) declares that a and b should be preserved.

The setupOutputs line will be removed from the setup code. It is really a marker during nimble-Function creation of what should be preserved.

simNodes	*Basic nimbleFunctions for calculate, simulate, and getLogProb with a set of nodes*

Description

simulate, calculate, or get existing log probabilities for the current values in a NIMBLE model

Usage

```
simNodes(model, nodes)

calcNodes(model, nodes)

getLogProbNodes(model, nodes)
```

Arguments

model	A NIMBLE model
nodes	A set of nodes. If none are provided, default is all model$getNodeNames()

Details

These are basic nimbleFunctions that take a model and set of nodes and return a function that will call `calculate`, `simulate`, or `getLogProb` on those nodes. Each is equivalent to a direct call from R, but in nimbleFunction form they can be be compiled and can be put into a nimbleFunction-List. For example, `myCalc <-calcNodes(model,nodes); ans <-myCalc()` is equivalent to ans `<-model$calculate(nodes)`, but one can also do `CmyCalc <-compileNimble(myCalc)` to get a faster version.

In nimbleFunctions, for only one set of nodes, it is equivalent or slightly better to simply use `model$calculate(nodes)` in the run-time code. The compiler will process the model-nodes combination in the same way as would occur by creating a specialized `calcNodes` in the setup code. However, if there are multiple sets of nodes, one can do the following:

Setup code: `myCalcs <-nimbleFunctionList(calcNodes); myCalcs[[1]] <-calcNodes(model,nodes[[1]]); myCalcs[[2]] <-calcNodes[[2]]`

Run code: `for(i in seq_along(myCalcs)) {ans[i] <-myCalcs[[i]]()}`

Author(s)

Perry de Valpine

| simNodesMV | *Basic nimbleFunctions for using a NIMBLE model with sets of stored values* |

Description

simulate, calculate, or get the existing log probabilities for values in a modelValues object using a NIMBLE model

Usage

```
simNodesMV(model, mv, nodes)

calcNodesMV(model, mv, nodes)

getLogProbNodesMV(model, mv, nodes)
```

Arguments

model	A nimble model.
mv	A modelValues object in which multiple sets of model variables and their corresponding logProb values are or will be saved. mv must include the nodes provided
nodes	A set of nodes. If none are provided, default is all `model$getNodeNames()`

Details

simNodesMV simulates values in the given nodes and saves them in `mv`. `calcNodesMV` calculates these nodes for each row of `mv` and returns a vector of the total log probabilities (densities) for each row. `getLogProbNodesMV` is like `calcNodesMV` without actually doing the calculations.

Each of these will expand variables or index blocks and topologically sort them so that each node's parent nodes are processed before itself.

`getLogProbMV` should be used carefully. It is generally for situations where the logProb values are guaranteed to have already been calculated, and all that is needed is to query them. The risk is that a program may have changed the values in the nodes, in which case `getLogProbMV` would collect logProb values that are out of date with the node values.

Value

from `simNodesMV`: NULL. from `calcNodesMV` and `getLogProbMV`: a vector of the sum of log probabilities (densities) from any stochastic nodes in `nodes`.

Run time arguments

- `m`

  (`simNodesMV` only). Number of simulations requested.

- `saveLP`

  (`calcNodesMV`only). Whether to save the logProb values in `mv`. Should be given as `TRUE` unless there is a good reason not to.

Author(s)

Clifford Anderson-Bergman

Examples

```
code <- nimbleCode({
for(i in 1:5)
x[i] ~ dnorm(0,1)
})

myModel <- nimbleModel(code)
myMV <- modelValues(myModel)

Rsim <- simNodesMV(myModel, myMV)
Rcalc <- calcNodesMV(myModel, myMV)
Rglp <- getLogProbNodesMV(myModel, myMV)
## Not run:
  cModel <- compileNimble(myModel)
  Csim <- compileNimble(Rsim, project = myModel)
  Ccalc <- compileNimble(Rcalc, project = myModel)
  Cglp <- compileNimble(Rglp, project = myModel)
  Csim$run(10)
  Ccalc$run(saveLP = TRUE)
  Cglp$run() #Gives identical answers to Ccalc because logProbs were saved
```

```
   Csim$run(10)
   Ccalc$run(saveLP = FALSE)
   Cglp$run()   #Gives wrong answers because logProbs were not saved

## End(Not run)
```

---

singleVarAccessClass-class
                              Class singleVarAccessClass

Description

Classes used internally in NIMBLE and not expected to be called directly by users.

StickBreakingFunction   *The Stick Breaking Function*

Description

Computes probabilities based on stick breaking construction.

Usage

```
stick_breaking(z, log = 0)
```

Arguments

z	vector argument.
log	logical; if TRUE, weights are returned on the log scale.

Details

The stick breaking function produces a vector of probabilities that add up to one, based on a series of individual probabilities in z, which define the breaking points relative to the remaining stick length. The first element of z determines the first probability based on breaking a proportion z[1] from a stick of length one. The second element of z determines the second probability based on breaking a proportion z[2] from the remaining stick (of length 1-z[1]), and so forth. Each element of z should be in $(0, 1)$. The returned vector has length equal to the length of z plus 1. If z[k] is equal to 1 for any k, then the returned vector has length smaller than z. If one of the components is smaller than 0 or greater than 1, NaNs are returned.

Author(s)

Claudia Wehrhahn

References

Sethuraman, J. (1994). A constructive definition of Dirichlet priors. *Statistica Sinica*, 639-650.

Examples

```
z <- rbeta(5, 1, 1)
stick_breaking(z)

## Not run:
cstick_breaking <- compileNimble(stick_breaking)
cstick_breaking(z)

## End(Not run)
```

---

svdNimbleList   *svdNimbleList definition*

Description

nimbleList definition for the type of nimbleList returned by [nimSvd](#).

Usage

```
svdNimbleList
```

Format

An object of class list of length 1.

Author(s)

NIMBLE development team

See Also

[nimSvd](#)

---

t *The t Distribution*

Description

Density, distribution function, quantile function and random generation for the t distribution with df degrees of freedom, allowing non-zero location, mu, and non-unit scale, sigma

Usage

```
dt_nonstandard(x, df = 1, mu = 0, sigma = 1, log = FALSE)

rt_nonstandard(n, df = 1, mu = 0, sigma = 1)

pt_nonstandard(q, df = 1, mu = 0, sigma = 1, lower.tail = TRUE, log.p = FALSE)

qt_nonstandard(p, df = 1, mu = 0, sigma = 1, lower.tail = TRUE, log.p = FALSE)
```

Arguments

x	vector of values.
df	vector of degrees of freedom values.
mu	vector of location values.
sigma	vector of scale values.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations.
q	vector of quantiles.
lower.tail	logical; if TRUE (default) probabilities are $P[X \leq x]$; otherwise, $P[X > x]$.
log.p	logical; if TRUE, probabilities p are given by user as log(p).
p	vector of probabilities.

Details

See Gelman et al., Appendix A or the BUGS manual for mathematical details.

Value

dt_nonstandard gives the density, pt_nonstandard gives the distribution function, qt_nonstandard gives the quantile function, and rt_nonstandard generates random deviates.

Author(s)

Christopher Paciorek

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

[Distributions](#) for other standard distributions

Examples

```
x <- rt_nonstandard(50, df = 1, mu = 5, sigma = 1)
dt_nonstandard(x, 3, 5, 1)
```

---

testBUGSmodel   *Tests BUGS examples in the NIMBLE system*

Description

`testBUGSmodel` builds a BUGS model in the NIMBLE system and simulates from the model, comparing the values of the nodes and their log probabilities in the uncompiled and compiled versions of the model

Usage

```
testBUGSmodel(
  example = NULL,
  dir = NULL,
  model = NULL,
  data = NULL,
  inits = NULL,
  useInits = TRUE,
  expectModelWarning = FALSE,
  debug = FALSE,
  verbose = nimbleOptions("verbose")
)
```

Arguments

example   (optional) example character vector indicating name of BUGS example to test; can be null if `model` is provided

dir   (optional) character vector indicating directory in which files are contained, by default the classic-bugs directory if the installed package is used; to use the current working directory, set this to ""

model	(optional) one of (1) a character string giving the file name containing the BUGS model code, (2) an R function whose body is the BUGS model code, or (3) the output of `nimbleCode`. If a file name, the file can contain a 'var' block and 'data' block in the manner of the JAGS versions of the BUGS examples but should not contain references to other input data files nor a const block. The '.bug' or '.txt' extension can be excluded.
data	(optional) one of (1) character string giving the file name for an R file providing the input constants and data as R code [assigning individual objects or as a named list] or (2) a named list providing the input constants and data. If neither is provided, the function will look for a file named `example-data` including extensions .R, .r, or .txt.
inits	(optional) (1) character string giving the file name for an R file providing the initial values for parameters as R code [assigning individual objects or as a named list] or (2) a named list providing the values. If neither is provided, the function will look for a file named `example-init` or `example-inits` including extensions .R, .r, or .txt.
useInits	boolean indicating whether to test model with initial values provided via `inits`.
expectModelWarning	
	boolean indicating whether `nimbleModel` is expected to produce a warning or character string giving part of expected warning.
debug	logical indicating whether to put the user in a browser for debugging when `testBUGSmodel` calls `readBUGSmodel`. Intended for developer use.
verbose	logical indicating whether to print additional logging information

Details

Note that testing without initial values may cause warnings when parameters are sampled from improper or fat-tailed distributions

Author(s)

Christopher Paciorek

Examples

```
## Not run:
testBUGSmodel('pump')

## End(Not run)
```

---

valueInCompiledNimbleFunction

get or set value of member data from a compiled nimbleFunction using a multi-interface

---

Description

Most nimbleFunctions written for direct user interaction allow standard R-object-like access to member data using $ or `[[`. However, sometimes compiled nimbleFunctions contained within other compiled nimbleFunctions are interfaced with a light-weight system called a multi-interface. valueInCompiledNimbleFunction provides a way to get or set values in such cases.

Usage

```
valueInCompiledNimbleFunction(cnf, name, value)
```

Arguments

cnf	Compiled nimbleFunction object
name	Name of the member data
value	If provided, the value to assign to the member data. If omitted, the value of the member data is returned.

Details

The member data of a nimbleFunction are the objects created in setup code that are used in run code or other member functions.

Whether multi-interfaces are used for nested nimbleFunctions is controlled by the buildInterfacesForCompiledNestedNim option in [nimbleOptions](#).

To see an example of a multi-interface, see samplerFunctions in a compiled MCMC interface object.

Author(s)

Perry de Valpine

---

values *Access or set values for a set of nodes in a model*

Description

Get or set values for a set of nodes in a model

Usage

```
values(model, nodes, accessorIndex)

values(model, nodes, accessorIndex) <- value
```

Arguments

`model`	a NIMBLE model object, either compiled or uncompiled
`nodes`	a vector of node names, allowing index blocks that will be expanded
`accessorIndex`	For internal NIMBLE use only
`value`	value to set the node(s) to

Details

Access or set values for a set of nodes in a NIMBLE model.

Calling `values(model,nodes)` returns a vector of the concatenation of values from the nodes requested P `<-values(model,nodes)` is a newer syntax for `getValues(P,model,values)`, which still works and modifies P in the calling environment.

Calling `values(model,nodes) <-P` sets the value of the nodes in the model, in sequential order from the vector P.

In both uses, when requested nodes are from matrices or arrays, the values will be handled following column-wise order.

The older function `getValues(P,model,nodes)` is equivalent to P `<-values(model,nodes)`, and the older function `setValues(P,model,nodes)` is equivalent to `values(model,nodes) <-P`

These functions work in R and in NIMBLE run-time code that can be compiled.

Value

A vector of values concatenated from the provided nodes in the model

Author(s)

NIMBLE development team

---

waic	*Using WAIC*

Description

Details of the WAIC measure for comparing models. NIMBLE implements an online WAIC algorithm, computed during the course of the MCMC iterations.

Details

To use WAIC, set `enableWAIC = TRUE` when configuring or (if not using `configureMCMC` building an MCMC) and set `WAIC = TRUE` when calling `nimbleMCMC` and optionally when calling `runMCMC`.

By default, NIMBLE calculates WAIC using an online algorithm that updates required summary statistics at each post-burnin iteration of the MCMC.

One can also use `calculateWAIC` to run an offline version of the WAIC algorithm after all MCMC sampling has been done. This allows calculation of WAIC from a matrix (or dataframe) of posterior

samples and also retains compatibility with WAIC in versions of NIMBLE before 0.12.0. However, the offline algorithm is less flexible than the online algorithm and only provides conditional WAIC without the ability to group data points. See `help(calculateWAIC)` for details.

`controlWAIC` **list**

The `controlWAIC` argument is a list that controls the behavior of the WAIC algorithm and is passed to either `configureMCMC` or (if not using `configureMCMC`) `buildMCMC`. One can supply any of the following optional components:

`online`: Logical value indicating whether to calculate WAIC during the course of the MCMC. Default is `TRUE` and setting to `FALSE` is primarily for backwards compatibility to allow use of the old `calculateWAIC` method that calculates WAIC from monitored values after the MCMC finishes.

`dataGroups`: Optional list specifying grouping of data nodes, one element per group, with each list element containing the node names for the data nodes in that group. If provided, the predictive density values computed will be the joint density values, one joint density per group. Defaults to one data node per 'group'. See details.

`marginalizeNodes`: Optional set of nodes (presumably latent nodes) over which to marginalize to compute marginal WAIC (i.e., WAIC based on a marginal likelihood), rather than the default conditional WAIC (i.e., WAIC conditioning on all parent nodes of the data nodes). See details.

`niterMarginal`: Number of Monte Carlo iterations to use when marginalizing (default is 1000).

`convergenceSet`: Optional vector of numbers between 0 and 1 that specify a set of shorter Monte Carlo simulations for marginal WAIC calculation as fractions of the full (`niterMarginal`) Monte Carlo simulation. If not provided, NIMBLE will use 0.25, 0.50, and 0.75. NIMBLE will report the WAIC, lppd, and pWAIC that would have been obtained for these smaller Monte Carlo simulations, allowing assessment of the number of Monte Carlo samples needed for stable calculation of WAIC.

`thin`: Logical value for specifying whether to do WAIC calculations only on thinned samples (default is `FALSE`). Likely only useful for reducing computation when using marginal WAIC.

Extracting WAIC

The calculated WAIC and related quantities can be obtained in various ways depending on how the MCMC is run. If using `nimbleMCMC` and setting `WAIC = TRUE`, see the `WAIC` component of the output list. If using `runMCMC` and setting `WAIC = TRUE`, either see the `WAIC` component of the output list or use the `getWAIC` method of the MCMC object (in the latter case `WAIC = TRUE` is not required). If using the `run` method of the MCMC object, use the `getWAIC` method of the MCMC object.

The output of running WAIC (unless one sets `online = FALSE`) is a list containing the following components:

`WAIC`: The computed WAIC, on the deviance scale. Smaller values are better when comparing WAIC for two models.

`lppd`: The log predictive density component of WAIC.

`pWAIC`: The pWAIC estimate of the effective number of parameters, computed using the *p*WAIC2 method of Gelman et al. (2014).

To get further information, one can use the `getWAICdetails` method of the MCMC object. The result of running `getWAICdetails` is a list containing the following components:

`marginal`: Logical value indicating whether marginal (`TRUE`) or conditional (`FALSE`) WAIC was calculated.

`niterMarginal`: Number of Monte Carlo iterations used in computing marginal likelihoods if using marginal WAIC.

`thin`: Whether WAIC was calculated based only on thinned samples.

`online`: Whether WAIC was calculated during MCMC sampling.

`WAIC_partialMC`, `lppd_partialMC`, `pWAIC_partialMC`: The computed marginal WAIC, lppd, and pWAIC based on fewer Monte Carlo simulations, for use in assessing the sensitivity of the WAIC calculation to the number of Monte Carlo iterations.

`niterMarginal_partialMC`: Number of Monte Carlo iterations used for the values in `WAIC_partialMC`, `lppd_partialMC`, `pWAIC_partialMC`.

`WAIC_elements`, `lppd_elements`, `pWAIC_elements`: Vectors of individual WAIC, lppd, and pWAIC values, one element per data node (or group of nodes in the case of specifying `dataGroups`). Of use in computing the standard error of the difference in WAIC between two models, following Vehtari et al. (2017).

Online WAIC

As of version 0.12.0, NIMBLE provides enhanced WAIC functionality, with user control over whether to use conditional or marginal versions of WAIC and whether to group data nodes. In addition, users are no longer required to carefully choose MCMC monitors. WAIC by default is now calculated in an online manner (updating the required summary statistics at each MCMC iteration), using all post-burnin samples. The WAIC (Watanabe, 2010) is calculated from Equations 5, 12, and 13 in Gelman et al. (2014) (i.e., using 'pWAIC2').

Note that there is not a unique value of WAIC for a model. By default, WAIC is calculated conditional on the parent nodes of the data nodes, and the density values used are the individual density values of the data nodes. However, by modifying the `marginalizeNodes` and `dataGroups` elements of the control list, users can request a marginal WAIC (using a marginal likelihood that integrates over user-specified latent nodes) and/or a WAIC based on grouping observations (e.g., all observations in a cluster) to use joint density values. See the MCMC Chapter of the NIMBLE User Manual for more details.

For more detail on the use of different predictive distributions, see Section 2.5 from Gelman et al. (2014) or Ariyo et al. (2019).

Note that based on a limited set of simulation experiments in Hug and Paciorek (2021) our tentative recommendation is that users only use marginal WAIC if also using grouping.

Author(s)

Joshua Hug and Christopher Paciorek

References

Watanabe, S. (2010). Asymptotic equivalence of Bayes cross validation and widely applicable information criterion in singular learning theory. *Journal of Machine Learning Research* 11: 3571-3594.

Gelman, A., Hwang, J. and Vehtari, A. (2014). Understanding predictive information criteria for Bayesian models. *Statistics and Computing* 24(6): 997-1016.

Ariyo, O., Quintero, A., Munoz, J., Verbeke, G. and Lesaffre, E. (2019). Bayesian model selection in linear mixed models for longitudinal data. *Journal of Applied Statistics* 47: 890-913.

Vehtari, A., Gelman, A. and Gabry, J. (2017). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing* 27: 1413-1432.

Hug, J.E. and Paciorek, C.J. (2021). A numerically stable online implementation and exploration of WAIC through variations of the predictive density, using NIMBLE. *arXiv e-print* <arXiv:2106.13359>.

See Also

calculateWAIC configureMCMC buildMCMC runMCMC nimbleMCMC

Examples

```
code <- nimbleCode({
  for(j in 1:J) {
    for(i in 1:n)
      y[j, i] ~ dnorm(mu[j], sd = sigma)
    mu[j] ~ dnorm(mu0, sd = tau)
  }
  sigma ~ dunif(0, 10)
  tau ~ dunif(0, 10)
})
J <- 5
n <- 10
groups <- paste0('y[', 1:J, ', 1:', n, ']')
y <- matrix(rnorm(J*n), J, n)
Rmodel <- nimbleModel(code, constants = list(J = J, n = n), data = list(y = y),
                      inits = list(tau = 1, sigma = 1))

## Various versions of WAIC available via online calculation.
## Conditional WAIC without data grouping:
conf <- configureMCMC(Rmodel, enableWAIC = TRUE)
## Conditional WAIC with data grouping
conf <- configureMCMC(Rmodel, enableWAIC = TRUE, controlWAIC = list(dataGroups = groups))
## Marginal WAIC with data grouping:
conf <- configureMCMC(Rmodel, enableWAIC = TRUE, controlWAIC =
            list(dataGroups = groups, marginalizeNodes = 'mu'))
## Not run:
Rmcmc <- buildMCMC(conf)
Cmodel <- compileNimble(Rmodel)
Cmcmc <- compileNimble(Rmcmc, project = Rmodel)
output <- runMCMC(Cmcmc, niter = 1000, WAIC = TRUE)
output$WAIC               # direct access
## Alternatively call via the `getWAIC` method; this doesn't require setting
## `waic=TRUE` in `runMCMC`
Cmcmc$getWAIC()
Cmcmc$getWAICdetails()

## End(Not run)
```

| waicDetailsList | *waicDetailsList definition* |

Description

waicDetailsList definition for the nimbleList type returned by WAIC computation.

Usage

```
waicDetailsList
```

Format

An object of class list of length 1.

Details

See help(waic) for details on the elements of the list.

Author(s)

NIMBLE development team

| waicList | *waicList definition* |

Description

waicList definition for the nimbleList type returned by WAIC computation.

Usage

```
waicList
```

Format

An object of class list of length 1.

Details

See help(waic) for details on the elements of the list.

Author(s)

NIMBLE development team

Wishart   *The Wishart Distribution*

Description

Density and random generation for the Wishart distribution, using the Cholesky factor of either the scale matrix or the rate matrix.

Usage

```
dwish_chol(x, cholesky, df, scale_param = TRUE, log = FALSE)

rwish_chol(n = 1, cholesky, df, scale_param = TRUE)
```

Arguments

x	vector of values.
cholesky	upper-triangular Cholesky factor of either the scale matrix (when `scale_param` is TRUE) or rate matrix (otherwise).
df	degrees of freedom.
scale_param	logical; if TRUE the Cholesky factor is that of the scale matrix; otherwise, of the rate matrix.
log	logical; if TRUE, probability density is returned on the log scale.
n	number of observations (only n=1 is handled currently).

Details

See Gelman et al., Appendix A or the BUGS manual for mathematical details. The rate matrix as used here is defined as the inverse of the scale matrix, S^{-1}, given in Gelman et al.

Value

`dwish_chol` gives the density and `rwish_chol` generates random deviates.

Author(s)

Christopher Paciorek

References

Gelman, A., Carlin, J.B., Stern, H.S., and Rubin, D.B. (2004) *Bayesian Data Analysis*, 2nd ed. Chapman and Hall/CRC.

See Also

[Distributions](#) for other standard distributions

Examples

```
df <- 40
ch <- chol(matrix(c(1, .7, .7, 1), 2))
x <- rwish_chol(1, ch, df = df)
dwish_chol(x, ch, df = df)
```

---

withNimbleOptions   *Temporarily set some NIMBLE options.*

Description

Temporarily set some NIMBLE options.

Usage

```
withNimbleOptions(options, expr)
```

Arguments

options	a list of options suitable for `nimbleOptions`.
expr	an expression or statement to evaluate.

Value

expr as evaluated with given options.

Examples

```
## Not run:
if (!(getNimbleOption('showCompilerOutput') == FALSE)) stop()
nf <- nimbleFunction(run = function(){ return(0); returnType(double()) })
cnf <- withNimbleOptions(list(showCompilerOutput = TRUE), {
    if (!(getNimbleOption('showCompilerOutput') == TRUE)) stop()
    compileNimble(nf)
})
if (!(getNimbleOption('showCompilerOutput') == FALSE)) stop()

## End(Not run)
```

Index